

# Databases

## Cap. 9. SQL Execution Plan. Introduction to Query Optimization



Textbook: Ramakrishnan, Gehrke, "Database Management Systems", McGraw Hill, 2003

2017 UPT

Conf.Dr. Dan Pescaru

# Relational operators

---

1. Useful for SQL query decomposition
  - Used for representing a query execution plan (close to expression tree in compiler lexical analysis)
2. A query is applied to relation instances
3. The result of a query is also a relation instance
4. The schema for the result of a given query is fixed

# Basic operators

---

## 1. Selection ( $\sigma$ )

- Selects a subset of rows from a relation

## 2. Projection ( $\pi$ )

- Removes unwanted columns from a relation

## 3. Cross-product ( $\times$ )

- Combine two relations

## 4. Set-difference ( $-$ )

- Tuples in Relation 1, but not in Relation 2

## 5. Union ( $\cup$ )

- Tuples in Relation 1 and in Relation 2

# Additional operators

---

## 1. Join ( $\circ$ )

- Corresponding tuples based on a relationship between certain columns in relations

## 2. Intersection ( $\cap$ )

- Tuples that exists in both relations

## 3. Renaming

- Change an attribute name

## 4. Since each operation returns a relation, operations can be composed. Relational algebra is “closed”

# Query evaluation

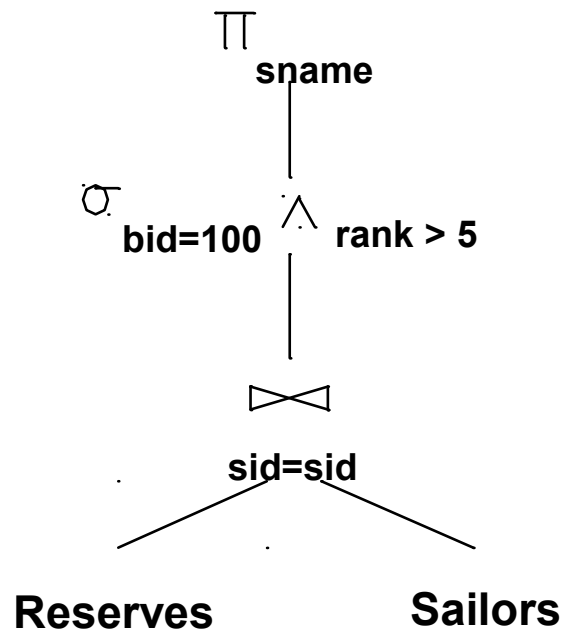
---

1. Plan: Tree of relational algebra operators, with choice of algorithms for each of them
2. Two main issues in query optimization:
  - a. For a query, what plans are considered?
  - b. Algorithm to search plan space for cheapest (estimated) plan
3. How is the cost of a plan estimated?
4. Ideally: Want to find best plan.  
Practically: Avoid worst plans!



# Execution plan example

- E.g. **SELECT** sname **FROM** Sailors s  
**INNER JOIN** Reserves r **ON** s.sid=r.sid  
**WHERE** r.bid=100 **AND** s.rank>5



# Basic optimization techniques

---

1. Algorithms for evaluating relational operators use some simple ideas
2. Iteration: sometimes, faster to scan all tuples even if there is an index. Sometimes, faster to scan the data entries in an index instead of the table
3. Indexing: can use WHERE conditions to retrieve small set of tuples ( $\sigma, \circ$ )
4. Partitioning: by using sorting or hashing, we can partition the input tuples and replace an expensive operation by similar operations on smaller inputs

# Statistics and catalogs

---

1. Catalogs contain metadata and statistics:
  - # tuples (NTuples) and # pages (NPages) for each relation
  - # distinct key values (NKeys) and NPages for each index
  - Index height, Low/High key values for each tree index
2. Catalogs updated periodically (efficiency)
3. Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok



# Cost estimation

---

1. For each execution plan considered, must estimate cost

- Cost of each operation in plan tree
  - Depends on input cardinalities
  - Depends on tuple size

2. Must also estimate size of result for each operation in tree!

- Use information about the input relations
- For selections and joins, assume independence of predicates

# Access paths

---

1. An access path is a method of retrieving tuples
  - File scan, or index that matches a selection (in the query)
2. A tree index matches a conjunction of terms that involve only attributes in a prefix of the search key
  - E.g., Tree index on  $\langle a, b, c \rangle$  matches the selection  $a=5$  AND  $b=3$ , and  $a=5$  AND  $b>6$ , but not  $b=3$
3. A hash index matches a conjunction of terms that has a term attribute = value for every attribute in the search key of the index
  - E.g., Hash index on  $\langle a, b, c \rangle$  matches  $a=5$  AND  $b=3$  AND  $c=5$ ; but it does not match  $b=3$  or  $a>5$  AND  $b=3$  AND  $c=5$

# Optimizing selection

---

1. Find the most selective access path, retrieve tuples using it, and apply any remaining terms that don't match the index:
  1. Most selective access path: an index or file scan that we estimate will require the fewest page I/Os
  2. Terms that match this index reduce the number of tuples retrieved; other terms are used to discard some retrieved tuples, but do not affect number of tuples/pages fetched

## Optimizing selection example

---

1. Consider:  $\text{day} < 8/9/94$  AND  $\text{bid} = 5$  AND  $\text{sid} = 3$
2. A B+ tree index on  $\text{day}$  can be used
  - $\text{bid} = 5$  and  $\text{sid} = 3$  must be checked for each retrieved tuple
3. Similarly, a hash index on  $\langle \text{bid}, \text{sid} \rangle$  could be used
  - $\text{day} < 8/9/94$  must then be checked

# Optimizing projection

---

1. The expensive part is removing duplicates
  - SQL systems do not remove duplicates unless the keyword **DISTINCT** is specified in a query
2. *Sorting* Approach: sort on keys and remove duplicates. Can optimize this by dropping unwanted information while sorting
3. *Hashing* Approach: Hash on keys to create partitions. Load partitions into memory one at a time, build in-memory hash structure, and eliminate duplicates



# Projection: sorting approach

---

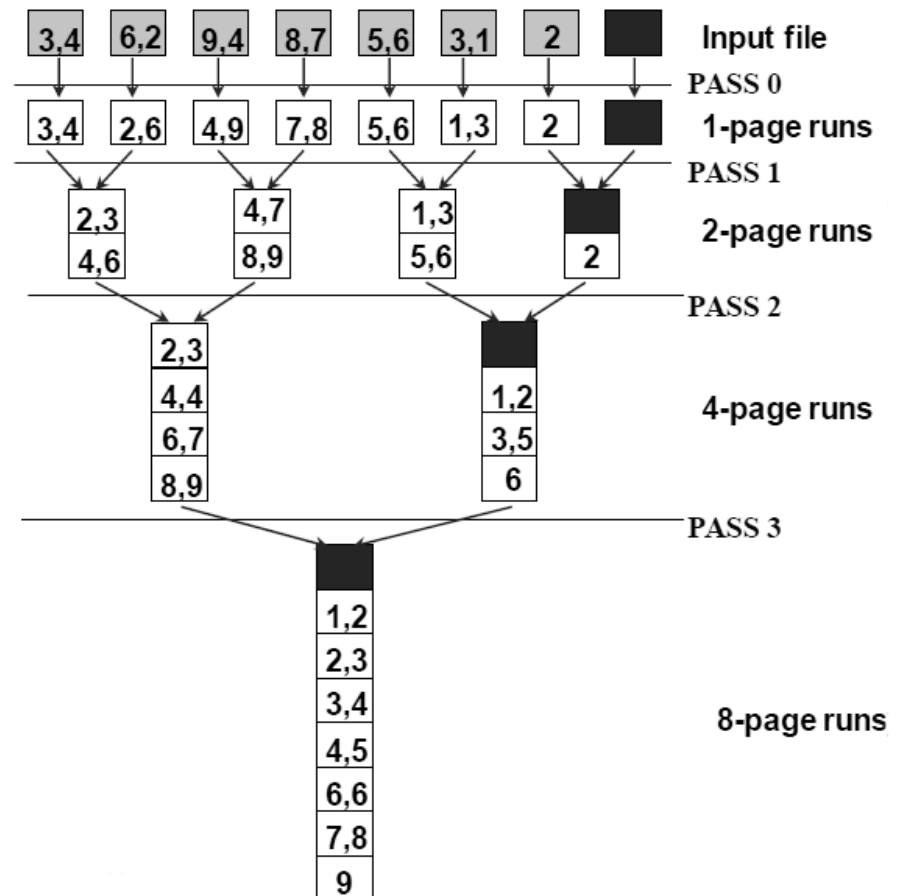
**SELECT DISTINCT** sid, bid **FROM** Reserves

## 1. An approach based on sorting:

- Two way external merge sort with elimination of unwanted fields
- Tuples in result are smaller than input tuples (size ratio depends on # and size of fields that are dropped)
- Modify merging passes to eliminate duplicates. Thus, number of result tuples is smaller than input (difference depends on # of duplicates)

# Two way external merge sort

1. Idea: sort partition files and merge (using divide and conquer), three buffers needed
2. In each pass: read + write each page
3. Total cost is approx.  
 $2N \times \log_2 N$



# Projection: hashing approach

---

**SELECT DISTINCT** sid, bid **FROM** Reserves

1. Partitioning phase using B buffers: Read R using one input buffer. For each tuple, discard unwanted fields, apply hash function h1 to choose one of B-1 output buffers
  - Result is B-1 partitions (of tuples with no unwanted fields). Two tuples from different partitions guaranteed to be distinct
2. Duplicate elimination phase: For each partition, read it and build an in-memory hash table, using hash function h2 ( $\neq$  h1) on all fields, while discarding duplicates by checking the equal hash-value

# Join: Index Nested Loops

---

SELECT r.\*, s.\*

FROM Reserves r, Sailors s WHERE r.sid=s.sid

1. Without an index: approx.  $\text{Rows}_{R1} \times \text{Pages}_{R2}$  I/Os (for  $50000 \times 1000 = 70h$  at  $5\text{ms}/(\text{I/O})$  on average 7200rpm HDD)!
2. If there is an index on the join column of one relation, can make it the inner and exploit the index
  - Cost:  $\text{Rows}_{R1} \times \text{cost of finding matching R2 tuples}$
3. For each R1 tuple, cost of probing R2 index is about 1.2 for hash index, 2-4 for B+ tree. Cost of then finding R2 tuples depends on clustering
4. Clustered index: 0 I/O
5. Unclustered: 1 I/O per matching S tuple

# Set operation optimization

---

1. Intersect and Cross-Product special cases of join
2. Sorting based approach to Union:
  - Sort both relations (on combination of all attributes)
  - Scan sorted relations and merge them
3. Hash based approach to Union:
  - Partition R1 and R2 using hash function  $h$
  - For each R2-partition, build in-memory hash table (using  $h2$ ), scan correlated R1-partition and add tuples to table while discarding duplicates



# Aggregate operators optimization (I)

---

## 1. Without grouping:

- In general, requires scanning the relation
- Given index whose search key includes all attributes in the **SELECT** or **WHERE** clauses, can do index-only scan

# Aggregate operators optimization

---

## 1. With grouping:

- Sort on group-by attributes, then scan relation and compute aggregate for each group. (Can improve upon this by combining sorting and aggregate computation)
- Similar approach based on hashing on group-by attributes
- Given tree index whose search key includes all attributes in **SELECT**, **WHERE** and **GROUP BY** clauses, can do index-only scan; if group-by attributes form prefix of search key, can retrieve data entries/tuples in group-by order and apply **HAVING** same time

# Query optimization example

