



# Expert Systems


---

**PROFESSOR DARIAN M. ONCHIS**

[HTTPS://STAFF.FMI.UVT.RO/~DARIAN.ONCHIS/](https://staff.fmi.uvt.ro/~darian.onchis/)

EUROPEAN LABORATORY FOR LEARNING AND INTELLIGENT SYSTEMS (ELLIS)

DARIAN.ONCHIS@E-UVT.RO



**SIMT** ≡ Signal, Image and Machine Learning Team




Is there a cat in this image?




Suppose we have a 'cat-filter'  $W$

Suppose we have a 'cat-filter'  $W$




Suppose we have a 'cat-filter'  $W$



## C-S Inequality

Suppose we have a 'cat-filter'  $W$



### C-S Inequality

For any


$$X = \begin{pmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \end{pmatrix} \text{ we}$$

have

$$X \cdot W \leq (X \cdot X)^{1/2} (W \cdot W)^{1/2}$$

with equality if and only if  $X$  is parallel to a cat.

Suppose we have a 'cat-filter'  $W$



### (Cat-Selection) C-S Inequality

For any

$$X = \begin{pmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \end{pmatrix} \text{ we}$$


have

$$X \cdot W \leq (X \cdot X)^{1/2} (W \cdot W)^{1/2}$$

with equality if and only if  $X$  is parallel to a cat.



Suppose we have a 'cat-filter'  $W$



### (Cat-Selection) C-S Inequality

For any

$$X = \begin{pmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \end{pmatrix} \text{ we}$$

have


$$X \cdot W \leq (X \cdot X)^{1/2} (W \cdot W)^{1/2}$$


with equality if and only if  $X$  is parallel to a cat.




perform cat-test on all  $3 \times 3$  image subpatches!











































# Convolution

## Definition

Suppose that  $X, Y \in \mathbb{R}^{n \times n}$ . Then  $Z = X * Y \in \mathbb{R}^{n \times n}$  is defined as

$$Z[i, j] = \sum_{k, l=0}^{n-1} X[i - k, j - l] Y[k, l],$$

where periodization or zero-padding of  $X, Y$  is used if  $i - k$  or  $j - l$  is not in  $\{0, \dots, n - 1\}$ .

# Convolution

## Definition

Suppose that  $X, Y \in \mathbb{R}^{n \times n}$ . Then  $Z = X * Y \in \mathbb{R}^{n \times n}$  is defined as

$$Z[i, j] = \sum_{k, l=0}^{n-1} X[i - k, j - l] Y[k, l],$$

where periodization or zero-padding of  $X, Y$  is used if  $i - k$  or  $j - l$  is not in  $\{0, \dots, n - 1\}$ .




Efficient computation possible via FFT (or directly if  $X$  or  $Y$  are sparse)!

## Example: Detecting Vertical Edges


Given 'vertical-edge-detection-filter'  $W = \begin{pmatrix} -1 & 2 & -1 \\ -1 & 2 & -1 \\ -1 & 2 & -1 \end{pmatrix}$

## Example: Detecting Vertical Edges


Given 'vertical-edge-detection-filter'  $W = \begin{pmatrix} -1 & 2 & -1 \\ -1 & 2 & -1 \\ -1 & 2 & -1 \end{pmatrix}$



## Example: Detecting Vertical Edges




## Example: Detecting Vertical Edges



\*




## Example: Detecting Vertical Edges





## Example: Detecting Vertical Edges




# Introducing Convolutional Nodes

- A convolutional node accepts as input a stack of images, e.g.  
 $X \in \mathbb{R}^{n_1 \times n_2 \times S}$ .

# Introducing Convolutional Nodes

- A convolutional node accepts as input a stack of images, e.g.  
 $X \in \mathbb{R}^{n_1 \times n_2 \times S}$ .



# Introducing Convolutional Nodes

- A convolutional node accepts as input a stack of images, e.g.  $X \in \mathbb{R}^{n_1 \times n_2 \times S}$ .
- Given a filter  $W \in \mathbb{R}^{F \times F \times S}$ , where  $F$  is the *spatial extent* and a *bias*  $b \in \mathbb{R}$ , it computes a matrix

$$Z = W *_{12} X := \sum_{i=1}^S X[:, :, i] * W[:, :, i] + b.$$

# Introducing Convolutional Layers

- A convolutional node accepts as input a stack of images, e.g.  $X \in \mathbb{R}^{n_1 \times n_2 \times S}$ .
- Given a filter  $W \in \mathbb{R}^{F \times F \times S}$ , where  $F$  is the *spatial extent* and a *bias*  $b \in \mathbb{R}$ , it computes a matrix

$$Z = W *_{12} X := \sum_{i=1}^S X[:, :, i] * W[:, :, i] + b.$$

- A convolutional layer consists of  $K$  convolutional nodes  $((W_i, b_i))_{i=1}^K \subset \mathbb{R}^{F \times F \times S} \times \mathbb{R}$  and produces as output a stack  $Z \in \mathbb{R}^{n_1 \times n_2 \times K}$  via

$$Z[:, :, i] = W_i *_{12} X + b_i.$$

# Introducing Convolutional Layers

- A convolutional node accepts as input a stack of images, e.g.  $X \in \mathbb{R}^{n_1 \times n_2 \times S}$ .
- Given a filter  $W \in \mathbb{R}^{F \times F \times S}$ , where  $F$  is the *spatial extent* and a *bias*  $b \in \mathbb{R}$ , it computes a matrix

$$Z = W *_{12} X := \sum_{i=1}^S X[:, :, i] * W[:, :, i] + b.$$

- A convolutional layer consists of  $K$  convolutional nodes  $((W_i, b_i))_{i=1}^K \subset \mathbb{R}^{F \times F \times S} \times \mathbb{R}$  and produces as output a stack  $Z \in \mathbb{R}^{n_1 \times n_2 \times K}$  via

$$Z[:, :, i] = W_i *_{12} X + b_i.$$



A convolutional layer can be written as a conventional neural network layer!

# Activation Layers

The activation layer is defined in the same way as before, e.g.,  $Z \in \mathbb{R}^{n_1 \times n_2 \times K}$  is mapped to

$$A = \text{ReLU}(Z)$$

where ReLU is applied component-wise.

# Pooling Layers



Reduce dimensionality after filtering.



# Pooling Layers



Reduce dimensionality after filtering.

## Definition

A *pooling operator*  $\mathbf{R}$  acts layer-wise on a tensor  $X \in \mathbb{R}^{n_1 \times n_2 \times S}$  to result in a tensor  $\mathbf{R}(X) \in \mathbb{R}^{m_1 \times m_2 \times S}$ , where  $m_1 < n_1$  and  $m_2 < n_2$ .


# Pooling Layers



Reduce dimensionality after filtering.

## Definition

A *pooling operator*  $\mathbf{R}$  acts layer-wise on a tensor  $X \in \mathbb{R}^{n_1 \times n_2 \times S}$  to result in a tensor  $\mathbf{R}(X) \in \mathbb{R}^{m_1 \times m_2 \times S}$ , where  $m_1 < n_1$  and  $m_2 < n_2$ .



Downsampling

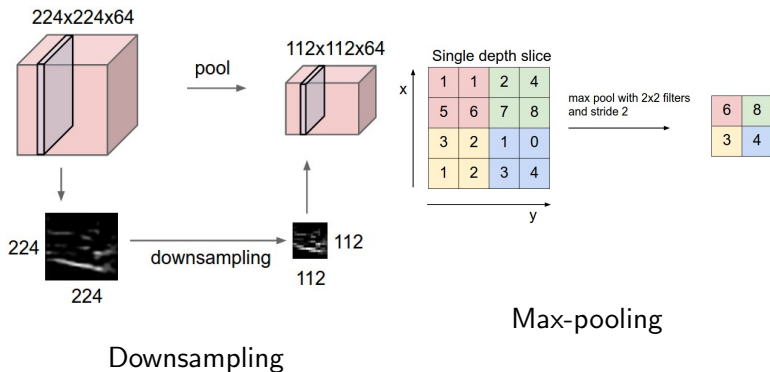
# Pooling Layers



Reduce dimensionality after filtering.

## Definition

A *pooling operator*  $\mathbf{R}$  acts layer-wise on a tensor  $X \in \mathbb{R}^{n_1 \times n_2 \times S}$  to result in a tensor  $\mathbf{R}(X) \in \mathbb{R}^{m_1 \times m_2 \times S}$ , where  $m_1 < n_1$  and  $m_2 < n_2$ .



# Convolutional Neural Networks (CNNs)

## Definition

A CNN with  $L$  layers consists of  $L$  iterative applications of a convolutional layer, followed by an activation layer, (possibly) followed by a pooling layer.

# Convolutional Neural Networks (CNNs)

## Definition

A CNN with  $L$  layers consists of  $L$  iterative applications of a convolutional layer, followed by an activation layer, (possibly) followed by a pooling layer.



Typical architectures consist of a CNN (as a *feature extractor*), followed by a fully connected NN (as a *classifier*)


# Convolutional Neural Networks (CNNs)

## Definition

A CNN with  $L$  layers consists of  $L$  iterative applications of a convolutional layer, followed by an activation layer, (possibly) followed by a pooling layer.



Typical architectures consist of a CNN (as a *feature extractor*), followed by a fully connected NN (as a *classifier*)



**Figure:** LeNet (1998, LeCun et al): the first successful CNN architecture, used for reading handwritten digits

# Feature Extractor vs. Classifier



# Feature Extractor vs. Classifier






# Feature Extractor vs. Classifier




D. Trump



B. Sanders



A. Merkel



B. Johnson



```
x_image = tf.reshape(x, [-1, 28, 28, 1])  
# First convolutional layer - maps one grayscale image to 32 feature maps.  
W_conv1 = weight_variable([5, 5, 1, 32])  
b_conv1 = bias_variable([32])  
h_conv1 = tf.nn.conv2d(x_image, W_conv1) + b_conv1  
# Pooling layer - downsamples by 2X.  
h_pool1 = max_pool_2x2(h_conv1)  
# Second convolutional layer -- maps 32 feature maps to 64.  
W_conv2 = weight_variable([5, 5, 32, 64])  
b_conv2 = bias_variable([64])  
h_conv2 = tf.nn.conv2d(h_pool1, W_conv2) + b_conv2  
# Second pooling layer.  
h_pool2 = max_pool_2x2(h_conv2)  
# Fully connected layer 1 -- after 2 round of downsampling, our 28x28 image  
W_fc1 = weight_variable([7 * 7 * 64, 1024])  
b_fc1 = bias_variable([1024])  
h_pool2_flat = tf.reshape(h_pool2, [-1, 7*7*64])  
h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)  
# Map the 1024 features to 10 classes, one for each digit  
W_fc2 = weight_variable([1024, 10])  
b_fc2 = bias_variable([10])  
y_conv = tf.matmul(h_fc1, W_fc2) + b_fc2
```

```

x_image = tf.reshape(x, [-1, 28, 28, 1])
# First convolutional layer - maps one grayscale image to 32 feature maps.
W_conv1 = weight_variable([5, 5, 1, 32])
b_conv1 = bias_variable([32])
h_conv1 = tf.nn.conv2d(x_image, W_conv1) + b_conv1
# Pooling layer - downsamples by 2X.
h_pool1 = max_pool_2x2(h_conv1)
# Second convolutional layer -- maps 32 feature maps to 64.
W_conv2 = weight_variable([5, 5, 32, 64])
b_conv2 = bias_variable([64])
h_conv2 = tf.nn.conv2d(h_pool1, W_conv2) + b_conv2
# Second pooling layer.
h_pool2 = max_pool_2x2(h_conv2)
# Fully connected layer 1 -- after: www.tensorflow.org.
W_fc1 = weight_variable([7 * 7 * 64, 1024])
b_fc1 = bias_variable([1024])
h_pool2_flat = tf.reshape(h_pool2, [-1, 7*7*64])
h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)
# Map the 1024 features to 10 classes, one for each digit
W_fc2 = weight_variable([1024, 10])
b_fc2 = bias_variable([10])
y_conv = tf.matmul(h_fc1, W_fc2) + b_fc2

```

Python Library *Tensor Flow*,  
developed by Google Brain, based  
on symbolic computational graphs

[www.tensorflow.org](http://www.tensorflow.org).

```

x_image = tf.reshape(x, [-1, 28, 28, 1])
# Flatten -- maps one grayscale image to 32 feature maps.
W_conv1 = tf.get_variable('W_conv1', [[5, 5, 1, 32]])
b_conv1 = tf.get_variable('b_conv1', [32])
h_conv1 = tf.nn.conv2d(x_image, W_conv1) + b_conv1
# Pooling -- maps 32 feature maps to 16.
h_pool1 = tf.nn.max_pool(h_conv1, [1, 2, 2, 1], [0, 1, 1, 0], [0.5, 0.5])
# Second convolutional layer -- maps 32 feature maps to 64.
W_conv2 = tf.get_variable('W_conv2', [[5, 5, 32, 64]])
b_conv2 = tf.get_variable('b_conv2', [64])
h_conv2 = tf.nn.conv2d(h_pool1, W_conv2) + b_conv2
# Pooling -- maps 64 feature maps to 16.
h_pool2 = tf.nn.max_pool(h_conv2, [1, 2, 2, 1], [0, 1, 1, 0], [0.5, 0.5])
# Fully connected layer -- maps 16 feature maps to 10 classes, one for each digit.
W_fc1 = tf.get_variable('W_fc1', [1024, 1024])
b_fc1 = tf.get_variable('b_fc1', [1024])
h_fc1 = tf.nn.matmul(h_pool2_flat, W_fc1) + b_fc1
# Map to 10 classes, one for each digit.
W_fc2 = tf.get_variable('W_fc2', [1024, 10])
b_fc2 = tf.get_variable('b_fc2', [10])
y_conv = tf.nn.matmul(h_fc1, W_fc2) + b_fc2

```

Python Library *Tensor Flow*,  
developed by Google Brain, based  
on symbolic computational graphs  
-- after: [www.tensorflow.org](http://www.tensorflow.org).