

Digital microsystems design

Marius Marcu

2020

Objectives

- Specific objectives
 - Internal architecture and external interfaces of x86 processors and their functional behavior

Outline

- Pins
- Machine cycles
- Data transfer
 - Normal cycles
 - Pipelined cycles
 - Burst cycles
- Pipeline
- System architecture
- Summary

80x86 microprocessors family

- 8086
 - Release year: 1978
 - First 16 bits microprocessor
- 80386
 - Release year: 1985
 - 32 bits microprocessor
- Pentium
 - 32 bits addresses/ 64 bits data
- Core
- x86 microprocessors family
 - IA-16
 - IA-32 (i386)
 - IA-32-64 (amd64)
- IA64 – Itanium
 - No IA32 compatible

8086

- Characteristics
 - 16 bits internal registers
 - 16 bits external data bus
 - 20 bits external address bus
 - Address space: 1 MB
 - 20 bits memory address space
 - 16 bits I/O address space
 - Multiplexed address and data lines
 - 5-10 MHz
 - 40 pins DIL package - Socket

80386



- Characteristics
 - 32 bits internal registers
 - 32 bits external data bus
 - 32 bits external address bus
 - Address space: 4 GB
 - 32 bits memory address space
 - 16 bits I/O address space
 - Non-multiplexed address and data lines
 - 12-33 MHz
 - 132 pins PGA package

Pentium


- Characteristics
 - 32 bits internal registers
 - 64 bits external data bus
 - 32 bits external address bus
 - Address space: 4 GB
 - 32 bits memory address space
 - 16 bits I/O address space
 - Non-multiplexed address and data lines
 - 60-66 MHz
 - 273 pins PGA package – Socket 4

Processor sockets

- PGA – pin grid array

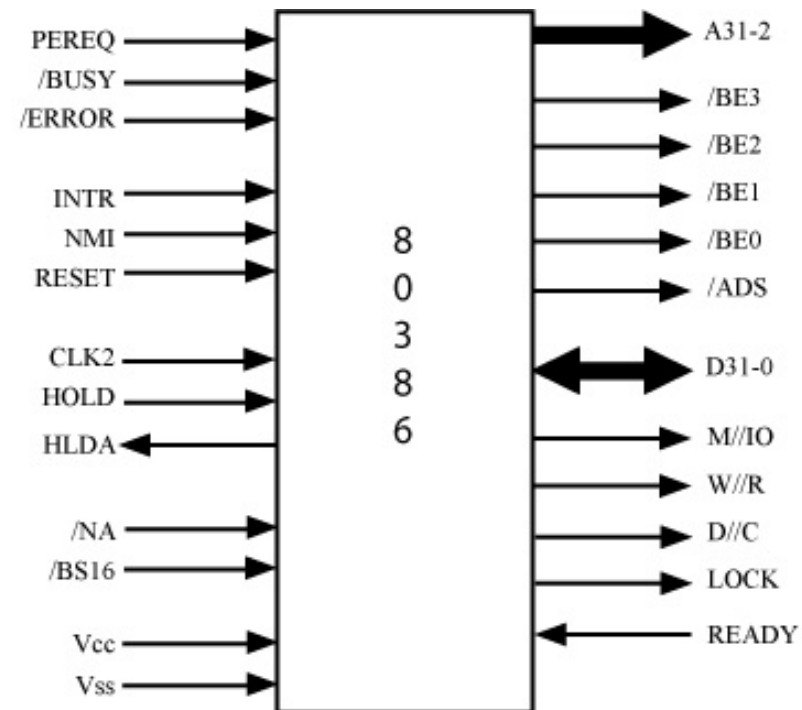


- LGA – land grid array



Processor pins

- 80386 local bus



Processor pins

- 80386

Name	Type	Description
A31-2	O, 3-state (Z)	Memory-I/O address bus
D31-0	I/O, Z	Data bus
/BE3-0	O, Z	Byte enables
/ADS	O, Z	Address status
D/ /C	O, Z	Data/Control indication
W/ /R	O, Z	Write-Read control indication - the processor is performing a memory or I/O write/read cycle
M/ /I/O	O, Z	Memory-I/O indication - distinguishes a memory access from an I/O access
/READY	I	Ready – acknowledgement from the addressed memory or I/O that the transfer will be performed

Processor pins

- 80386

Name	Type	Description
INTR	I	Maskable interrupt request – level triggered input sampled during last clock cycle of each instruction, when active the processor acknowledge the interrupt and execute the interrupt service routine
NMI	I	Non-masking interrupt – an edge triggered input which causes an interrupt all the time
/LOCK	O, Z	Bus lock - indicates that other bus masters are not allowed to gain the control of system bus (used for synchronization)
RESET	I	Reset – stops immediately the current execution and restarts the processor
CLK2	I	Double the internal clock

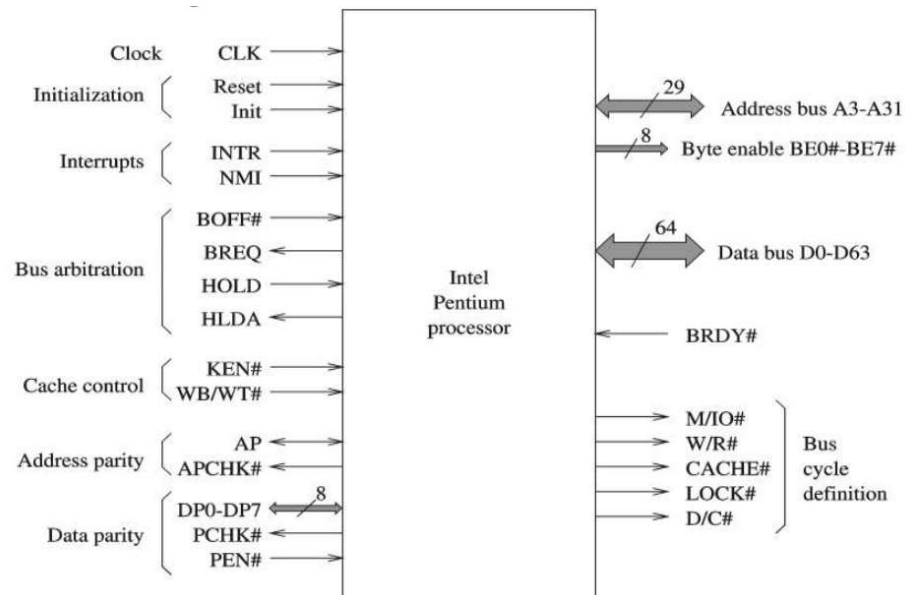
Processor pins

- 80386

Name	Type	Description
PEREQ	I	Coprocessor request (fetch data from coprocessor)
/BUSY	I	Coprocessor busy
/ERROR	I	Coprocessor error
/NA	I	Next address
/BS16	I	Bus size 16 bits
HOLD	I	Bus hold request of another bus master
HLDA	O	Bus hold acknowledgment

Processor pins

- Pentium
 - 32bits processor
 - 64 bits data bus width



Processor pins

- Advanced operations
 - Bus lock (LOCK#)
 - Used in read-modify-write cycle
 - Useful in implementing semaphores
 - Cacheability (CACHE#)
 - Read cycle: indicates internal cacheability
 - Write cycle: burst write-back
 - Cache enable (KEN#)
 - If asserted, the current cycle is transformed into cache line fill
 - WB/ /WT – write-back / write-through

Machine cycles

- Low level external operation of a microprocessor is based on machine cycles
- A machine cycle defines the basic operations that a processor performs to transfer data with external components:
 - memory
 - I/O
- Bus cycles are the external result of executing machine cycles by a microprocessor
- Every interaction between the processor and external components is implemented by a specific machine cycle

Machine cycles

- Every external operation has a dedicated machine cycle, implemented as a bus cycle on microprocessor's bus:
 - Instruction fetch
 - Memory read
 - Memory write
 - I/O read
 - I/O write
 - Interrupt acknowledge


Machine cycles

- Example – 80386 machine cycles encoding

M/ /IO	D/ /C	W/ /R	Bus cycle	Locked
0	0	0	Interrupt acknowledge	Yes
0	0	1	Not possible	-
0	1	0	I/O data read	No
0	1	1	I/O data write	No
1	0	0	Memory code read (Fetch)	No
1	0	1	Processor control - HALT - SHUTDOWN	No
1	1	0	Memory data read	If LOCK
1	1	1	Memory data write	If LOCK


Machine cycles

- Every machine cycle is implemented by a bus cycle on processor external bus
- A bus cycle corresponds to a sequence of events (signal activation and deactivation), called bus protocol
- A bus cycle can be described using time diagrams
 - Typical read/write bus cycles (conceptual)




Machine cycles

- A machine cycle is implemented by a specific number of states
 - Called T-States
 - 80386: minimum 2 clock periods, T1-T2, clock period – twice of CLK2
 - Timing diagram for normal cycles




Machine cycles

- Normal cycles (386)




Machine cycles

- Pipelined bus cycles (conceptual)
 - Address for the next cycle is provided in advance, starting in the current cycle

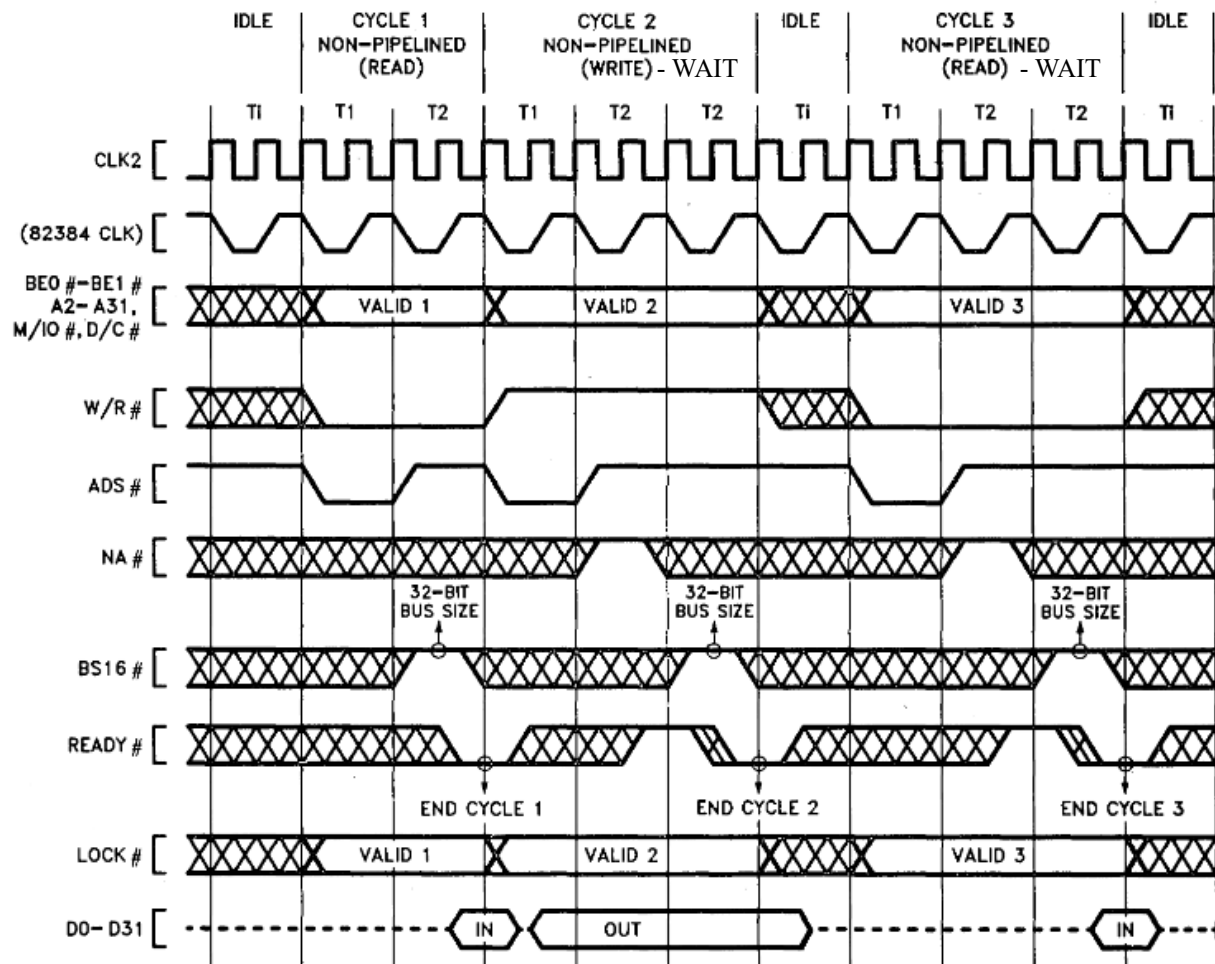


Machine cycles

- Address generated in advance – pipelined cycle (386)



- 386 datasheet
 - Normal cycle
 - Wait states

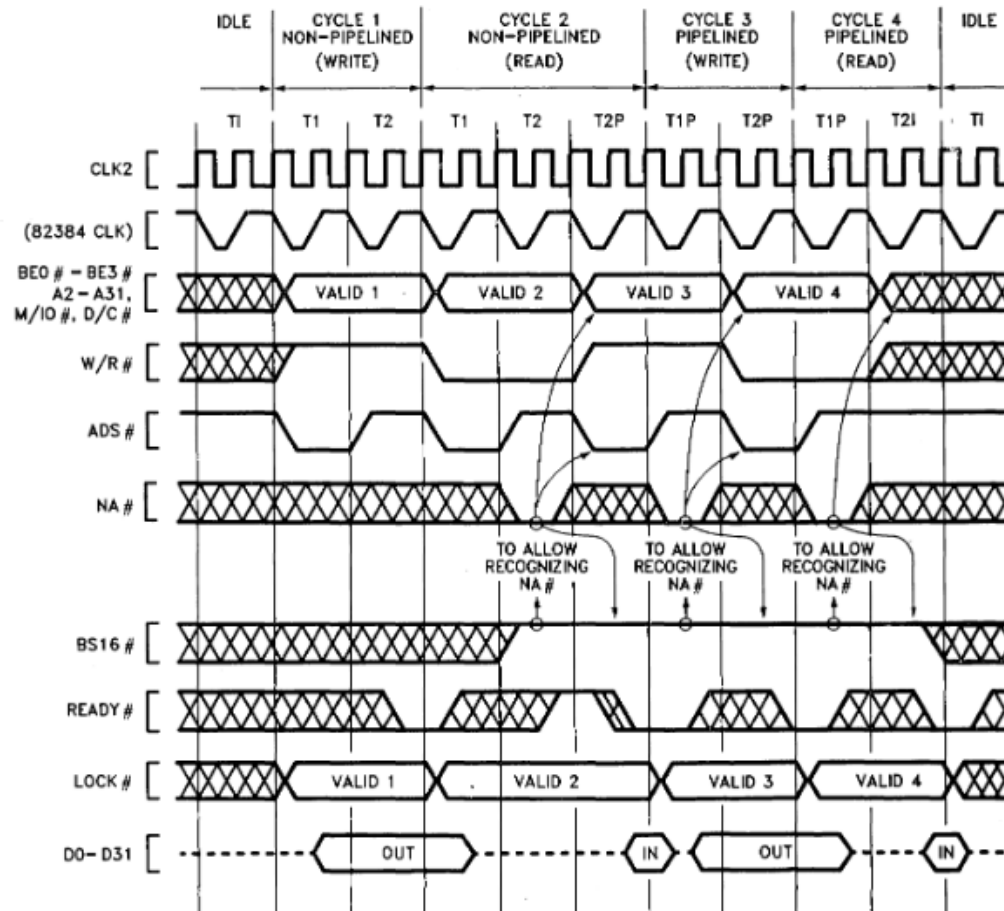


Machine cycles

- Wait states
 - Used by slower components as a mean to ask the processor to wait for the current transaction to be finished
 - Activated according to memory access time

Machine cycles

- 386 datasheet
– Pipelined cycle



Machine cycles

- How a processor can execute an operation?

```
int a,b;
```

```
a = 10;
```

```
b = a;
```

Machine cycles

- How many bus/machine cycles a processor will perform to execute one instruction?
 - Every instruction is encoded on several bytes – the machine code of the instruction
 - MOV dest, src

```
MOV AX, 10                ; constant value
; read from a memory location
MOV AX, [BX]              ; [xx] - memory location
; write to a memory location
MOV [BX], AX              ; xx - register content
```

Machine cycles

– MOV dest, src

MOV AX, [BX]

- Instruction fetch (reads the code of the next instruction to be decoded and executed by the processor)
- Memory read (reads the memory content of the location addressed by BX)

Machine cycles

– MOV dest, src

MOV [BX], AX

- Instruction fetch (reads the code of the next instruction to be de coded and executed by the processor)
- Memory write (writes the content of the AX register into the memory location addressed by BX)

Machine cycles


- The execution of one instruction is achieved by a sequence of several machine cycles:
 - Instruction fetch,
 - Read operands,
 - Internal execution
 - Write results
- A machine cycle can be extended by inserting wait states, a state is a clock cycle or clock period.
 - When the memory is slow

Machine cycles

- Clock and reset generator
 - Generates clock signals
 - Processor
 - Specialized circuits
 - Generates READY signal
 - Synchronized with clock signal
 - When inactive the current machine cycle will wait for a slower component involved in current transaction
 - Will add wait states to current machine cycle
 - Generated RESET signal
 - Synchronized with the clock

Machine cycles

- Clock and reset generator



Processor startup

- Reset
 - Input pin – hardware reset
 - When active the processor stops immediately its execution
 - Reset signal is synchronized with clock signal
 - Reset generator module
 - It should be kept active for minimum n clock cycles (e.g. $n = 4$ - 8086)

Processor startup

- Reset
 - After reset internal registers are initialized with predefined values:

Register	Content
FLAGS	0000H
IP	0000H
CS	FFFFH
DS	0000H
SS	0000H
ES	0000H
Instruction buffer	Empty

Processor startup

- Reset
 - After reset the processor will start executing code from logical address FFFFH:0000H
 - Real mode
 - ROM should be mapped at this address
 - Physical address FFFF0H
 - Initialization and boot loader code is executed from ROM

Machine cycles

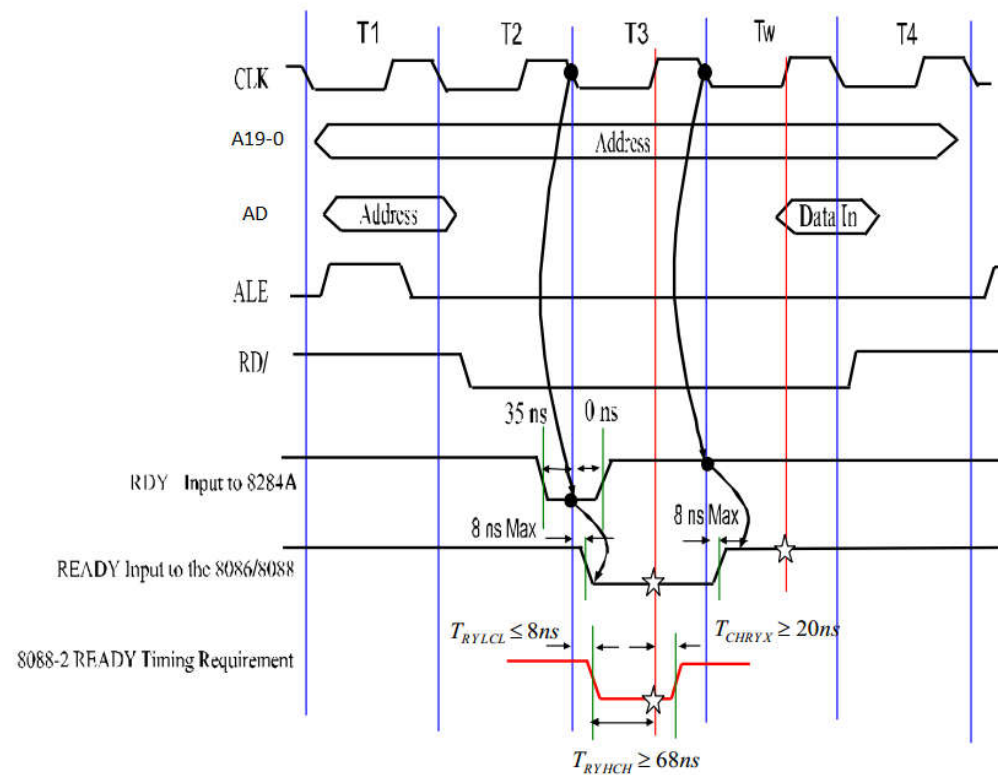
- What happens whether the memory is not fast enough to provide or store data to/from a processor?

- Zero wait states



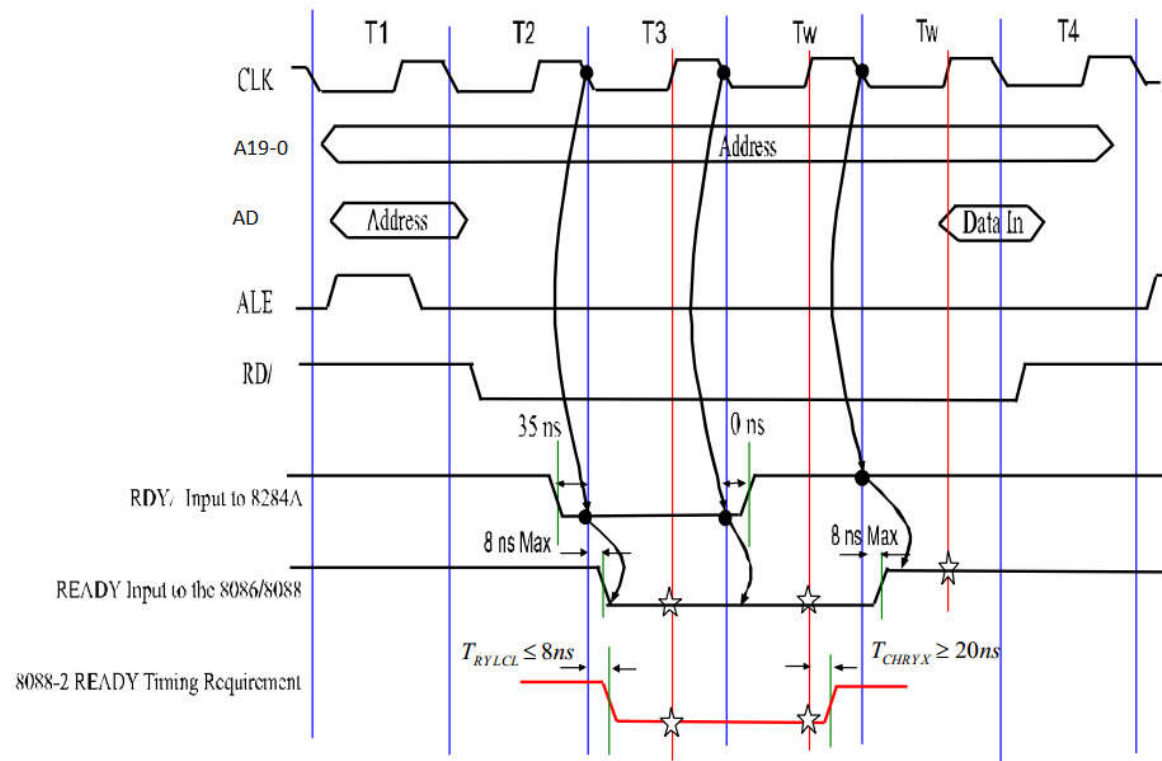
Machine cycles

- One wait state



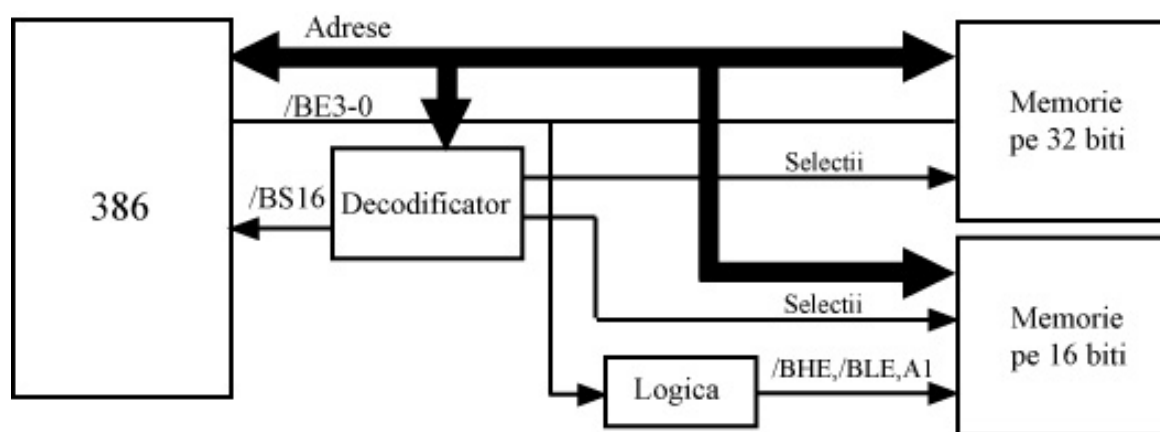
Machine cycles

- Two wait states



Data transfer

- 80386
 - Transfers on 32 bits and 16 bits

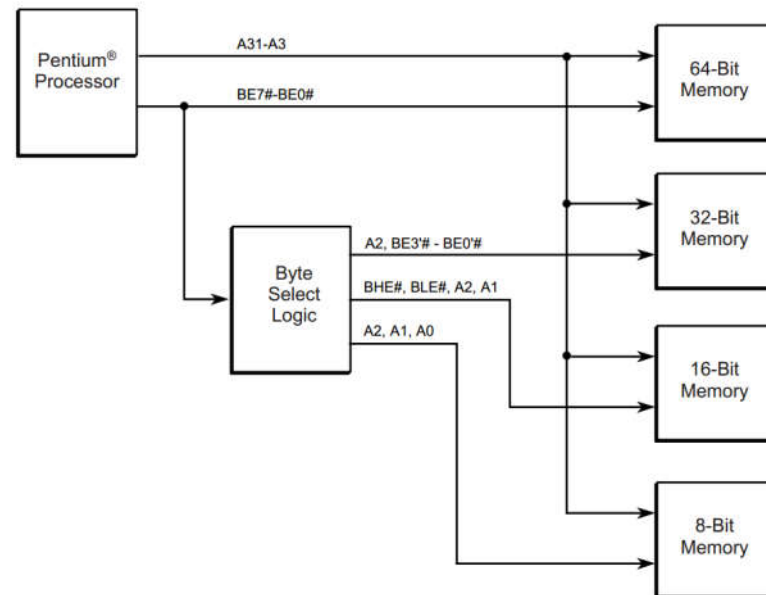


Data transfer

- Exercise
 - Design the memory map and memory decoder for a 32 bits microprocessor system with 32 address lines using the following memory requirements:
 - 16MB EEPROM, using 4M x 16 bits memories
 - 1MB SRAM, using 256K x 32 bits circuits
 - 1GB DRAM, using 128M x 8 bits memories

Data transfer

- Bytes enable
 - 32 bits uP (e.g. 80386) does not contain address connections A1 and A0 because these have been encoded as the byte enable signals (BE3-0).
 - 64 bits uP does not contain address connections A2-A0 because these have been encoded as the byte enable signals (BE7-0).




Data transfer

- 32 bits
 - Data Bus: $D_{31} - D_0$
 - Bi-directional Bus
 - Dynamic Bus size : BS16 input
 - $BE_0 : D_0 - D_7$
 - $BE_1 : D_8 - D_{15}$
 - $BE_2 : D_{16} - D_{23}$
 - $BE_3 : D_{24} - D_{31}$
- 64 bits – Data Bus D63-D0

Byte Enable Signal	Associated Data Bus Signals
BE0#	D0–D7 (byte 0 — least significant)
BE1#	D8–D15 (byte 1)
BE2#	D16–D23 (byte 2)
BE3#	D24–D31 (byte 3)
BE4#	D32–D39 (byte 4)
BE5#	D40–D47 (byte 5)
BE6#	D48–D55 (byte 6)
BE7#	D56–D63 (byte 7 — most significant)


Data transfer

- 32 bits data bus




Data transfer

- Hardware organization of physical address space (32 bits)
 - 32 bits memory block




Data transfer

- 1 byte
 - Address $4xk$




- Address $4xk+1$




Data transfer

- 1 word
 - Address $4xk$
- 1 double word
 - Address $4xk$



Data transfer

- Misaligned double-word data transfer
 - Address $4xk+2$ – the processor will perform 2 bus cycles



Pipeline

- Most instructions have common operations
 - Fetch instruction code from memory
 - Decode instruction code
 - Select operand registers
 - Use ALU for arithmetic and logic operations

Pipeline

- For each instructions the processor performs several tasks:
 - Instruction fetch (IF)
 - Instruction decode (ID)
 - Execution (EX)
 - Memory access (MEM)
 - Write-back (WB)
- However, each instruction uses different subsets of existing functional units
- Hence the instructions need various times to complete
 - Some instructions are faster than others

Pipeline


- If the processor executes instructions internally in one clock cycle, the operating frequency is established based on the slowest instruction
- Sequential execution of all tasks required by each instruction is called single-cycle execution

Pipeline

- Single cycle data path is not efficient in time
 - Clock cycle time is determined by the instruction taking the longest time
 - Variable clock cycle time is too complicated
- Alternative design approaches
 - Multiple clock cycles per instruction

Pipeline

- Pipelining is a microprocessor implementation technique in which multiple instructions are overlapped in execution
- One instruction requires several clock cycles to execute



Pipeline

- The pipeline is divided in stages
- Each stage completes a part of an instruction in parallel with the rest of the stages
- The stages are connected one to the next to form a pipe
 - Outputs of one stage are inputs for the next stage
 - The instructions enter at one end, progress through the stages, and exit at the other end

Pipeline

- Pipelining
 - Multiple tasks operating simultaneously using different resources
 - It does not reduce latency of single instruction execution, it increase the throughput of entire workload (several instruction which feed the pipeline)
 - Pipeline rate is limited by slowest pipeline stage
 - Unbalanced lengths of pipeline stages reduces speedup

Pipeline


- Pipelining speedup
 - The potential speedup is determined by the number of the stages
 - The effective speedup is limited by
 - Time needed to fill the pipeline with instructions
 - Time needed to finish last instructions
 - Stalls for dependencies
 - Cleaning for branches

Pipeline

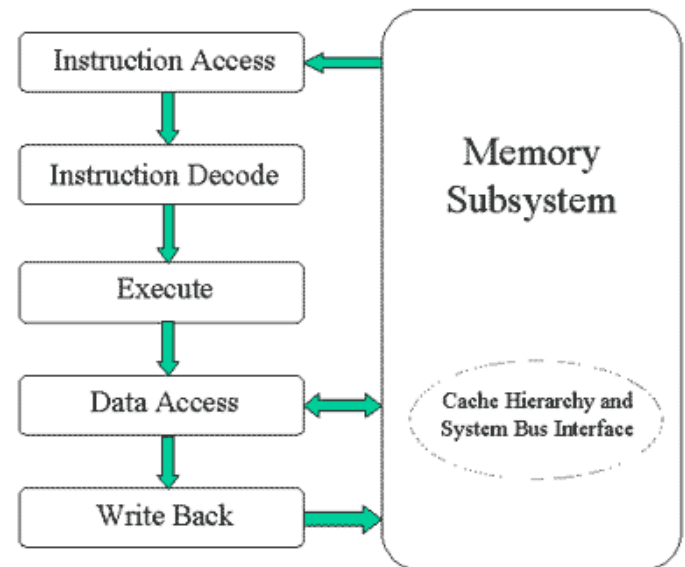
- Pipelining implementation
 - Instruction execution tasks becomes pipeline stages:
 - Instruction fetch (IF)
 - Instruction decode (ID)
 - Execution (EX)
 - Memory access (MEM)
 - Write-back (WB)
 - Each stage performs its operation in one clock cycle
 - The clock speed is determined by the slowest stage

Pipeline

- Five stage pipelining – concept



Four Pipelined Instructions



Pipeline

- Instruction Fetch (IF) stage
 - Is responsible for reading the requested instruction from memory.
 - The instruction and the program counter (which is incremented to the next instruction) are stored in the IF/ID pipeline register as temporary storage so that may be used in the next stage at the start of the next clock cycle.
- Instruction Decode (ID) stage
 - Is responsible for decoding the instruction and sending out the various control lines to the other parts of the processor.
 - The instruction is sent to the control unit where it is decoded and the registers are fetched from the register file.

Pipeline

- Execution (EX) stage
 - Is where any calculations are performed.
 - The main component in this stage is the ALU.
 - The ALU is made up of arithmetic, logic capabilities.
- Memory and IO (MEM) stage
 - Is responsible for storing and loading values to and from memory.
 - It also responsible for input or output from the processor.
 - If the current instruction is not of Memory or IO type than the result from the ALU is passed through to the write back stage

Pipeline

- Write-Back (WB) stage
 - Is responsible for writing the result of a calculation, memory access or input into the register file.

Pipeline

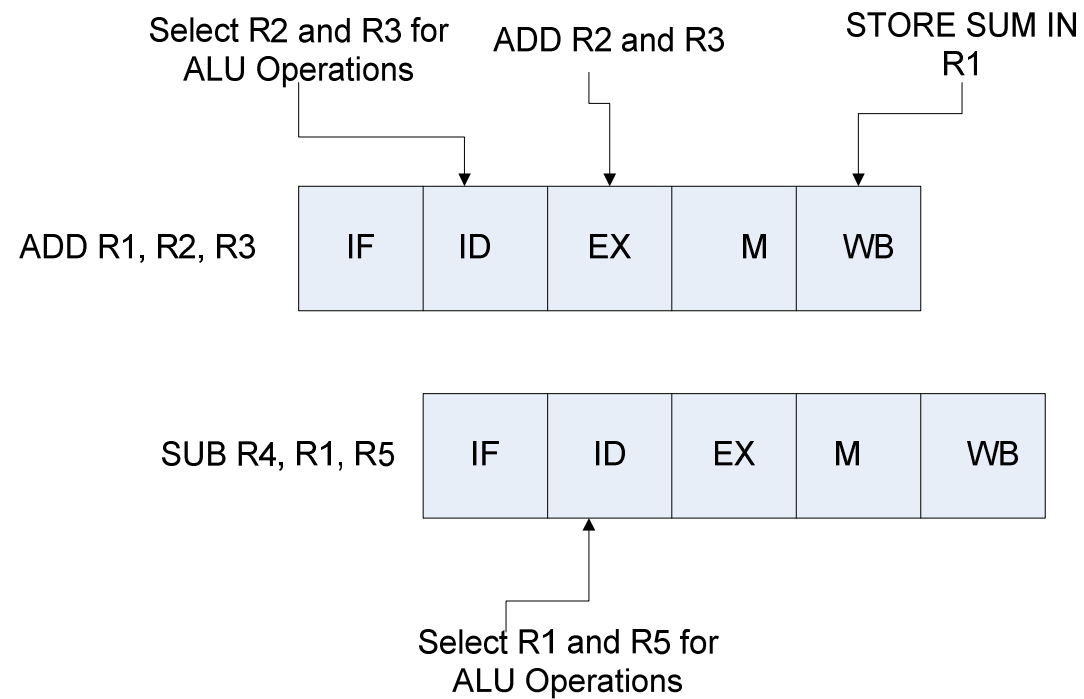
- Advantages:
 - More efficient use of processor's resources
 - Decrease execution time of blocks of instructions
 - Support for ILP
- Disadvantages:
 - Pipelining involves adding hardware to the chip
 - Inability to continuously run the pipeline at full speed because of pipeline hazards which disrupt the smooth execution of the pipeline.
 - Increase execution time of one instruction

Pipeline

- Situations that prevent starting the next instruction in the next cycle are called hazards
 - Structural Hazards – two instructions need to access the same resource
 - Data Hazards – an instruction uses the result of the previous instruction
 - Control Hazards – the location (address) of an instruction depends on previous instruction

Pipeline

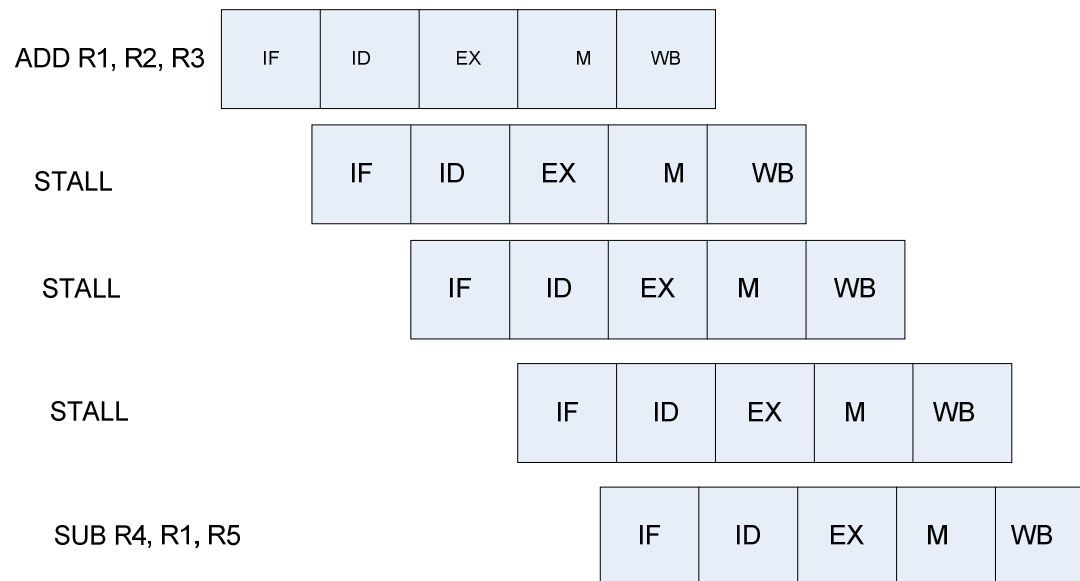
- Data hazards



Pipeline

- Data hazards

- Stalling - involves halting the flow of instructions until the required result is ready to be used. However stalling wastes processor time by doing nothing while waiting for the result

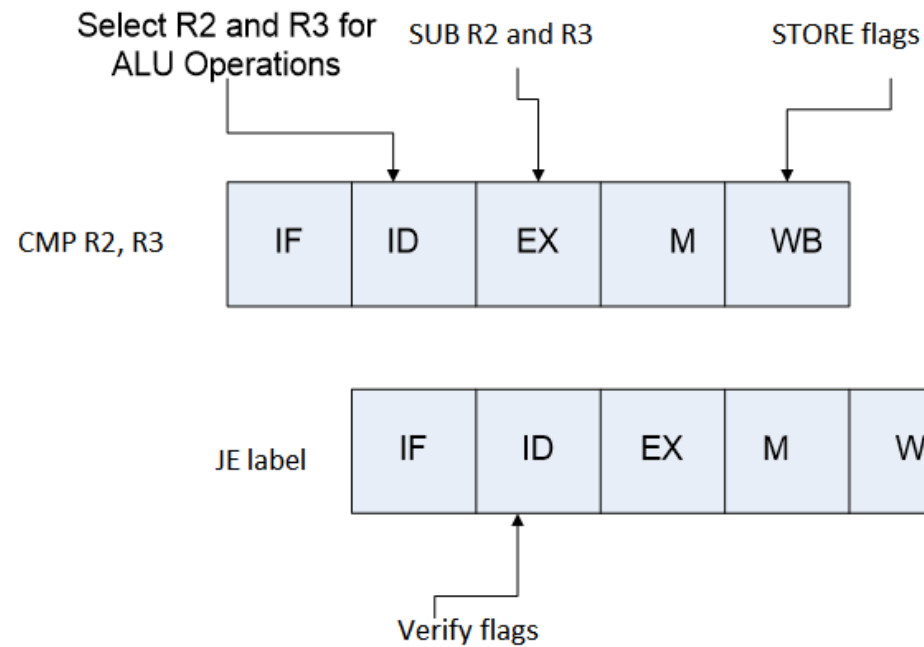


Pipeline

- Control hazards
 - Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch

Pipeline

- Control hazards



Pipeline

- Control hazards solutions
 - Stalls
 - Branch prediction
 - Longer pipelines cannot determine branch outcome early
 - Stall penalty becomes unacceptable
 - Predict: guess one direction then back up if wrong
 - Predict outcome of branch
 - Only stall if prediction is wrong
 - Empty the pipeline
 - Avoid branches
 - Conditional instructions (e.g. CMOV)
 - Predicative instructions

Pipeline

- Static branch prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- Dynamic branch prediction
 - Hardware measures actual branch behavior
 - record recent history of each branch
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history


Pipeline

- Branch Target Buffer – BTB
 - Cache memory which stores target addresses of branch instructions
 - Predicted addresses are decided at code execution
 - Additional flags encoding branch address probability

Instr address	Predicted PC

System architecture

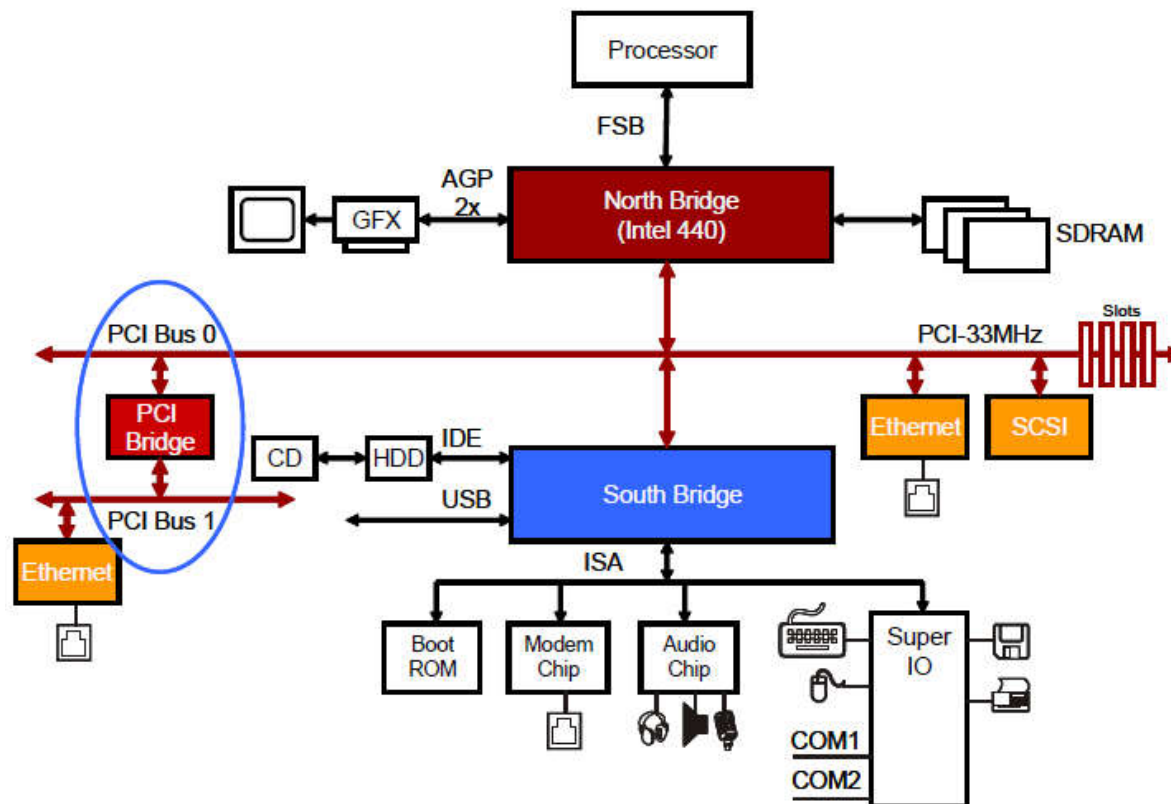
- FSB
- Memory bus
- Chipset
- I/O bus



System architecture


- Chipsets
 - Designed for use with a specific family of processors
 - Northbridge: connects the CPU to high-speed components like RAM, video card
 - Southbridge: connects to slower peripheral devices like PCI slots, USB, IDE, ...

System architecture

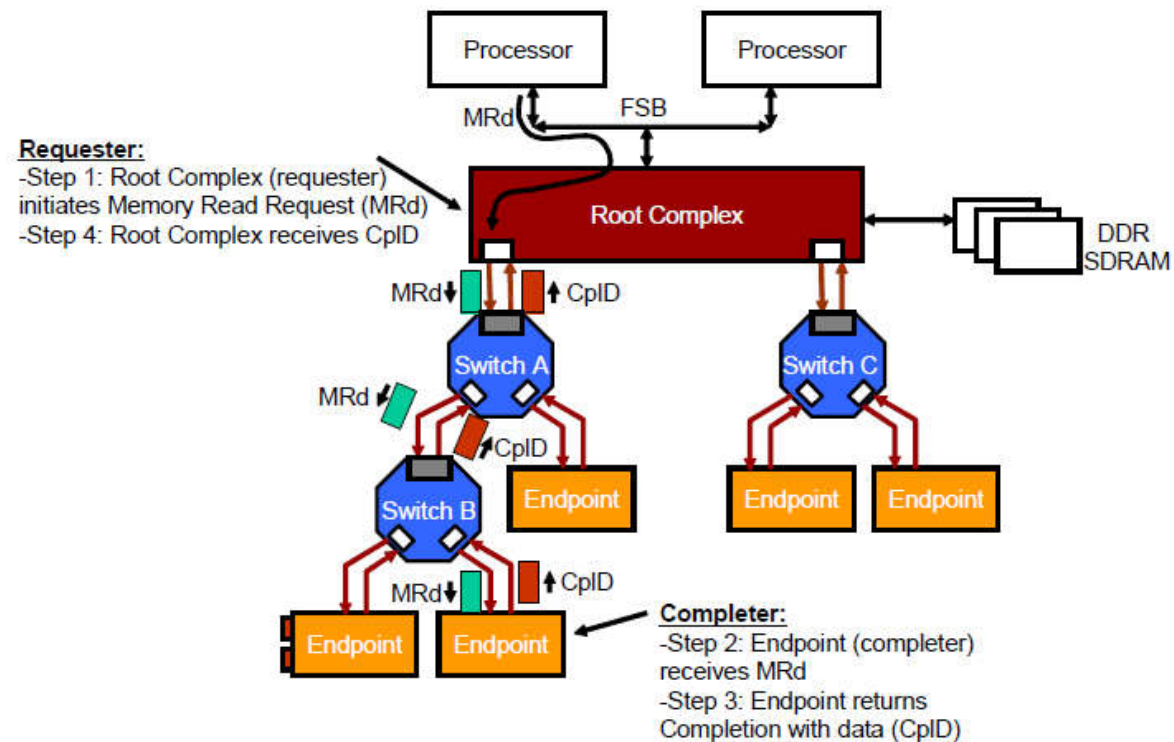


System architecture

- PCI express

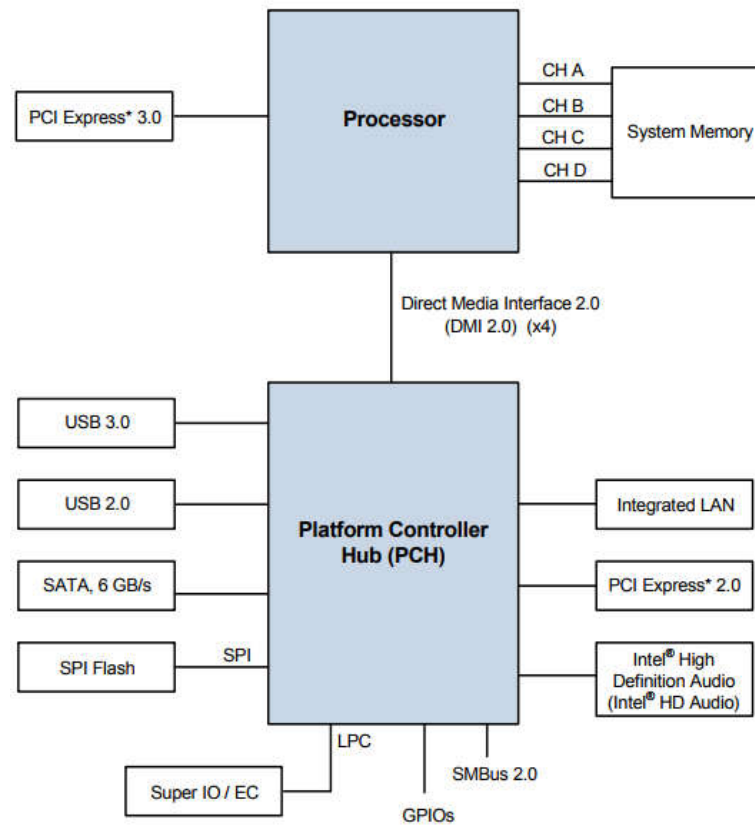


System architecture



System architecture

- iCore7



Interconnect

- FSB

Figure 3. Shared Front-side Bus, up until 2004

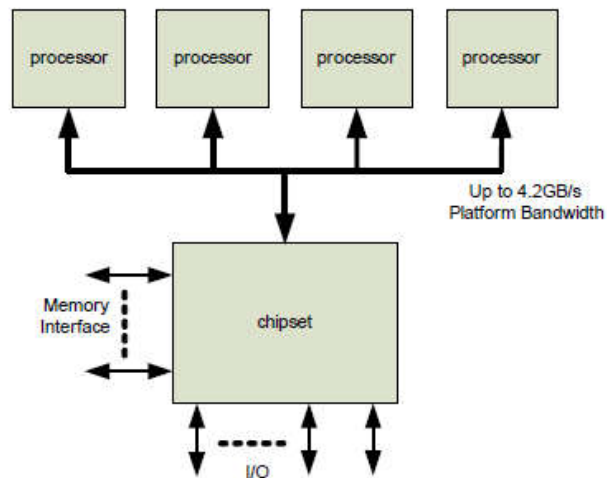



Figure 4. Dual Independent Buses, circa 2005



Interconnect

- DHSI, QPI
 - Point-to-point

Figure 5. Dedicated High-speed Interconnects, 2007





Figure 6. Intel® QuickPath Interconnect



Interconnect

- QPI
 - High-speed, packetized, point-to-point interconnect
 - Distributed shared memory architecture (NUMA)

Figure 7. Block Diagram of Processor with Intel® QuickPath Interconnects

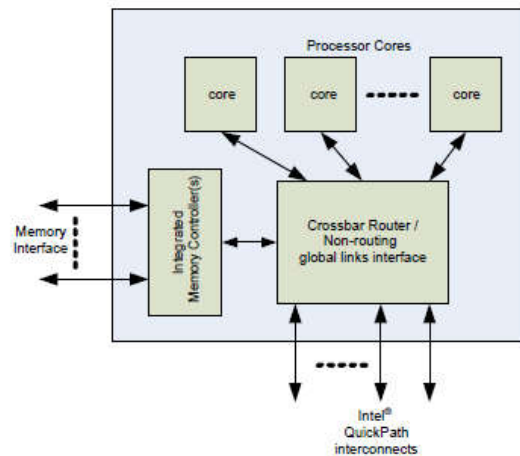

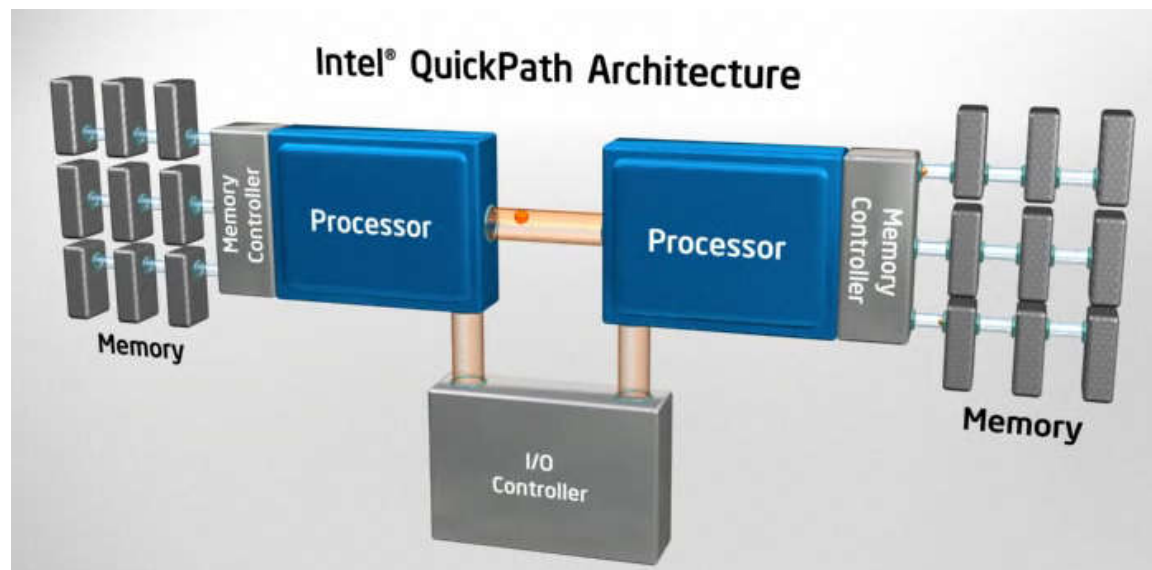


Figure 9. Physical Layer Diagram




Interconnect

- Quick Path Interconnect
 - iCore
 - NUMA



Interconnect

- Quick Path Interconnect
 - FSB - ~13GBps
 - QPI - ~26GBps

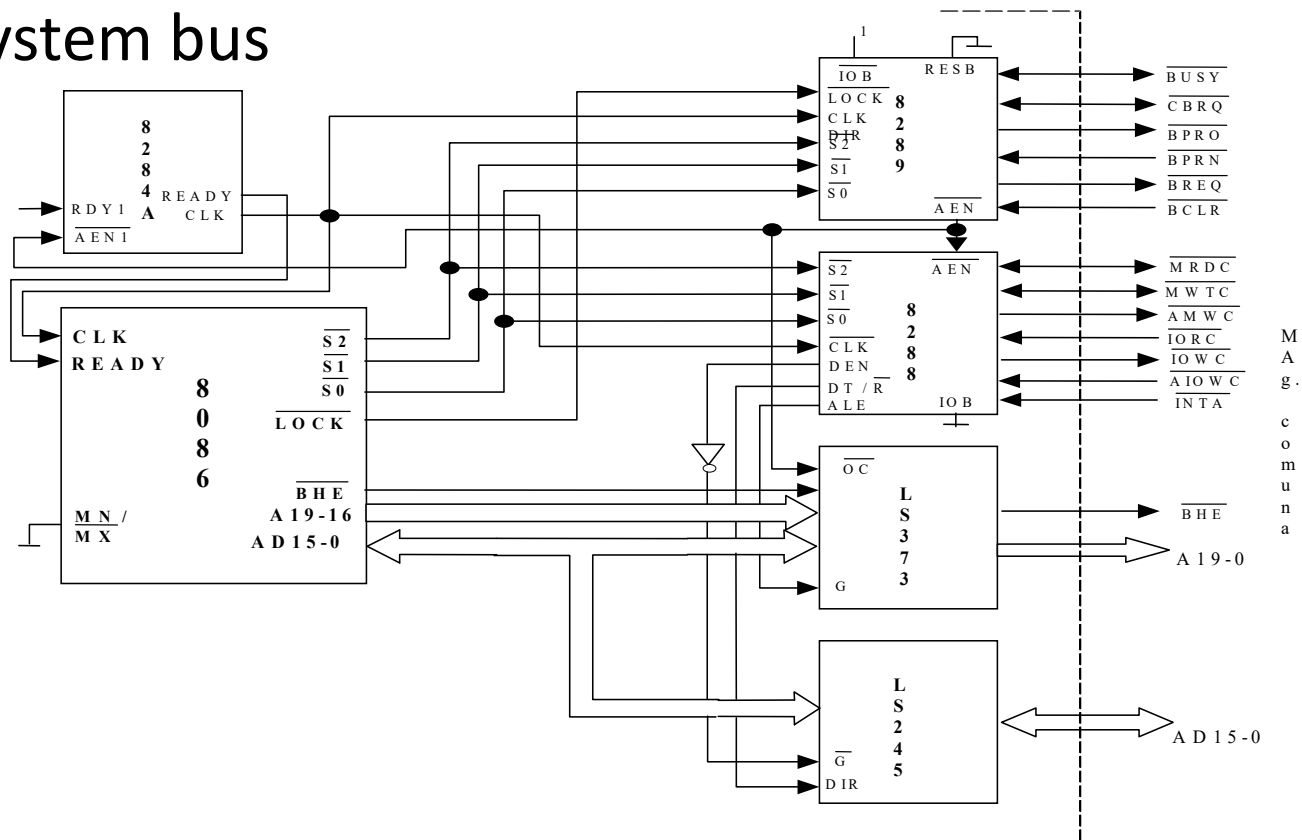


Summary

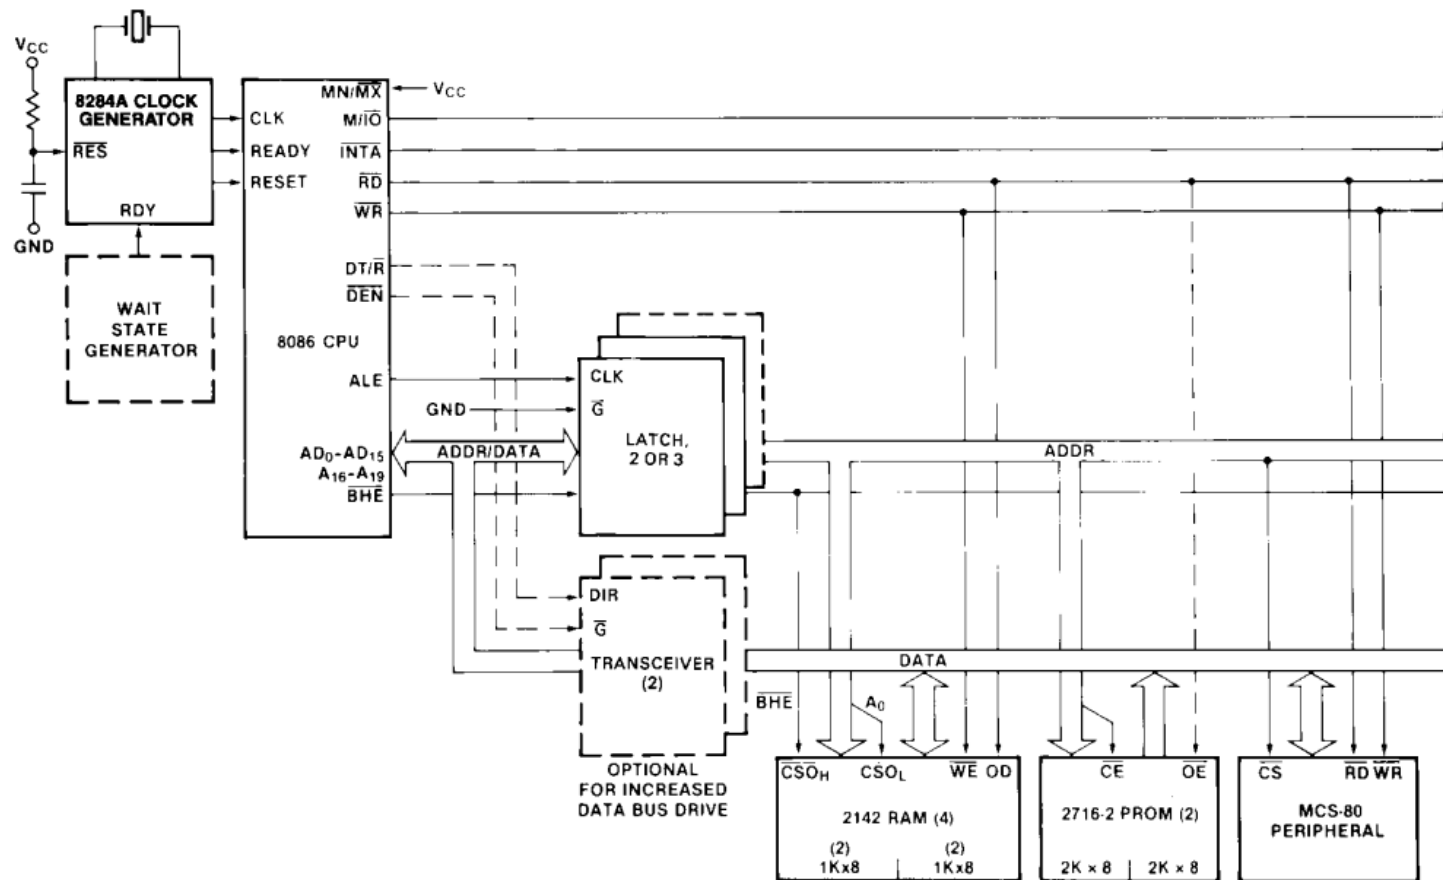
- Machine cycles
- Microsystem design
 - Clock generator
 - Memory connection
 - Data transfer
- Main board architecture

Summary

- Single system bus

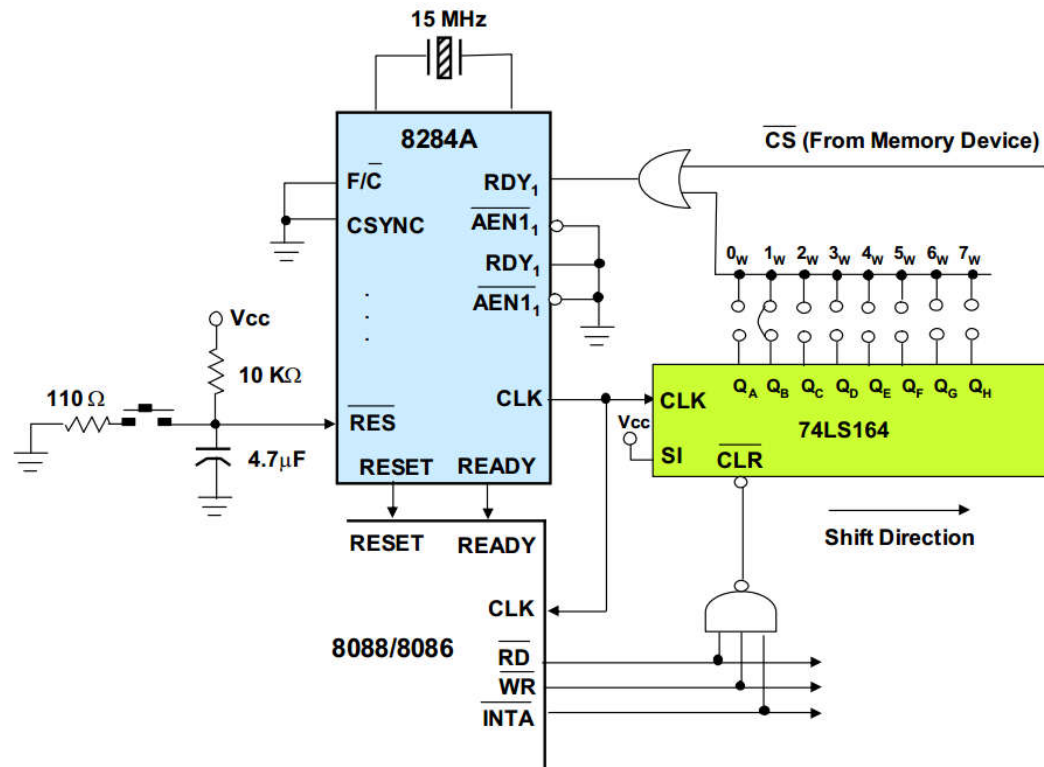


Summary



Machine cycles

- Wait states generator



Pipeline

- Pipelining speedup

