


Reflection and Environment Maps

Sorin Babii
sorin.babii@cs.upt.ro

Objectives

- Texture Mapping Applications
- Reflection (Environment) Maps
 - Cube Maps
 - Spherical Maps
- Bump Maps


Mapping Variations



smooth shading



environment mapping




bump mapping


Environment Mapping

- Environmental (reflection) mapping is the way to create the appearance of highly reflective surfaces without ray tracing, which requires global calculations
- Introduced in movies such as *The Abyss* and *Terminator 2*
- Prevalent in video games
- It is a form of texture mapping

Reflecting the Environment




Mapping to a Sphere



Hemisphere Map as a Texture


- If we map all objects to hemisphere, we cannot tell if they are on the sphere or anywhere else along the reflector
- Use the map on the sphere as a texture that can be mapped onto the object
- Can use other surfaces as the intermediate
 - Cube maps
 - Cylinder maps

Cube Map



Indexing into Cube Map

- Compute $R = 2(N \cdot V)N - V$
- Object at origin
- Use largest magnitude component of R to determine face of cube
- Other two components give texture coordinates



OpenGL Implementation

- WebGL supports only cube maps
 - desktop OpenGL also supports sphere maps
- First must form map
 - Use images from a real camera
 - Form images with WebGL
- Texture map it to object

Cube Maps

- We can form a cube map texture by defining six 2D texture maps that correspond to the sides of a box
- Supported by WebGL through cubemap sampler

```
vec4 texColor = textureCube(mycube, texcoord);
```

- Texture coordinates must be 3D
 - usually are given by the vertex location \Rightarrow we don't need to compute separate tex coords

Environment Maps with Shaders



- Environment maps are usually computed in world coordinates which can differ from object coordinates because of the modeling matrix
 - May have to keep track of modeling matrix and pass it to the shaders as a uniform variable
- Can also use reflection map or refraction map for effects such as simulating water

Issues



- Must assume environment is very far from object
(equivalent to the difference between near and distant lights)
- Object cannot be concave (no self reflections possible)
- No reflections between objects
- Need a reflection map for each object
- Need a new map if viewer moves...

Building Cube Map...

Use six cameras, each with a 90 degree angle of view



... and the Cube Image



WebGL style

```
gl.textureMap2D(  
    gl.TEXTURE_CUBE_MAP_POSITIVE_X,  
    level, rows, columns, border, gl.RGBA,  
    gl.UNSIGNED_BYTE, image1)
```

- Same for other five images
- Make one texture object out of the six images

Example

- Consider rotating cube inside a cube that reflects the color of the walls
- Each wall is a solid color (red, green, blue, cyan, magenta, yellow)
 - Each face of room can be a texture of one texel

```
var red = new Uint8Array([255, 0, 0, 255]);  
var green = new Uint8Array([0, 255, 0, 255]);  
var blue = new Uint8Array([0, 0, 255, 255]);  
var cyan = new Uint8Array([0, 255, 255, 255]);  
var magenta = new Uint8Array([255, 0, 255, 255]);  
var yellow = new Uint8Array([255, 255, 0, 255]);
```

Texture Object

```
cubeMap = gl.createTexture();
gl.bindTexture(gl.TEXTURE_CUBE_MAP, cubeMap);
gl.texImage2D(gl.TEXTURE_CUBE_MAP_POSITIVE_X, 0, gl.RGBA,
    1, 1, 0, gl.RGBA, gl.UNSIGNED_BYTE, red);
gl.texImage2D(gl.TEXTURE_CUBE_MAP_NEGATIVE_X, 0, gl.RGBA,
    1, 1, 0, gl.RGBA, gl.UNSIGNED_BYTE, green);
gl.texImage2D(gl.TEXTURE_CUBE_MAP_POSITIVE_Y, 0, gl.RGBA,
    1, 1, 0, gl.RGBA, gl.UNSIGNED_BYTE, blue);
gl.texImage2D(gl.TEXTURE_CUBE_MAP_NEGATIVE_Y, 0, gl.RGBA,
    1, 1, 0, gl.RGBA, gl.UNSIGNED_BYTE, cyan);
gl.texImage2D(gl.TEXTURE_CUBE_MAP_POSITIVE_Z, 0, gl.RGBA,
    1, 1, 0, gl.RGBA, gl.UNSIGNED_BYTE, yellow);
gl.texImage2D(gl.TEXTURE_CUBE_MAP_NEGATIVE_Z, 0, gl.RGBA,
    1, 1, 0, gl.RGBA, gl.UNSIGNED_BYTE, magenta);
gl.activeTexture( gl.TEXTURE0 );
gl.uniform1i(gl.getUniformLocation(program, "texMap"),0);
```

Vertex Shader

```
varying vec3 R;  
attribute vec4 vPosition;  
attribute vec4 vNormal;  
uniform mat4 modelViewMatrix;  
uniform mat4 projectionMatrix;  
uniform vec3 theta;  
void main(){  
    vec3 angles = radians( theta );  
    // compute rotation matrices rx, ry, rz here  
    mat4 ModelViewMatrix = modelViewMatrix*rz*ry*rx;  
    gl_Position = projectionMatrix*ModelViewMatrix*vPosition;  
    vec4 eyePos = ModelViewMatrix*vPosition;  
    vec4 N = ModelViewMatrix*vNormal;  
    R = reflect(eyePos.xyz, N.xyz);  }
```

Fragment Shader

```
precision mediump float;
```

```
varying vec3 R;
```

```
uniform samplerCube texMap;
```

```
void main()
```

```
{
```

```
    vec4 texColor = textureCube(texMap, R);
```


```
    gl_FragColor = texColor;
```

```
}
```

Sphere Mapping

- Original environmental mapping technique proposed by Blinn and Newell based in using lines of longitude and latitude to map parametric variables to texture coordinates
- OpenGL supports sphere mapping which requires a circular texture map equivalent to an image taken with a fisheye lens

Sphere Map



Video

<http://staff.cs.upt.ro/~sorin/webgl/Code/w07/reflectionMap.html>


Bump Maps

Sorin Babii

sorin.babii@cs.upt.ro

Objectives

- Introduce bump mapping




Modeling an Orange


- Consider modeling an orange
- Texture map a photo of an orange onto a surface
 - Captures dimples
 - Will not be correct if we move viewer or light
 - We have shades of dimples rather than their correct orientation
- Ideally we need to perturb normal across surface of object and compute a new color at each interior point

Bump Mapping (Blinn)


Consider a smooth surface



Rougher Version



Tangent Plane



Equations

$$\mathbf{p}(u,v) = [x(u,v), y(u,v), z(u,v)]^T$$

$$\mathbf{p}_u = [\partial x / \partial u, \partial y / \partial u, \partial z / \partial u]^T$$

$$\mathbf{p}_v = [\partial x / \partial v, \partial y / \partial v, \partial z / \partial v]^T$$

$$\mathbf{n} = (\mathbf{p}_u \times \mathbf{p}_v) / |\mathbf{p}_u \times \mathbf{p}_v|$$

Displacement Function

$$\mathbf{p}' = \mathbf{p} + d(u,v) \mathbf{n}$$

$d(u,v)$ is the bump or displacement function

$$|d(u,v)| \ll 1$$

Perturbed Normal

$$\mathbf{n}' = \mathbf{p}'_u \times \mathbf{p}'_v$$

$$\mathbf{p}'_u = \mathbf{p}_u + (\partial d / \partial u) \mathbf{n} + d(u, v) \mathbf{n}_u$$

$$\mathbf{p}'_v = \mathbf{p}_v + (\partial d / \partial v) \mathbf{n} + d(u, v) \mathbf{n}_v$$

If d is small, we can neglect last term

Approximating the Normal

$$\mathbf{n}' = \mathbf{p}'_u \times \mathbf{p}'_v$$

$$\approx \mathbf{n} + (\partial \mathbf{d} / \partial u) \mathbf{n} \times \mathbf{p}_v + (\partial \mathbf{d} / \partial v) \mathbf{n} \times \mathbf{p}_u$$



The vectors $\mathbf{n} \times \mathbf{p}_v$ and $\mathbf{n} \times \mathbf{p}_u$ lie in the tangent plane
Hence the normal is displaced in the tangent plane
Must precompute the arrays $\partial \mathbf{d} / \partial u$ and $\partial \mathbf{d} / \partial v$
Finally, we perturb the normal during shading

Image Processing

- Suppose that we start with a function $d(u,v)$
- We can sample it to form an array $D=[d_{ij}]$
- Then $\partial d / \partial u \approx d_{ij} - d_{i-1,j}$
and $\partial d / \partial v \approx d_{ij} - d_{i,j-1}$
- Embossing: multipass approach using floating point buffer

Example

Single Polygon and a Rotating Light Source



How to do this?

- The problem is that we want to apply the perturbation at all points on the surface
- Cannot solve by vertex lighting (unless polygons are very small)
- Really want to apply to every fragment
- Can't do that in fixed function pipeline
- But can do with a fragment program!!
- See `bumpmap.html` and `bumpmap.js`

Compositing and Blending


Sorin Babii
sorin.babii@cs.upt.ro

Objectives

- Learn to use the A component in RGBA color for
 - Blending for translucent surfaces
 - Compositing images
 - Antialiasing


Opacity and Transparency

- Opaque surfaces permit no light to pass through
- Transparent surfaces permit all light to pass
- Translucent surfaces pass some light
translucency = $1 - \text{opacity } (\alpha)$



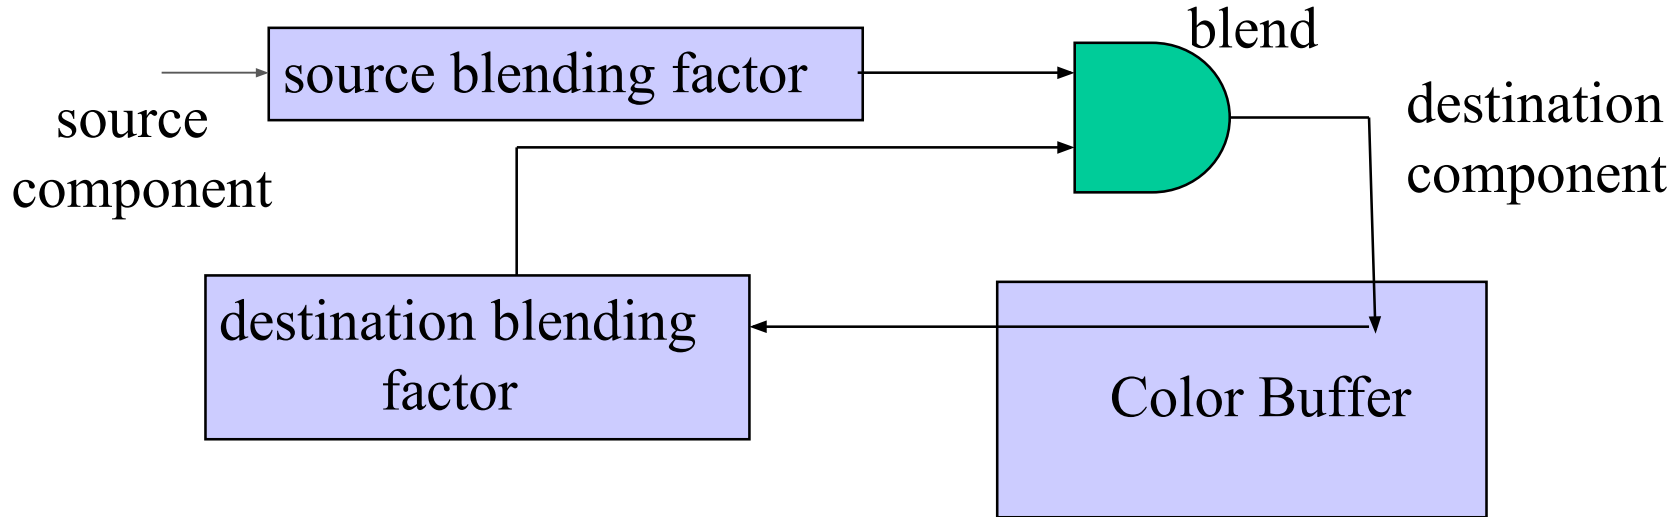
Physical Models

- Dealing with translucency in a physically correct manner is difficult due to:
 - the complexity of the internal interactions of light and matter
 - using a pipeline renderer



Writing Model

- Use A component of RGBA (or RGB α) color to store opacity
- During rendering we can expand our writing model to use RGBA values



Blending Equation

We can define source and destination blending factors for each RGBA component

$$S = [s_r, s_g, s_b, s_a]$$

$$d = [d_r, d_g, d_b, d_a]$$

Suppose that the source and destination colors are

$$b = [b_r, b_g, b_b, b_a]$$

$$c = [c_r, c_g, c_b, c_a]$$

Blend as

$$c' = [b_r s_r + c_r d_r, b_g s_g + c_g d_g, b_b s_b + c_b d_b, b_a s_a + c_a d_a]$$

OpenGL Blending and Compositing

- Must enable blending and pick source and destination factors
 `gl.enable(gl.BLEND)`
 `gl.blendFunc(source_factor,`
 `destination_factor)`
- Only certain factors supported:
 `gl.ZERO`, `gl.ONE`
 `gl.SRC_ALPHA`, `gl.ONE_MINUS_SRC_ALPHA`
 `gl.DST_ALPHA`, `gl.ONE_MINUS_DST_ALPHA`

See Redbook for complete list

Example

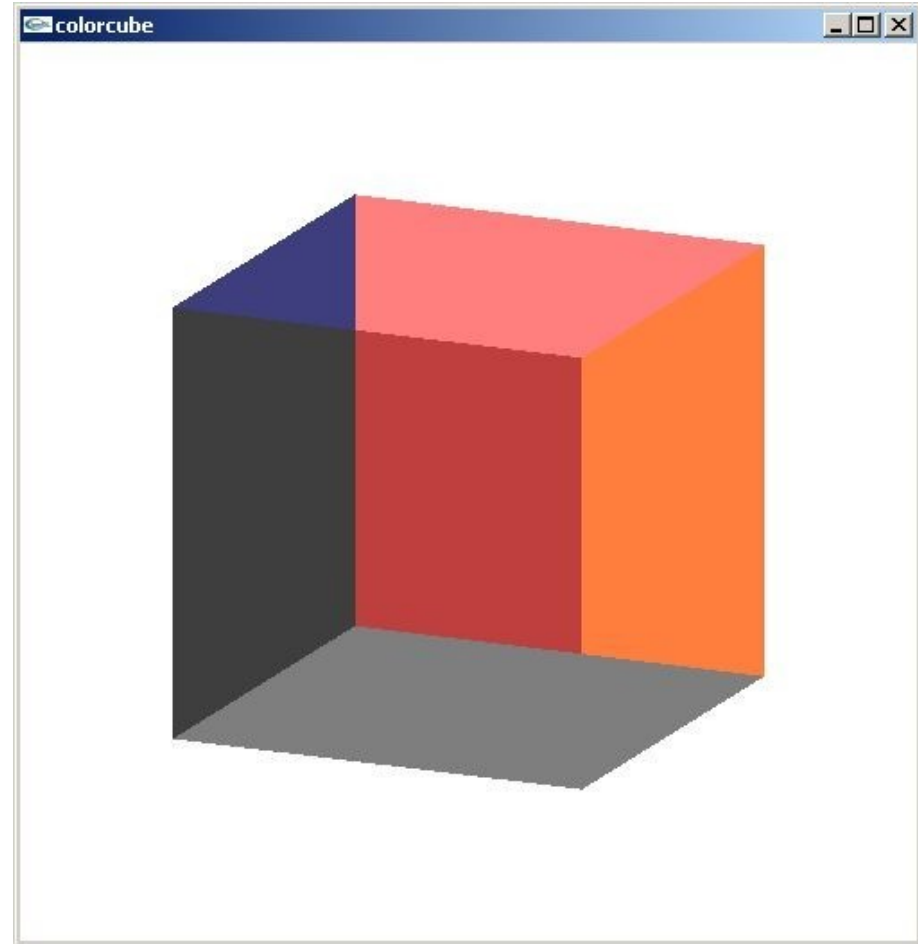
- Suppose that we start with the opaque background color $(R_0, G_0, B_0, 1)$
- This color becomes the initial destination color
- We now want to blend in a translucent polygon with color $(R_1, G_1, B_1, \alpha_1)$
- Select `gl.SRC_ALPHA` and `gl.ONE_MINUS_SRC_ALPHA` as the source and destination blending factors
$$R'_1 = \alpha_1 R_1 + (1 - \alpha_1) R_0, \dots\dots$$
- Note this formula is correct if polygon is either opaque or transparent

Clamping and Accuracy

- All the components (RGBA) are clamped and stay in the range (0,1)
- However, in a typical system, RGBA values are only stored to 8 bits
 - Can easily lose accuracy if we add many components together
 - Example: add together n images
 - Divide all color components by n to avoid clamping
 - Blend with source factor = 1, destination factor = 1
 - But division by n loses bits

Order Dependency

- Is this image correct?
 - Probably not
 - Polygons are rendered in the order they pass down the pipeline
 - Blending functions are order dependent



Opaque and Translucent Polygons

- Suppose that we have a group of polygons, some of which are opaque and some translucent
- How do we use hidden-surface removal?
- Opaque polygons block all polygons behind them and affect the depth buffer
- Translucent polygons should not affect depth buffer
 - Render with `gl.depthMask(false)` which makes depth buffer read-only
- Sort polygons first to remove order dependency


Video

<http://staff.cs.upt.ro/~sorin/webgl/Code/w07/cubet.html>

Fog

- We can composite with a fixed color and have the blending factors depend on depth
 - Simulates a fog effect
- Blend source color C_s and fog color C_f
$$C'_s = f C_s + (1-f) C_f$$
- f is the fog factor
 - Exponential
 - Gaussian
 - Linear (depth cueing)
- Deprecated but can recreate

Fog Functions




Compositing and HTML

- Desktop OpenGL: the A component has no effect unless blending is enabled
- WebGL: an A other than 1.0 has an effect because WebGL works with the HTML5 Canvas element
- $A = 0.5$ will cut the RGB values by $\frac{1}{2}$ when the pixel is displayed
- Allows other applications to be blended into the canvas along with the graphics



Line Aliasing

- Ideal raster line is one pixel wide
- All line segments, other than vertical and horizontal segments, partially cover pixels
- Simple algorithms color only whole pixels
- Lead to the “jaggies” or aliasing
- Similar issue for polygons




Antialiasing

- Can try to color a pixel by adding a fraction of its color to the frame buffer
 - Fraction depends on percentage of pixel covered by fragment
 - Fraction depends on whether there is overlap



Area Averaging

Use average area $\alpha_1 + \alpha_2 - \alpha_1 \alpha_2$ as blending factor



OpenGL Antialiasing

- Not (yet) supported in WebGL
- Can enable separately for points, lines, or polygons

```
glEnable(GL_POINT_SMOOTH);  
glEnable(GL_LINE_SMOOTH);  
glEnable(GL_POLYGON_SMOOTH);
```

```
glEnable(GL_BLEND);  
glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA);
```

- Note: most hardware will automatically antialias

Imaging Applications

Sorin Babii
sorin.babii@cs.upt.ro

Objectives

- Use the fragment shader to do image processing
 - Image filtering
 - Pseudo Color
- Use multiple textures
 - matrix operations
- Introduce GPGPU

Accumulation Techniques (OpenGL)

- Compositing and blending are limited by resolution of the frame buffer
 - Typically 8 bits per color component
- The accumulation buffer was a high resolution buffer (16 or more bits per component) that avoided this problem
- Could write into it or read from it with a scale factor
- Slower than direct compositing into the frame buffer
- Now deprecated but can do techniques with floating point frame buffers

Multirendering

- Composite multiple images
- Image Filtering (convolution)
 - add shifted and scaled versions of an image
- Whole scene antialiasing
 - move primitives a little for each render
- Depth of Field
 - move viewer a little for each render keeping one plane unchanged
- Motion effects

Fragment Shaders and Images

- Suppose that we send a rectangle (two triangles) to the vertex shader and render it with an $n \times m$ texture map
- Suppose that in addition we use an $n \times m$ canvas
- There is now a one-to-one correspondence between each texel and each fragment
- Hence we can regard fragment operations as imaging operations on the texture map

GPGPU

- Looking back at these examples, we can note that the only purpose of the geometry is to trigger the execution of the imaging operations in the fragment shader
- Consequently, we can look at what we have done as large matrix operations rather than graphics operations
- Leads to the field of General Purpose Computing with a GPU (GPGPU)

Examples

- Add two matrices
- Multiply two matrices
- Fast Fourier Transform
- Uses speed and parallelism of GPU
- But how do we get out results?
 - Floating point frame buffers
 - OpenCL (WebCL)
 - Compute shaders

Using Multiple Texels

Suppose we have a 1024×1024 texture in the texture object “image”

`sampler2D(image, vec2(x,y))` returns the value of the texture at (x,y)

`sampler2D(image, vec2(x+1.0/1024.0), y);`
returns the value of the texel to the right of (x,y)

We can use any combination of texels surrounding (x, y) in the fragment shader

Image Enhancer


```
precision mediump float;
varying vec2 fTexCoord;
uniform sampler2D texture;
void main()
{
    float d = 1.0/256.0; //spacing between texels
    float x = fTexCoord.x;
    float y = fTexCoord.y;

    gl_FragColor = 10.0*abs(
        texture2D( texture, vec2(x+d, y))
        - texture2D( texture, vec2(x-d, y)))
        + 10.0*abs( texture2D( texture, vec2(x, y+d))
        - texture2D( texture, vec2(x, y-d)));
    gl_FragColor.w = 1.0;
}
```


Honolulu Image



original



enhanced


Sobel Edge Detector

- Nonlinear
- Find approximate gradient at each point
- Compute smoothed finite difference approximations to x and y components separately
- Display magnitude of approximate gradient
- Simple with fragment shader

Sobel Edge Detector

```
vec4 gx = 3.0*texture2D( texture, vec2(x+d, y))
          + texture2D( texture, vec2(x+d, y+d))
          + texture2D( texture, vec2(x+d, y-d))
          - 3.0*texture2D( texture, vec2(x-d, y))
          - texture2D( texture, vec2(x-d, y+d))
          - texture2D( texture, vec2(x-d, y-d));
vec4 gy = 3.0*texture2D( texture, vec2(x, y+d))
          + texture2D( texture, vec2(x+d, y+d))
          + texture2D( texture, vec2(x-d, y+d))
          - 3.0*texture2D( texture, vec2(x, y-d))
          - texture2D( texture, vec2(x+d, y-d))
          - texture2D( texture, vec2(x-d, y-d));
gl_FragColor = vec4(sqrt(gx*gx + gy*gy), 1.0);
gl_FragColor.w = 1.0;
```

Sobel Edge Detector



Using Multiple Textures

- Example: matrix addition
- Create two samplers, texture1 and texture2 , that contain the data
- In fragment shader

```
gl_FragColor =  
    sampler2D(texture1, vec2(x, y))  
    +sampler2D(texture2, vec2(x, y));
```

Using 4 Way Parallelism

- Recent GPUs and graphics cards support textures up to $8K \times 8K$
- For scalar imaging, we can do twice as well using all four color components



$$\begin{bmatrix} R & G \\ B & A \end{bmatrix}$$

Indexed and Pseudo Color

- Display luminance (2D) image as texture map
- Treat pixel value as independent variable for separate functions for each color component

```
void main(){  
    vec4 color = texture2D(texture, fTexCoord);  
    if(color.g<0.5) color.g = 2.0*color.g;  
    else color.g = 2.0 - 2.0*color.g;  
    color.b = 1.0-color.b;  
    gl_FragColor = color;  
}
```

Top View of 2D Sinc



The Next Step

- Need more storage for most GPGPU calculations
- Example: filtering
- Example: iteration
- Need shared memory
- Solution: Use texture memory and off-screen rendering

Fractals - Mandelbrot


Sorin Babii
sorin.babii@cs.upt.ro

Objectives

- Introduce the most famous fractal object
 - more about fractal curves and surfaces later
- Imaging calculation
 - Must compute value for each pixel on display
 - Shows power of fragment processing

Sierpinski Gasket

Rule:



Repeat n times.

As $n \rightarrow \infty$

Area $\rightarrow 0$

Perimeter $\rightarrow \infty$

Not a normal geometric object

More about fractal curves and surfaces later

Reminder: Complex Arithmetic

- Complex number defined by two scalars

$$z = x + jy$$

$$j^2 = -1$$

- Addition and Subtraction


$$z_1 + z_2 = x_1 + x_2 + j(y_1 + y_2)$$

$$z_1 * z_2 = x_1 * x_2 - y_1 * y_2 + j(x_1 * y_2 + x_2 * y_1)$$

- Magnitude

$$|z|^2 = x^2 + y^2$$

Iteration in the Complex Plane



Mandelbrot Set

iterate on $z_{k+1} = z_k^2 + c$
with $z_0 = 0 + j0$

Two cases as $k \rightarrow \infty$

$$|z_k| \rightarrow \infty$$


$|z_k|$ remains finite

If for a given c , $|z_k|$ remains finite, then c belongs to the Mandelbrot set

Computing the Mandelbrot Set

- Pick a rectangular region
- Map each pixel to a value in this region
- Do an iterative calculation for each pixel
 - If magnitude is greater than 2, we know sequence will diverge and point does not belong to the set
 - Stop after a fixed number of iterations
 - Points with small magnitudes should be in set
 - Color each point based on its magnitude


Mandelbrot Set



Exploring the Mandelbrot Set

- Most interesting parts are centered near $(-0.5, 0.0)$
- Really interesting parts are where we are uncertain if points are in or out of the set
- Repeated magnification these regions reveals complex and beautiful patterns
- We use color maps to enhance the detail

Mandelbrot Set



Computing in the JS File (1)

Form a texture map of the set and map to a rectangle

```
var height = 0.5; var width = 0.5;  
    // size of window in complex plane  
var cx = -0.5; var cy = 0.5;  
    // center of window in complex plane  
var max = 100;  
    // number of iterations per point  
var n = 512;  
var m = 512;  
var texImage = new Uint8Array(4*n*m);
```

Computing in JS File (2)

```
for ( var i = 0; i < n; i++ )
  for ( var j = 0; j < m; j++ ) {
    var x = i * ( width / ( n - 1 ) ) + cx - width / 2;
    var y = j * ( height / ( m - 1 ) ) + cy - height / 2;
    var c = [ 0.0, 0.0 ];
    var p = [ x, y ];

    for ( var k = 0; k < max; k++ ) {
// compute  $c = c^2 + p$ 
      c = [c[0]*c[0]-c[1]*c[1], 2*c[0]*c[1]];
      c = [c[0]+p[0], c[1]+p[1]];
      v = c[0]*c[0]+c[1]*c[1];
      if ( v > 4.0 ) break;
    }
  }
```


/* assume not in set if mag > 2 */

Computing in JS File (3)

```
// assign gray level to point based on its magnitude
if ( v > 1.0 ) v = 1.0;          /* clamp if > 1 */
texImage[4*i*m+4*j] = 255*v;
texImage[4*i*m+4*j+1] =
    255*( 0.5* (Math.sin( v*Math.PI/180 ) + 1.0));
texImage[4*i*m+4*j+2] = 255*(1.0 - v);
texImage[4*i*m+4*j+3] = 255;
}
```

- Set up two triangles to define a rectangle
- Set up texture object with the set as data
- Render the triangles

Example



Fragment Shader

- Our first implementation is incredibly inefficient and makes no use of the power of the fragment shader
- Note the calculation is “embarrassingly parallel”
 - computation for the color of each fragment is completely independent
 - Why not have each fragment compute membership for itself?
 - Each fragment would then determine its own color

Interactive Program

- JS file sends window parameters obtained from sliders to the fragment shader as uniforms
- Only geometry is a rectangle
- No need for a texture map since shader will work on individual pixels

Fragment Shader (1)

```
precision mediump float;  
uniform float cx;  
uniform float cy;  
uniform float scale;  
float height;  
float width;
```

```
void main() {
```

```
    const int max = 100;      /* number of iterations per point */  
    const float PI = 3.14159;  
    float n = 1000.0;  
    float m = 1000.0;
```

Fragment Shader (2)

```
float v;  
float x = gl_FragCoord.x/(n*scale) + cx - 1.0/(2.0*scale);  
float y = gl_FragCoord.y/(m*scale) + cy - 1.0/(2.0*scale);  
float ax=0.0, ay=0.0;  
float bx, by;  
for ( int k = 0; k < max; k++ ) {  
// compute  $c = c^2 + p$   
    bx = ax*ax-ay*ay;  
    by = 2.0*ax*ay;  
    ax = bx+x;  
    ay = by+y;  
    v = ax*ax+ay*ay;  
    if ( v > 4.0 ) break; // assume not in set if mag > 2  
}
```

Fragment Shader

```
// assign gray level to point based on its magnitude
```

```
// clamp if > 1
```

```
    v = min(v, 1.0);
```

```
    gl_FragColor.r = v;
```

```
    gl_FragColor.g = 0.5*sin( 3.0*PI*v) + 1.0;
```

```
    gl_FragColor.b = 1.0-v;
```

```
    gl_FragColor.b = 0.5*cos( 19.0*PI*v) + 1.0;
```

```
    gl_FragColor.a = 1.0;
```

```
}
```

Analysis

- This implementation will use as many fragment processors as are available concurrently
- Note: if an iteration ends early, the GPU will use that processor to work on another fragment
- Note also the absence of loops over x and y
- Still not using the full parallelism, since we are really computing a luminance image

Video

<http://staff.cs.upt.ro/~sorin/webgl/Code/w10/mandelbrot2.html>