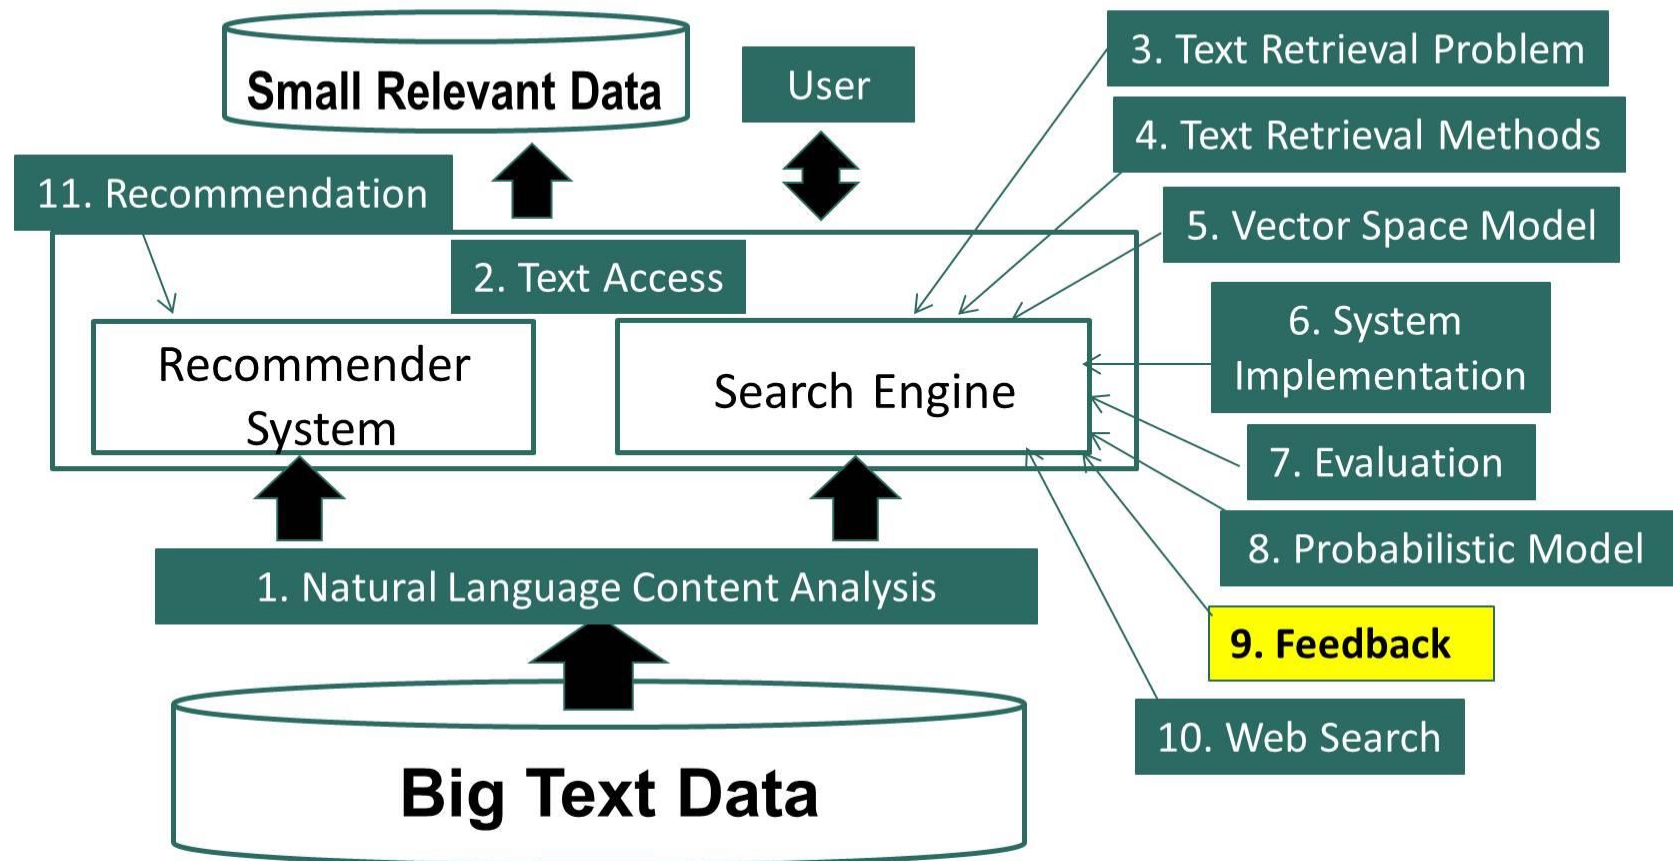


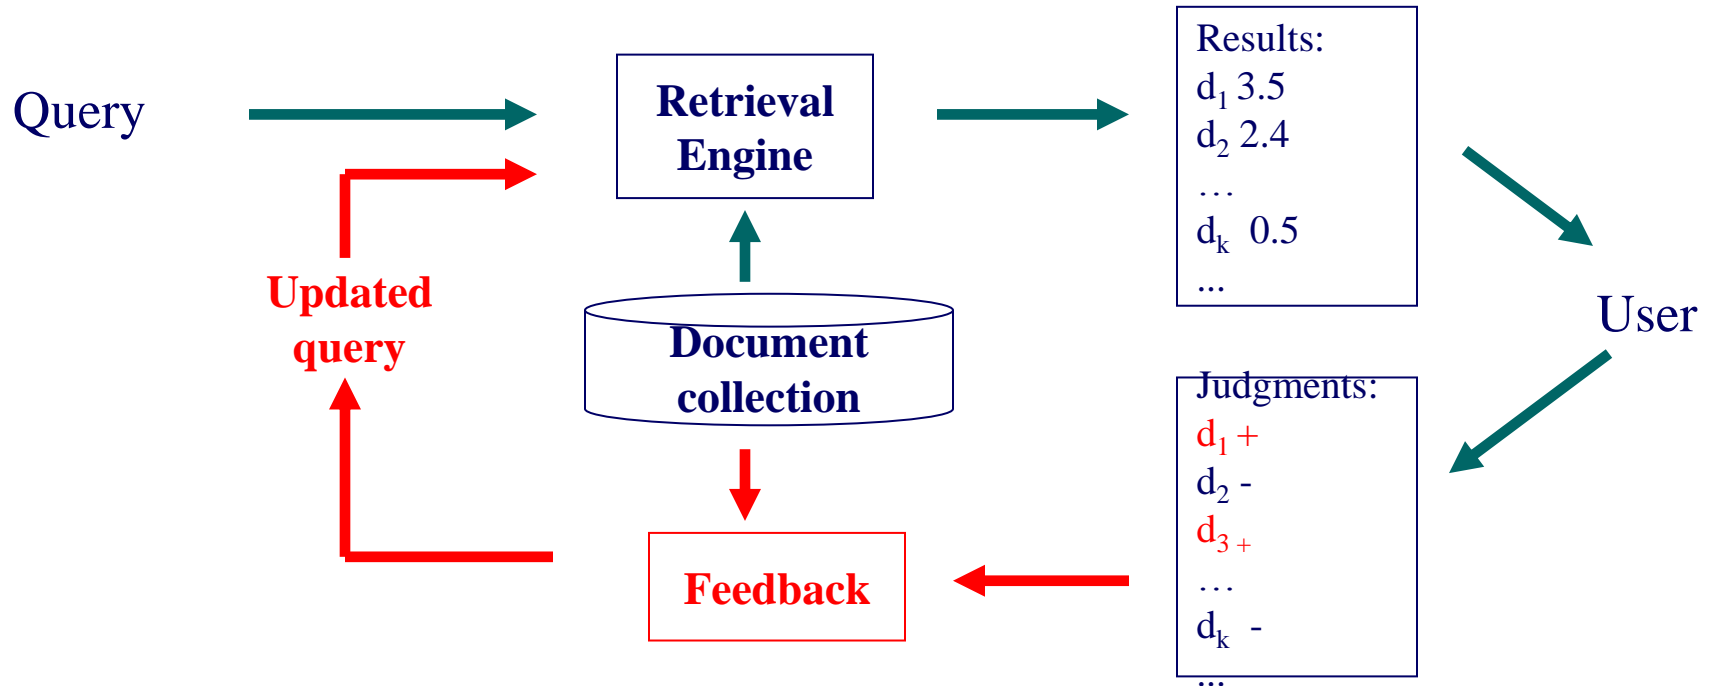
# Retrieval Methods: Feedback in Text Retrieval

# Text Retrieval Methods: Feedback in TR



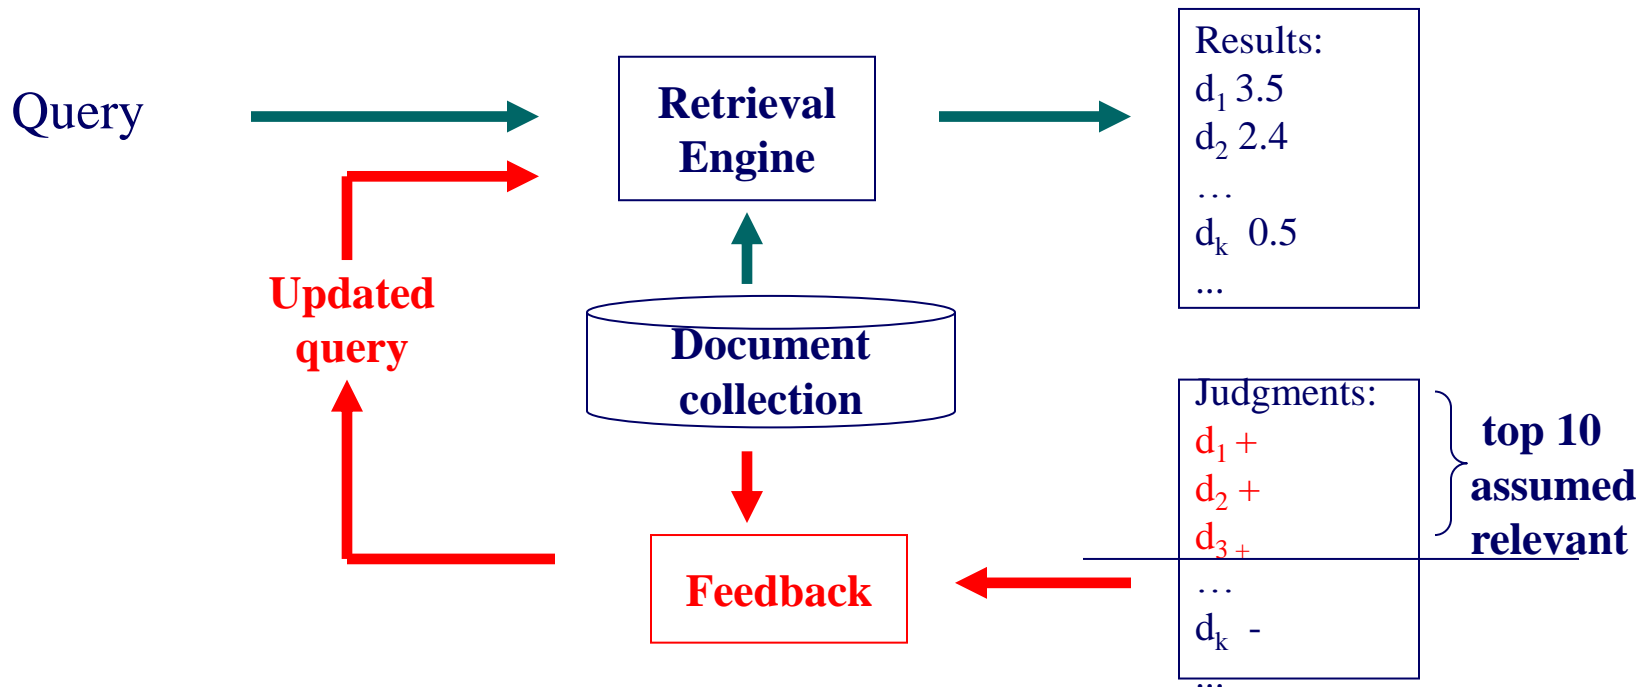
# Relevance Feedback

Users make explicit relevance judgments on the initial results  
(judgments are reliable, but users don't want to make extra effort)



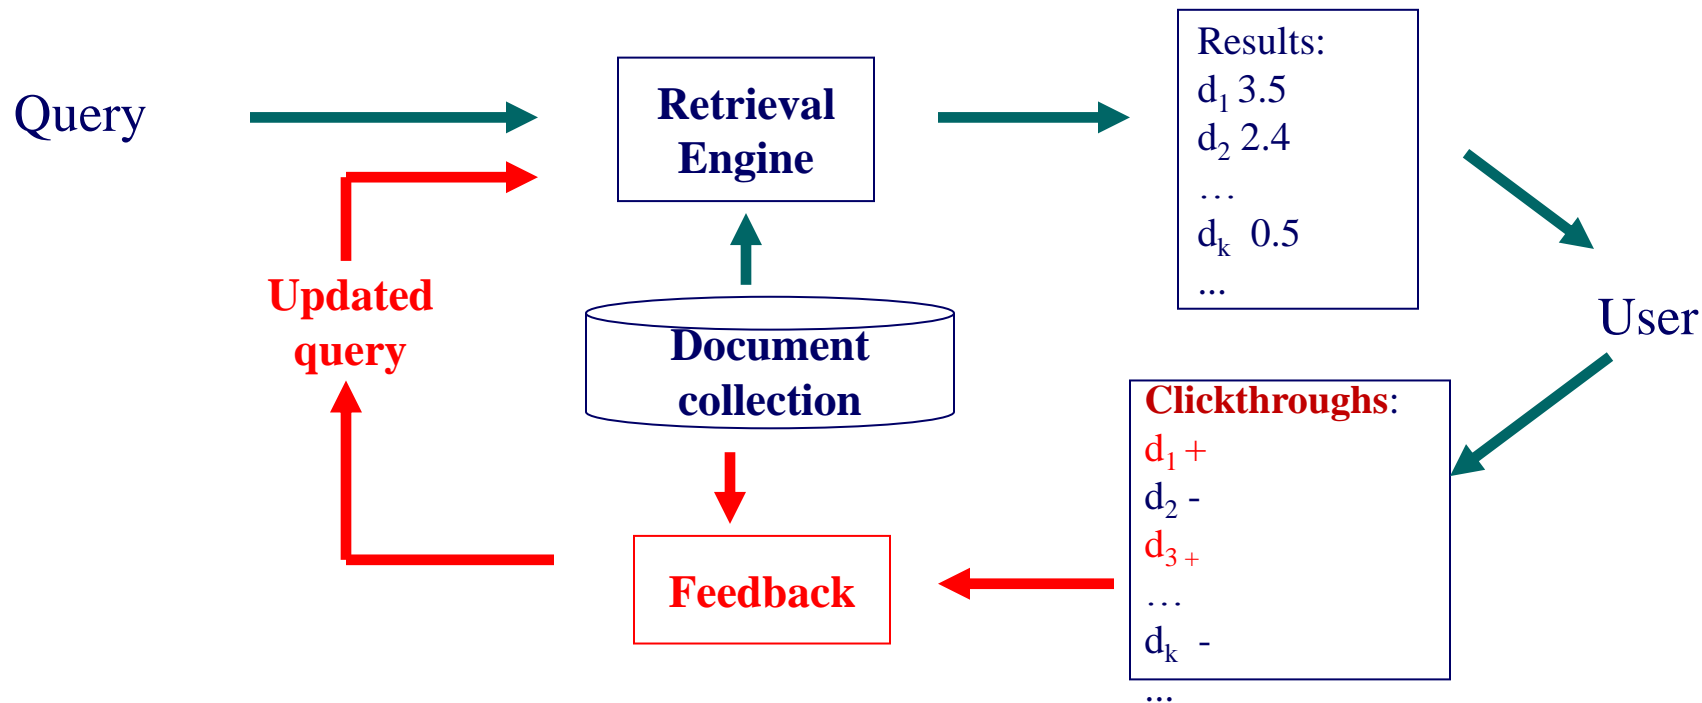
# Pseudo/Blind/Automatic Feedback

Top-k initial results are simply assumed to be relevant  
(judgments aren't reliable, but no user activity is required)



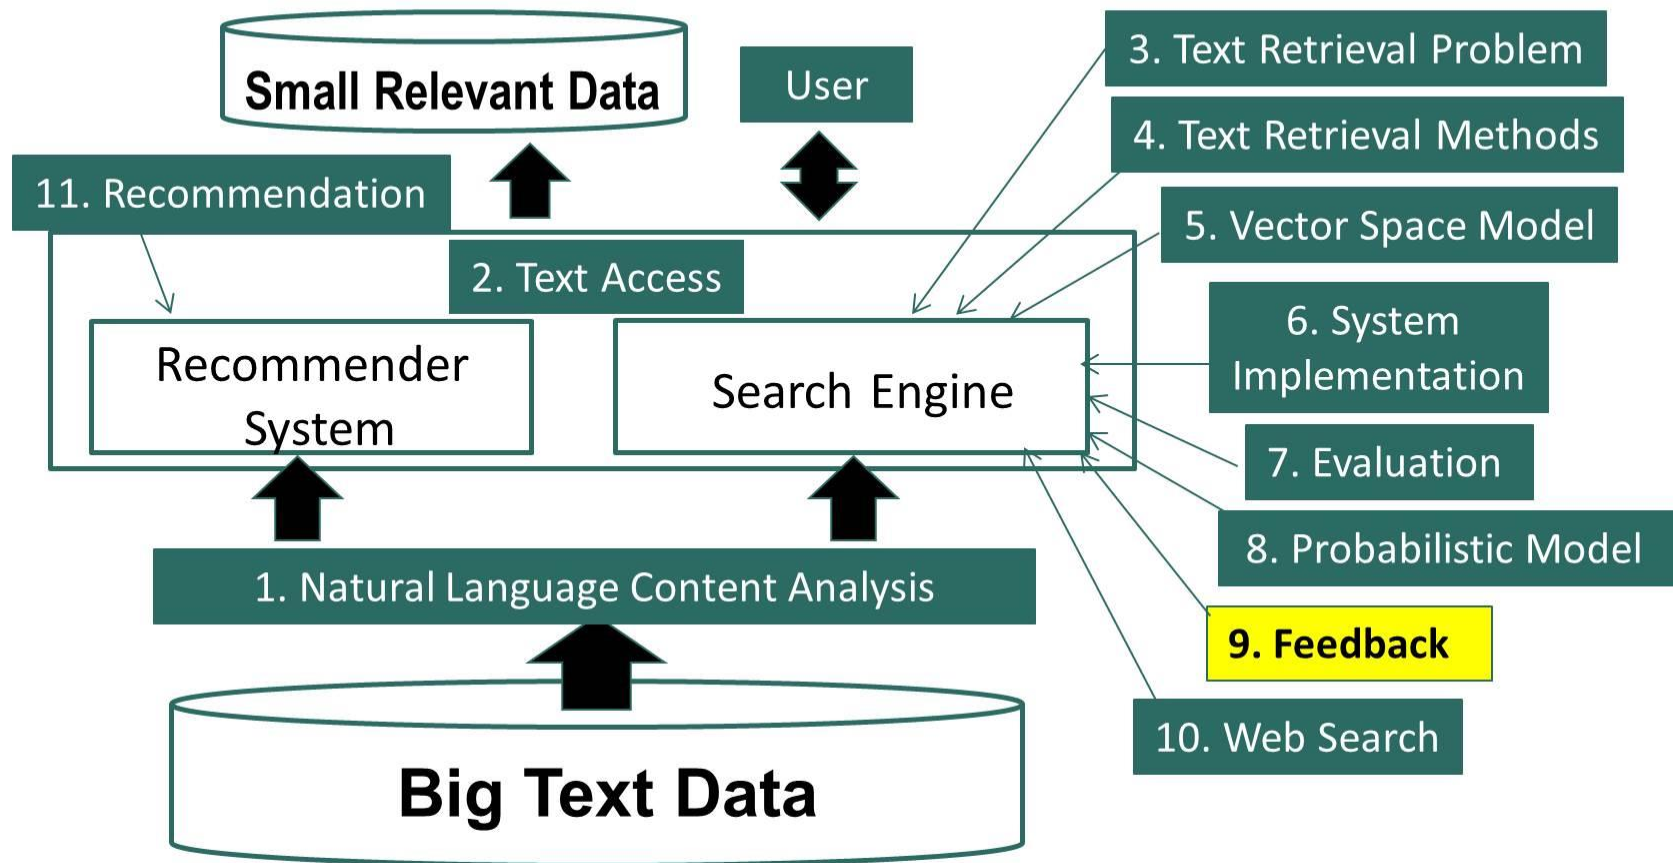
# Implicit Feedback

User-clicked docs are assumed to be relevant; skipped ones non-relevant  
(judgments aren't completely reliable, but no extra effort from users)



# Feedback in Text Retrieval: Feedback in VSM

# Feedback in Text Retrieval: Feedback in VSM




# Feedback in Vector Space Model

- How can a TR system learn from examples to improve retrieval accuracy?
  - Positive examples: docs known to be relevant
  - Negative examples: docs known to be non-relevant
- General method: query modification
  - Adding new (weighted) terms (query expansion)
  - Adjusting weights of old terms



# Rocchio Feedback: Illustration



# Rocchio Feedback: Formula

The diagram illustrates the Rocchio Feedback formula. At the top, the word "Parameters" has three arrows pointing to the coefficients  $\alpha$ ,  $\frac{\beta}{|D_r|}$ , and  $\frac{\gamma}{|D_n|}$  in the formula. On the left, "New query" has an arrow pointing to  $\vec{q}_m$ . Below the formula, "Original query" has an arrow pointing to  $\vec{q}$ . "Rel docs" has an arrow pointing to the summation  $\sum_{\forall \vec{d}_j \in D_r} \vec{d}_j$ . "Non-rel docs" has an arrow pointing to the summation  $\sum_{\forall \vec{d}_j \in D_n} \vec{d}_j$ .

$$\vec{q}_m = \alpha \vec{q} + \frac{\beta}{|D_r|} \sum_{\forall \vec{d}_j \in D_r} \vec{d}_j - \frac{\gamma}{|D_n|} \sum_{\forall \vec{d}_j \in D_n} \vec{d}_j$$

Annotations:

- New query  $\rightarrow \vec{q}_m$
- Parameters  $\rightarrow \alpha, \frac{\beta}{|D_r|}, \frac{\gamma}{|D_n|}$
- Original query  $\rightarrow \vec{q}$
- Rel docs  $\rightarrow \sum_{\forall \vec{d}_j \in D_r} \vec{d}_j$
- Non-rel docs  $\rightarrow \sum_{\forall \vec{d}_j \in D_n} \vec{d}_j$

# Example of Rocchio Feedback

$V = \{\text{news about presidential camp. food ....}\}$

Query = "news about presidential campaign"

$O = (1, 1, 1, 1, 0, 0, \dots)$

New Query  $Q' = (\alpha*1 + \beta*1.5 - \gamma*1.5, \alpha*1 - \gamma*0.067, \alpha*1 + \beta*3.5, \alpha*1 + \beta*2.0 - \gamma*2.6, -\gamma*1.3, 0, 0, \dots)$

-  $D1 = (1.5, 0.1, 0, 0, 0, 0, \dots)$

D2

... news about organic food campaign...

-  $D2 = (1.5, 0.1, 0, 2.0, 2.0, 0, \dots)$

D3

... news of presidential campaign ...

+  $D3 = (1.5, 0, 3.0, 2.0, 0, 0, \dots)$

D4

+ Centroid Vector =  $((1.5+1.5)/2, 0, (3.0+4.0)/2, (2.0+2.0)/2, 0, 0, \dots)$   
 $= (1.5, 0, 3.5, 2.0, 0, 0, \dots)$

+  $D4 = (1.5, 0, 4.0, 2.0, 0, 0, \dots)$

- Centroid Vector =  $((1.5+1.5+1.5)/3, (0.1+0.1+0)/3, 0, (0+2.0+6.0)/3, (0+2.0+2.0)/3, 0, \dots)$   
 $= (1.5, 0.067, 0, 2.6, 1.3, 0, \dots)$

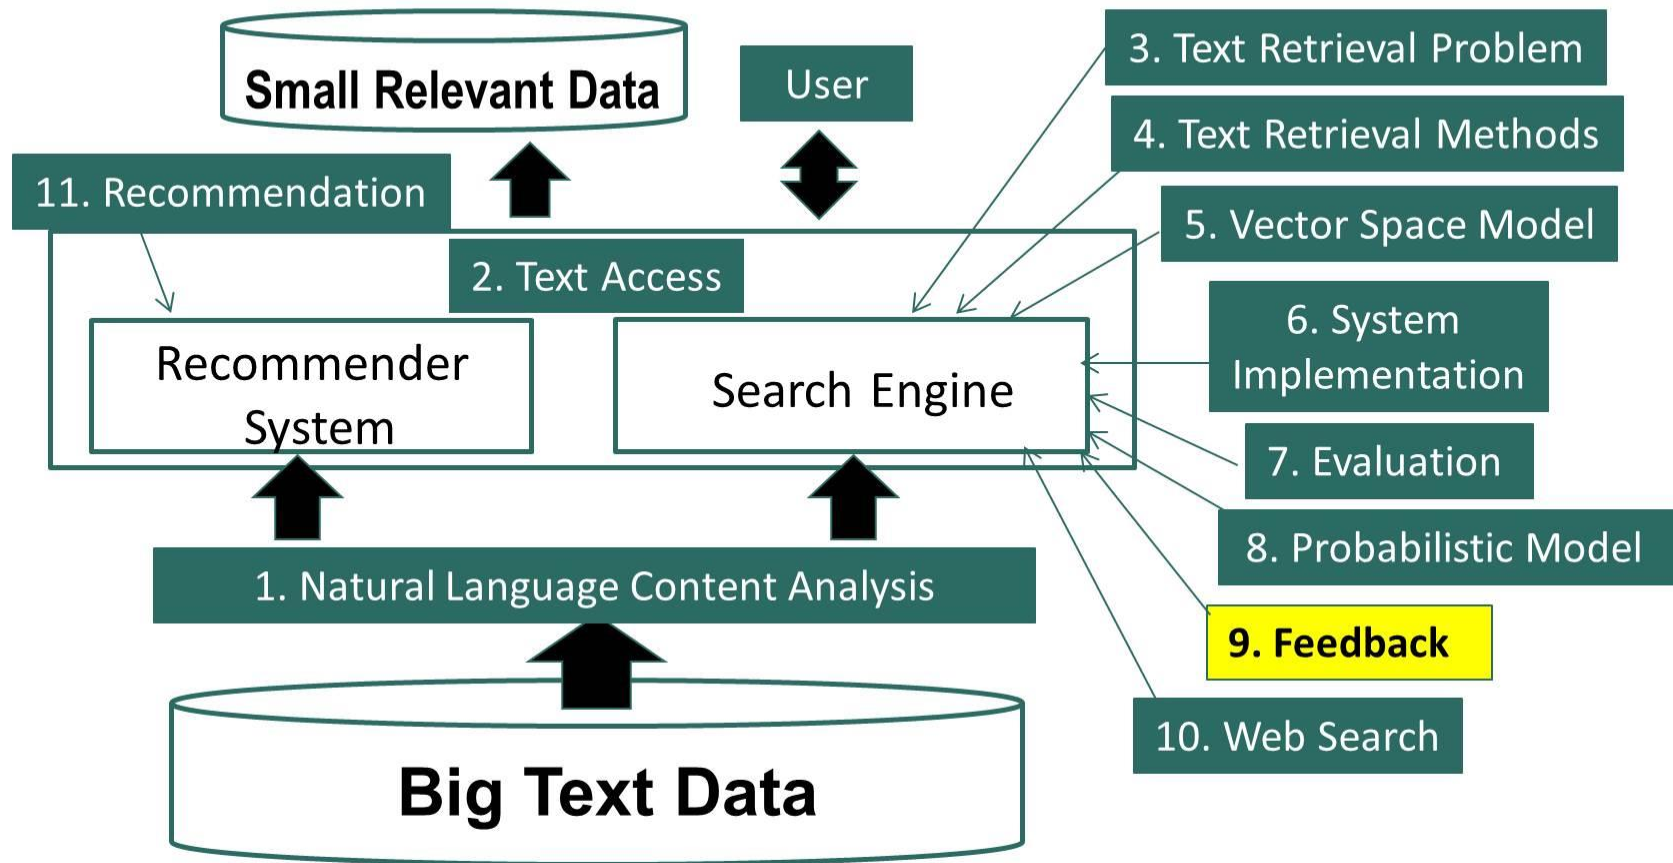
-  $D5 = (1.5, 0, 0, 6.0, 2.0, 0, \dots)$

# Rocchio in Practice

- Negative (non-relevant) examples are not very important (why?)
- Often truncate the vector (i.e., consider only a small number of words that have highest weights in the centroid vector) (efficiency concern)
- Avoid “over-fitting” (keep relatively high weight on the original query weights) (why?)
- Can be used for relevance feedback and pseudo feedback ( $\beta$  should be set to a larger value for relevance feedback than for pseudo feedback)
- Usually robust and effective

# Feedback in Text Retrieval: Feedback in LM

# Feedback in Text Retrieval: Feedback in LM




# Feedback with Language Models

- Query likelihood method can't naturally support relevance feedback
- Solution:
  - Kullback-Leibler (KL) divergence retrieval model as a generalization of query likelihood
  - Feedback is achieved through query model estimation/updating

# Kullback-Leibler (KL) Divergence Retrieval Model

Query Likelihood

$$f(q, d) = \sum_{\substack{w_i \in d \\ w_i \in q}} \boxed{c(w, q)} \left[ \log \frac{p_{\text{Seen}}(w_i | d)}{\alpha_d p(w_i | C)} \right] + n \log \alpha_d$$


KL-divergence  
(cross entropy)

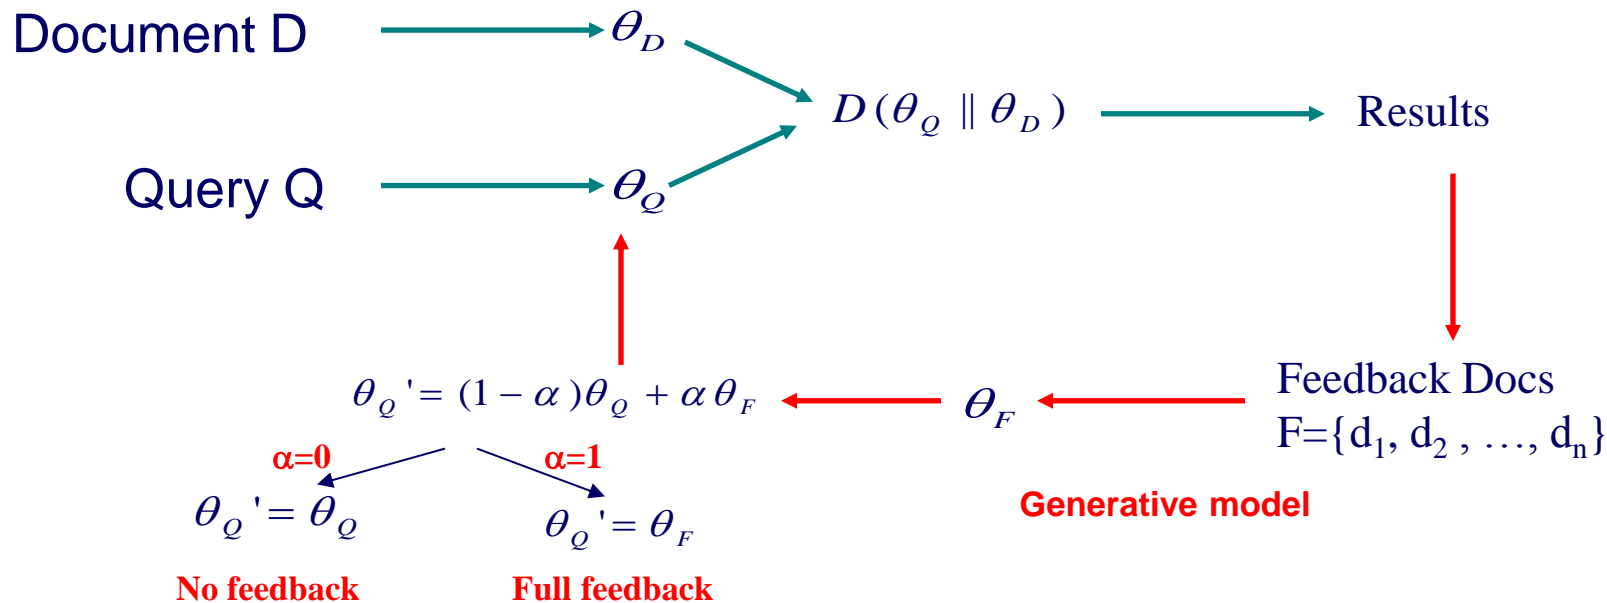
$$f(q, d) = \sum_{w \in d, p(w | \theta_Q) > 0} \boxed{[p(w | \hat{\theta}_Q)]} \log \frac{p_{\text{seen}}(w | d)}{\alpha_d p(w | C)} + \log \alpha_d$$

Query LM


$$p(w | \hat{\theta}_Q) = \frac{c(w, Q)}{|Q|}$$



# Feedback as Model Interpolation



# Generative Mixture Model



$$\log p(F | \theta) = \sum_i \sum_w c(w; d_i) \log[(1 - \lambda) p(w | \theta) + \lambda p(w | C)]$$

**Maximum Likelihood**  $\theta_F = \arg \max_{\theta} \log p(F | \theta)$

$\lambda$  = Noise in feedback documents

# Example of Pseudo-Feedback Query Model

Query: “**airport security**”

$\lambda=0.9$

W	$p(W \theta_F)$
security	0.0558
airport	0.0546
beverage	0.0488
alcohol	0.0474
bomb	0.0236
terrorist	0.0217
author	0.0206
license	0.0188
bond	0.0186
counter-terror	0.0173
terror	0.0142
newsnet	0.0129
attack	0.0124
operation	0.0121
headline	0.0121

Mixture model  
approach

Web database

Top 10 docs

$\lambda=0.7$

W	$p(W \theta_F)$
the	0.0405
security	0.0377
airport	0.0342
beverage	0.0305
alcohol	0.0304
to	0.0268
of	0.0241
and	0.0214
author	0.0156
bomb	0.0150
terrorist	0.0137
in	0.0135
license	0.0127
state	0.0127
by	0.0125

# Summary of Feedback in Text Retrieval

- Feedback = learn from examples
- Three major feedback scenarios
  - Relevance, pseudo, and implicit feedback
- Rocchio for VSM
- Query model estimation for LM
  - Mixture model
  - Many other methods (e.g., relevance model) have been proposed [Zhai 08]

# Additional Readings


- ChengXiang Zhai, *Statistical Language Models for Information Retrieval* (Synthesis Lectures Series on Human Language Technologies), Morgan & Claypool Publishers, 2008.

<http://www.morganclaypool.com/doi/abs/10.2200/S00158ED1V01Y200811HLT001>

- Victor Lavrenko and W. Bruce Croft. 2001. Relevance based language models. In *Proceedings of ACM SIGIR 2001*.

# Web Search

# Course Schedule



# Web Search: Challenges & Opportunities

- Challenges

- Scalability

**→ Parallel indexing & searching (MapReduce)**

- How to handle the size of the Web and ensure completeness of coverage?
    - How to serve many user queries quickly?

- Low quality information and spams

**→ Spam detection  
& Robust ranking**

- Dynamics of the Web

- New pages are constantly created and some pages may be updated very quickly

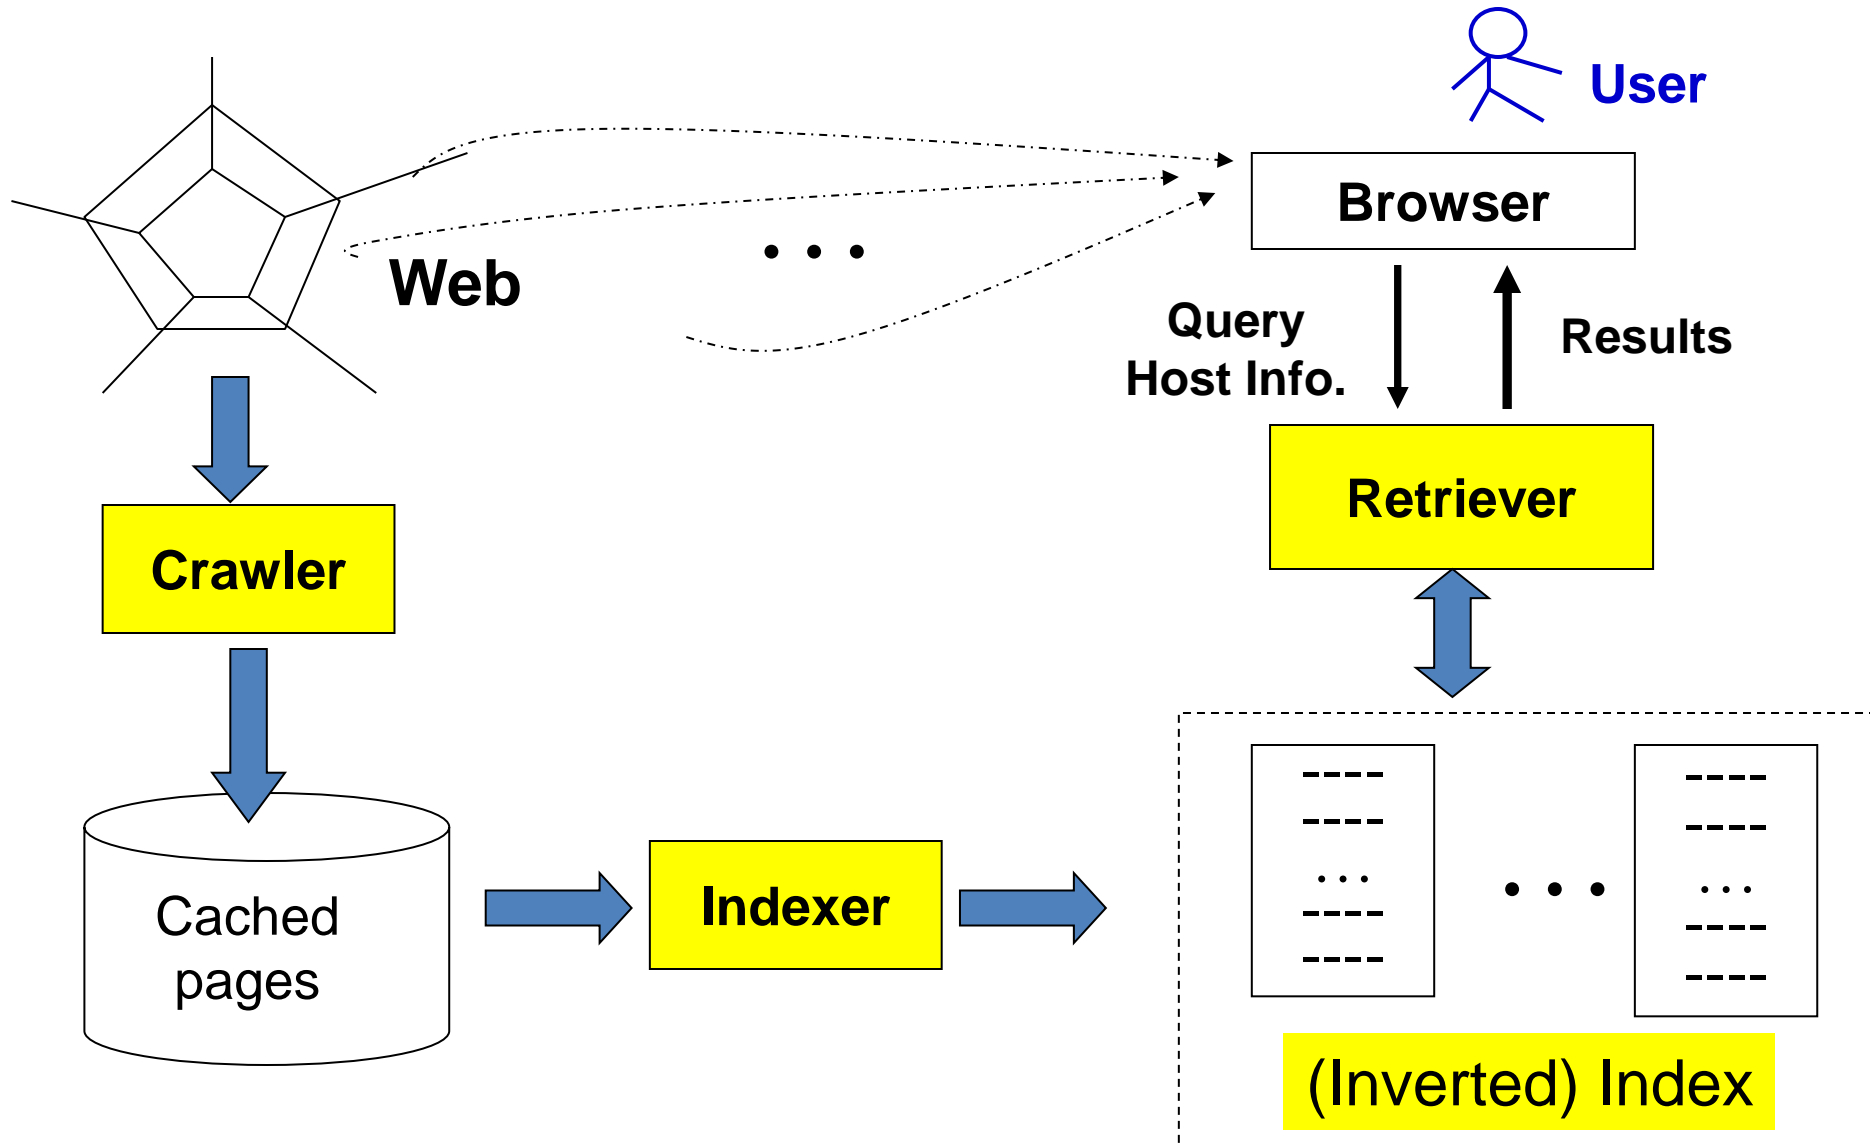
- Opportunities

- many additional heuristics (e.g., links) can be leveraged to improve search accuracy

**→ Link analysis & multi-feature ranking**



# Basic Search Engine Technologies



# Component I: Crawler/Spider/Robot

- Building a “toy crawler” is easy
  - Start with a set of “seed pages” in a priority queue
  - Fetch pages from the web
  - Parse fetched pages for hyperlinks; add them to the queue
  - Follow the hyperlinks in the queue
- A real crawler is much more complicated...
  - Robustness (server failure, trap, etc.)
  - Crawling courtesy (server load balance, robot exclusion, etc.)
  - Handling file types (images, PDF files, etc.)
  - URL extensions (cgi script, internal references, etc.)
  - Recognize redundant pages (identical and duplicates)
  - Discover “hidden” URLs (e.g., truncating a long URL )

# Major Crawling Strategies

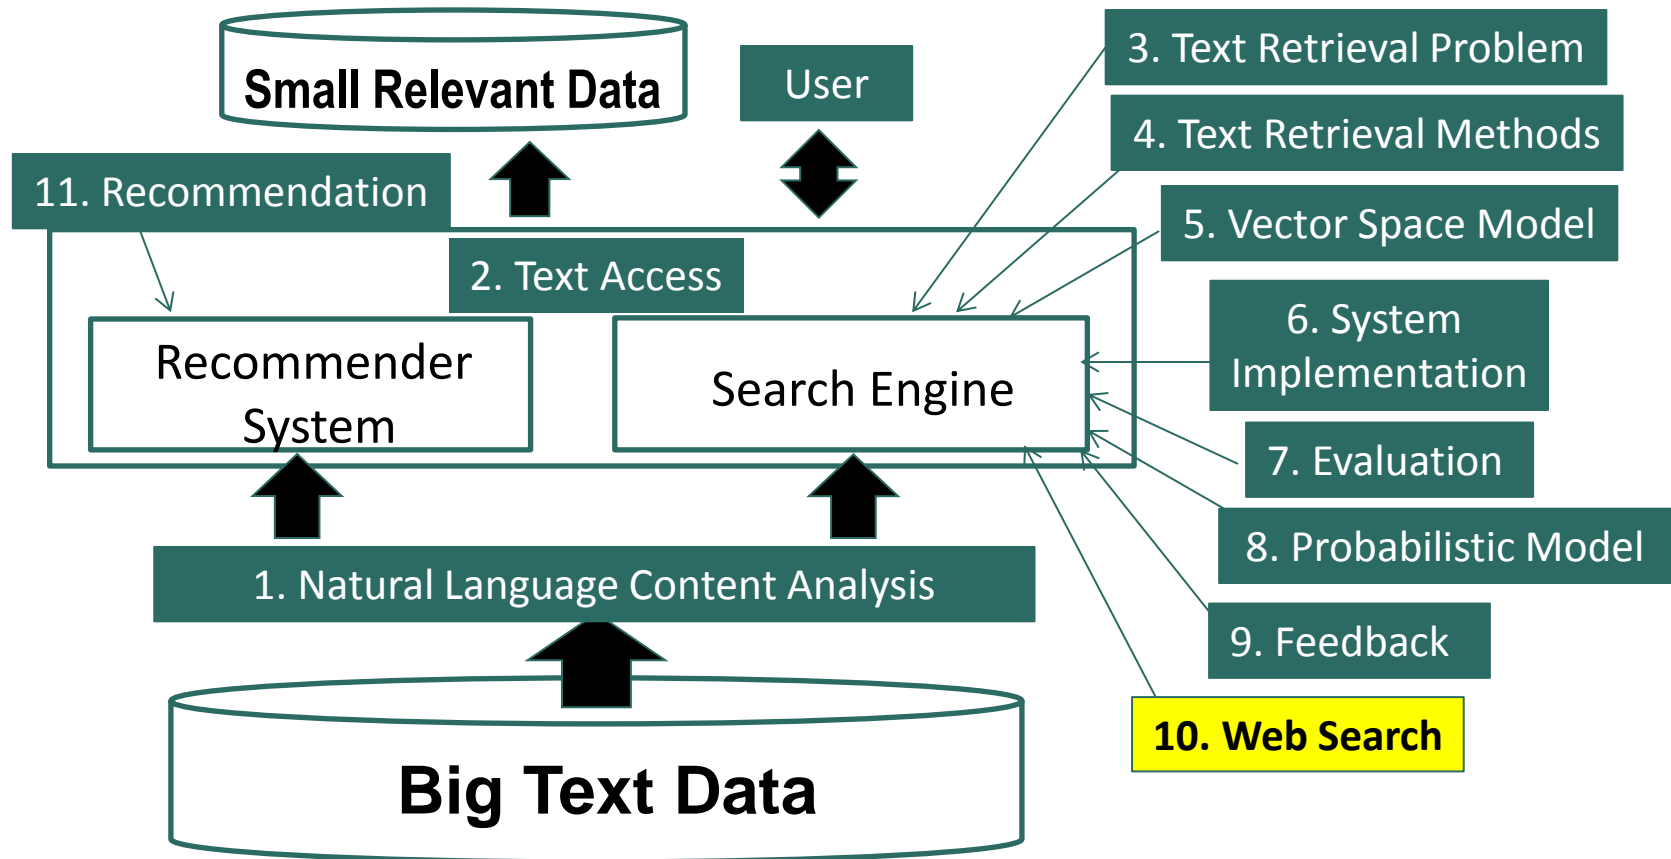
- Breadth-First is common (balance server load)
- Parallel crawling is natural
- Variation: focused crawling
  - Targeting at a subset of pages (e.g., all pages about “automobiles” )
  - Typically given a query
- How to find new pages (they may not linked to an old page!)
- Incremental/repeated crawling
  - Need to minimize resource overhead
  - Can learn from the past experience (updated daily vs. monthly)
  - Target at : 1) frequently updated pages; 2) frequently accessed pages

# Summary

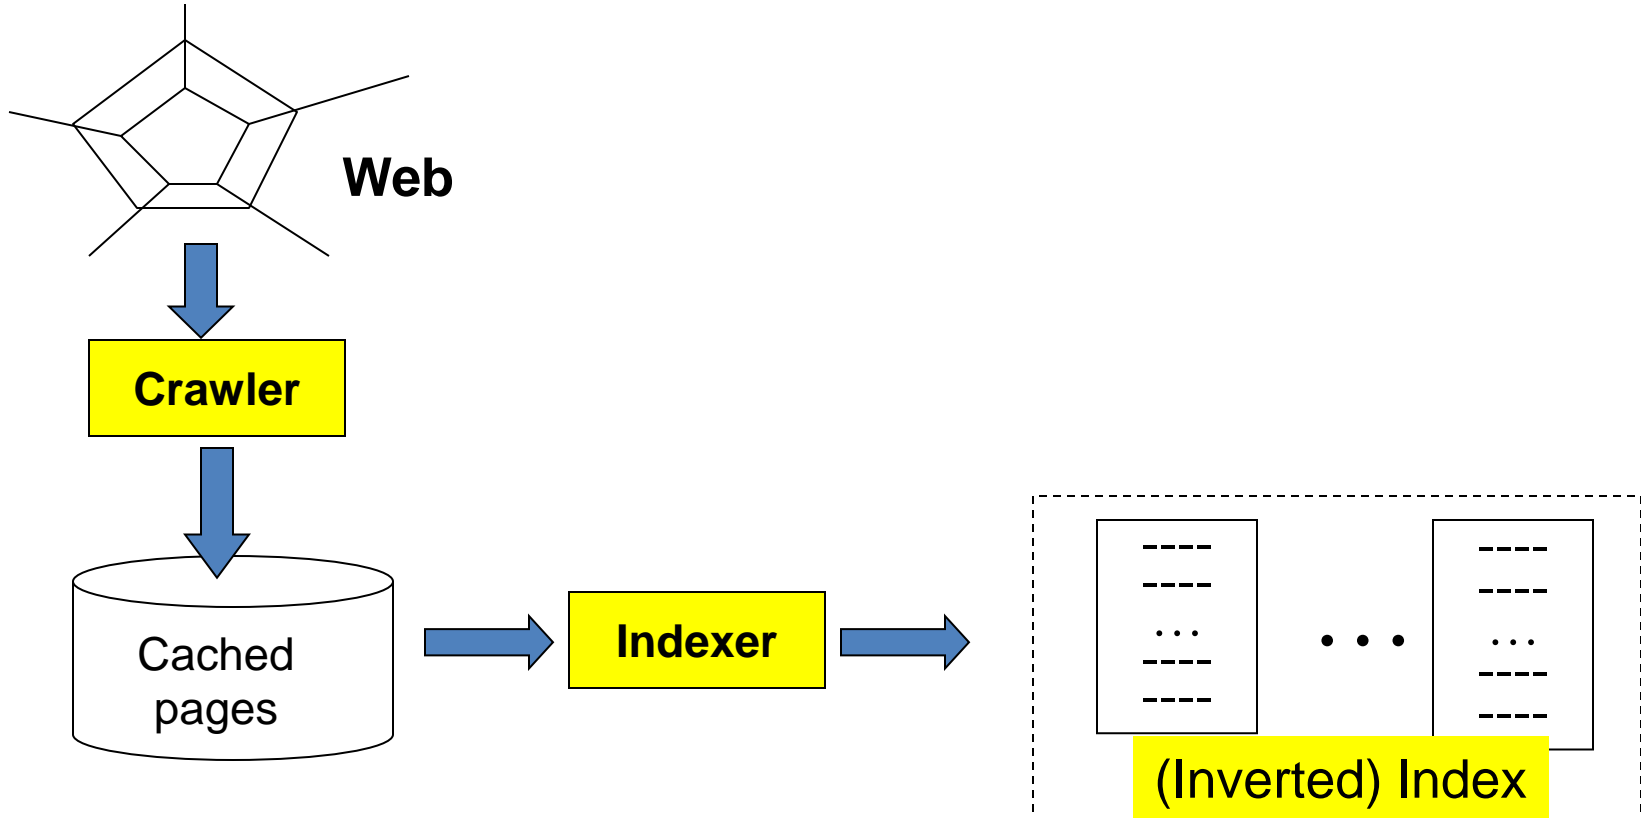
- Web search is one of the most important applications of text retrieval
  - New challenges: scalability, efficiency, quality of information
  - New opportunities: rich link information, layout, etc
- Crawler is an essential component of Web search applications
  - Initial crawling: complete vs. focused
  - Incremental crawling: resource optimization

# Web Search: Web Indexing

# Web Search: Web Indexing



# Basic Search Engine Technologies



# Overview of Web Indexing

- Standard IR techniques are the basis, but insufficient
  - Scalability
  - Efficiency
- Google's contributions:
  - Google File System (GFS): distributed file system
  - MapReduce: Software framework for parallel computation
  - Hadoop: Open source implementation of MapReduce



# GFS Architecture

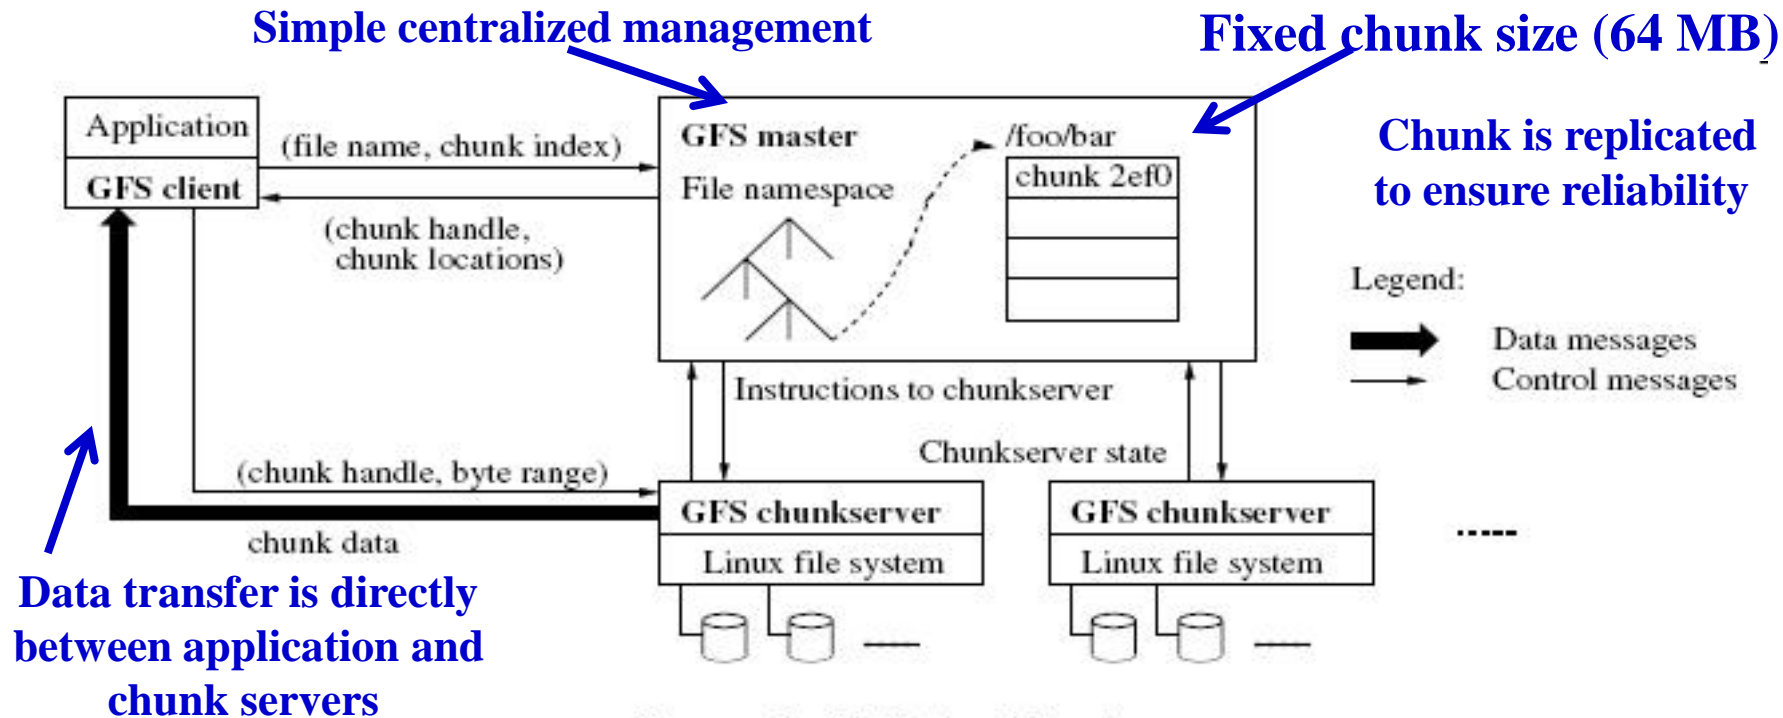


Figure 1: GFS Architecture

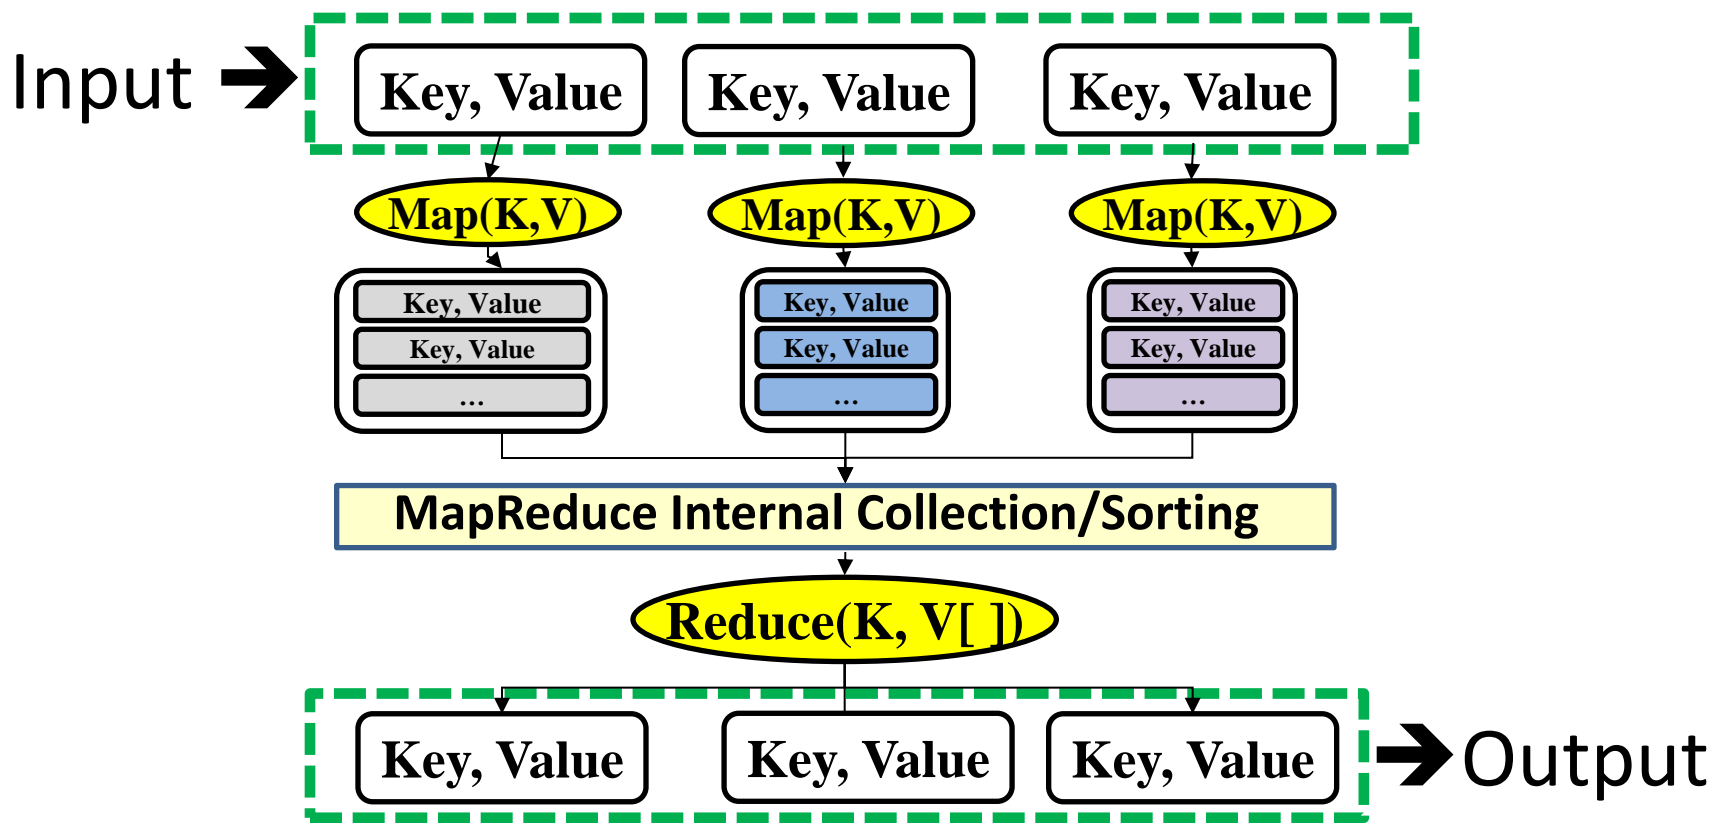
GHEMAWAT, S., GOBIOFF, H., AND LEUNG, S.-T. The google file system. In SOSP '03: Proceedings of the nineteenth ACM symposium on Operating systems principles (New York, NY, USA, 2003), ACM, pp. 29–43.

<http://static.googleusercontent.com/media/research.google.com/en/us/archive/gfs-sosp2003.pdf>

# MapReduce: A Framework for Parallel Programming

- Minimize effort of programmer for simple parallel processing tasks
- Features
  - Hide many low-level details (network, storage)
  - Built-in fault tolerance
  - Automatic load balancing

# MapReduce: Computation Pipeline



Slide adapted from Alexander Behm & Ajey Shah's presentation (<http://www.slideshare.net/gothicane/behm-shah-pagerank>)

# Word Counting

## Input: Text Data

Hello World Bye World  
Hello Hadoop Bye Hadoop  
Bye Hadoop Hello Hadoop  
... ..



## Output: Count of each word

Bye 3  
Hadoop 4  
Hello 3  
World 2  
...

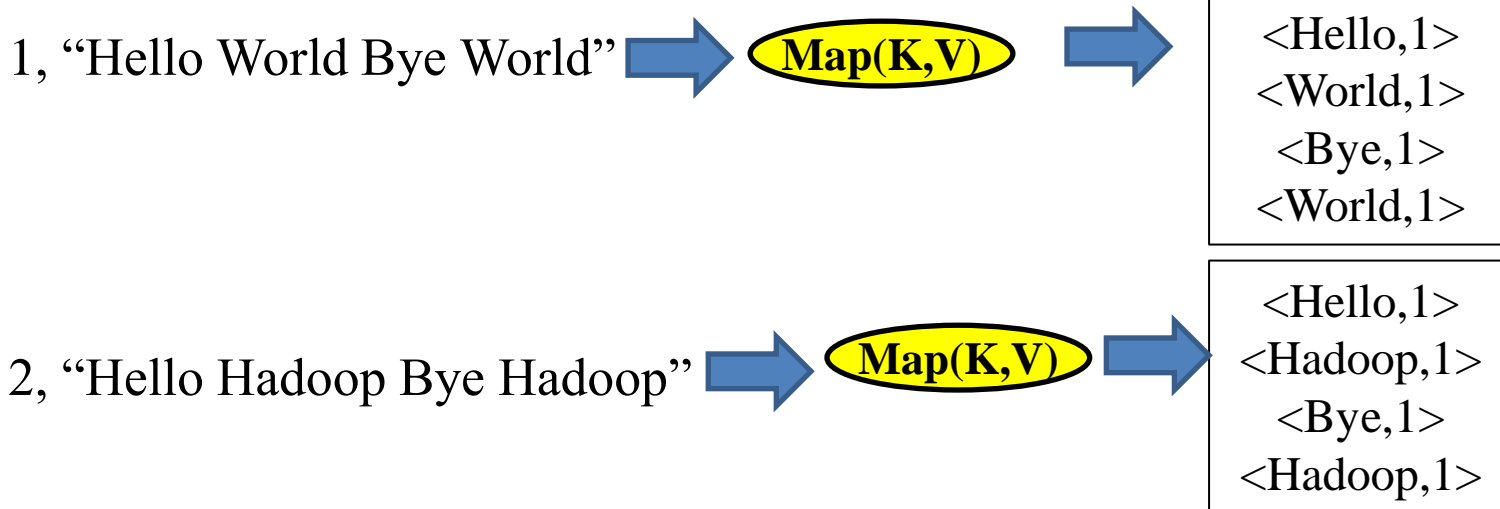
**How can we do this within the MapReduce framework?**

Slide adapted from Alexander Behm & Ajey Shah's presentation (<http://www.slideshare.net/gothicane/behm-shah-pagerank>)

# Word Counting: Map Function

**Input**

**Output**

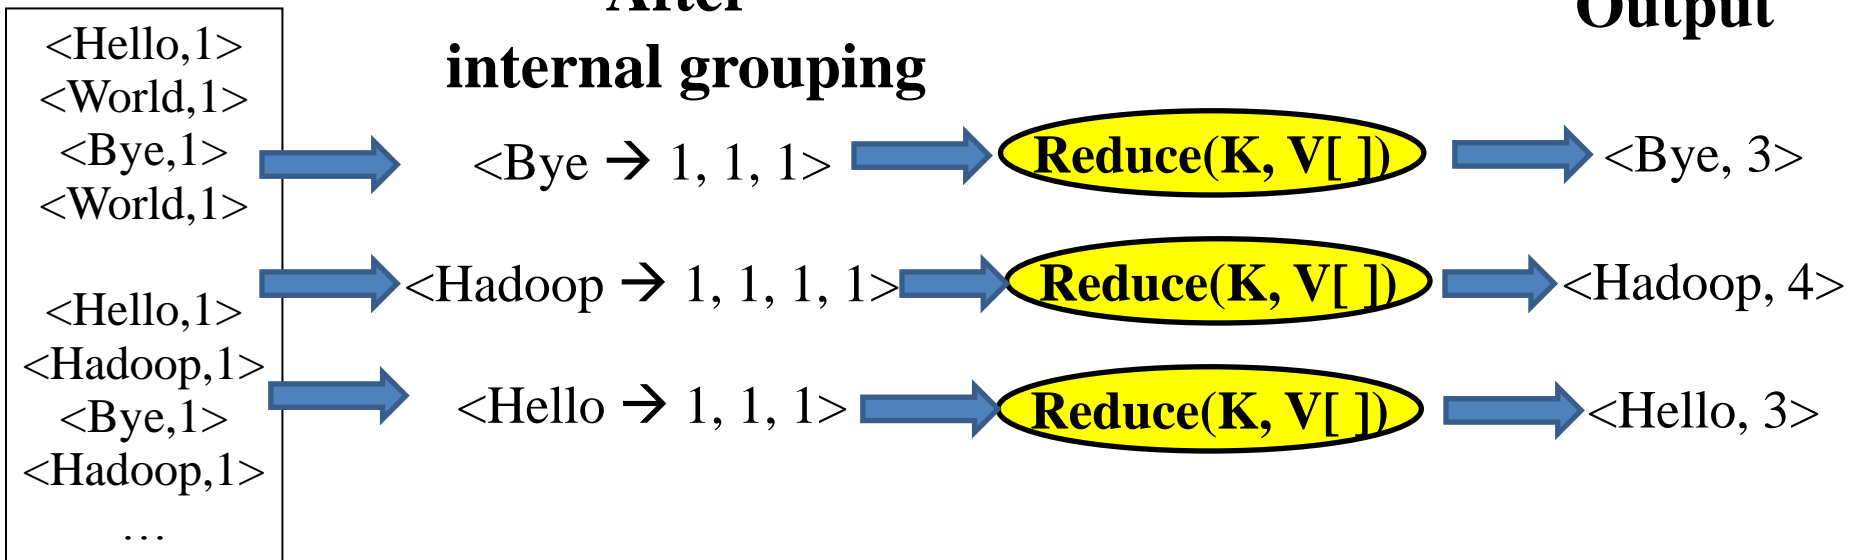


....

**Map(K, V)**  
**{ For each word w in V, Collect(w, 1);}**

# Word Counting: Reduce Function

## Map Output

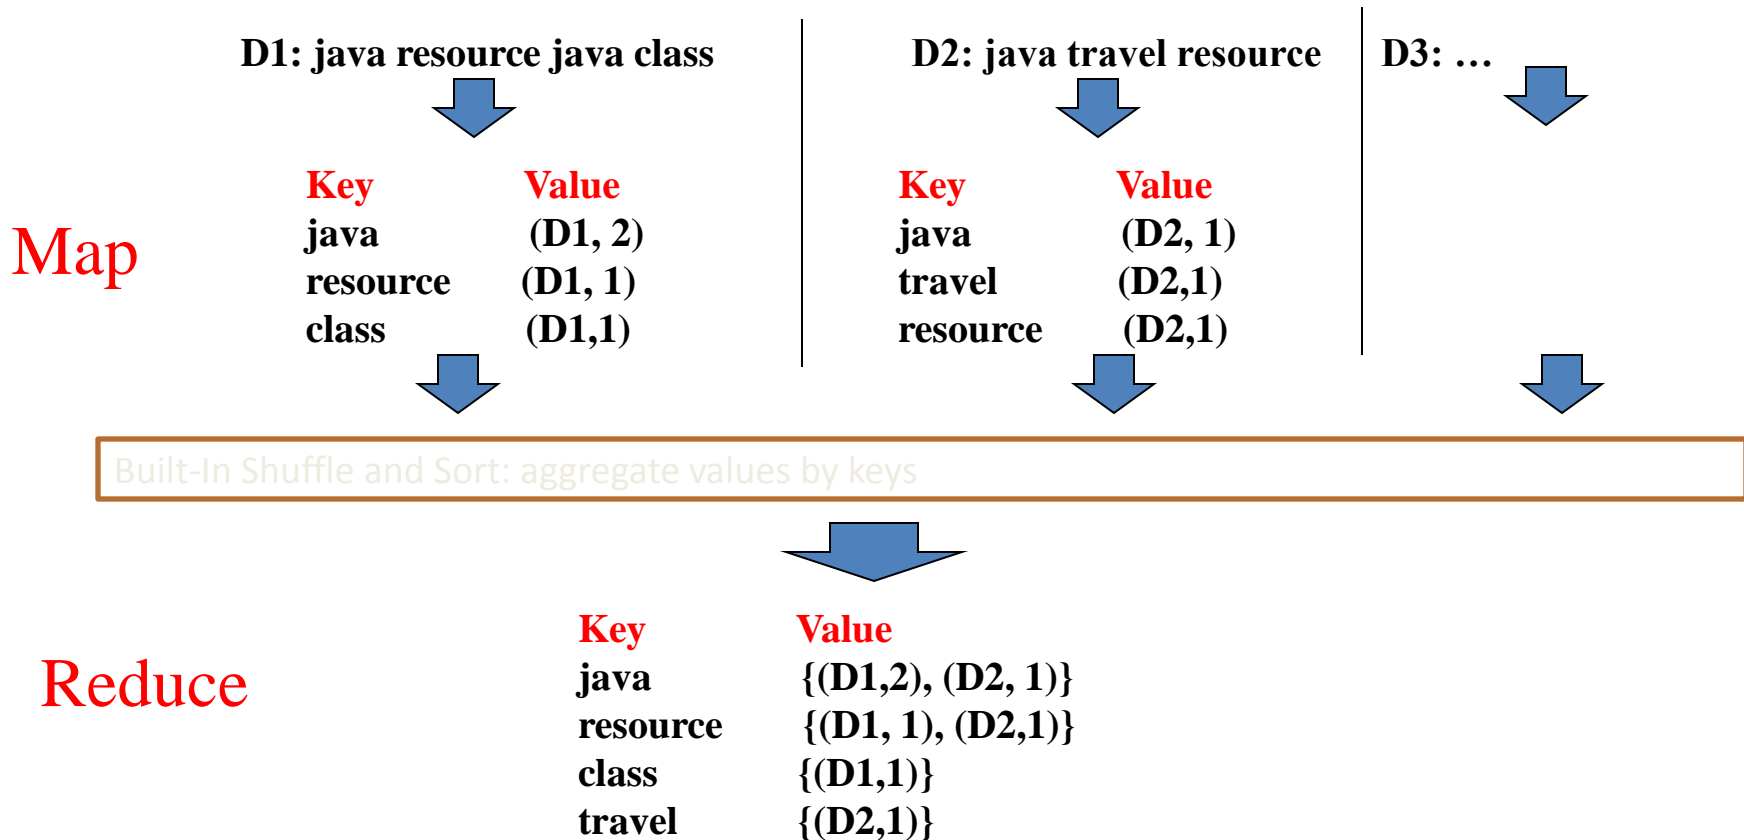


**Reduce(K, V[ ])**

```
{ Int count = 0; For each v in V, count += v; Collect(K, count); }
```

Slide adapted from Alexander Behm & Ajey Shah's presentation (<http://www.slideshare.net/gothicane/behm-shah-pagerank>)

# Inverted Indexing with MapReduce



# Inverted Indexing: Pseudo-Code

```
1: class MAPPER
2:   procedure MAP(docid  $n$ , doc  $d$ )
3:      $H \leftarrow$  new ASSOCIATIVEARRAY
4:     for all term  $t \in$  doc  $d$  do
5:        $H\{t\} \leftarrow H\{t\} + 1$ 
6:     for all term  $t \in H$  do
7:       EMIT(term  $t$ , posting  $\langle n, H\{t\} \rangle$ )

1: class REDUCER
2:   procedure REDUCE(term  $t$ , postings  $[\langle a_1, f_1 \rangle, \langle a_2, f_2 \rangle \dots]$ )
3:      $P \leftarrow$  new LIST
4:     for all posting  $\langle a, f \rangle \in$  postings  $[\langle a_1, f_1 \rangle, \langle a_2, f_2 \rangle \dots]$  do
5:       APPEND( $P, \langle a, f \rangle$ )
6:     SORT( $P$ )
7:     EMIT(term  $t$ , postings  $P$ )
```

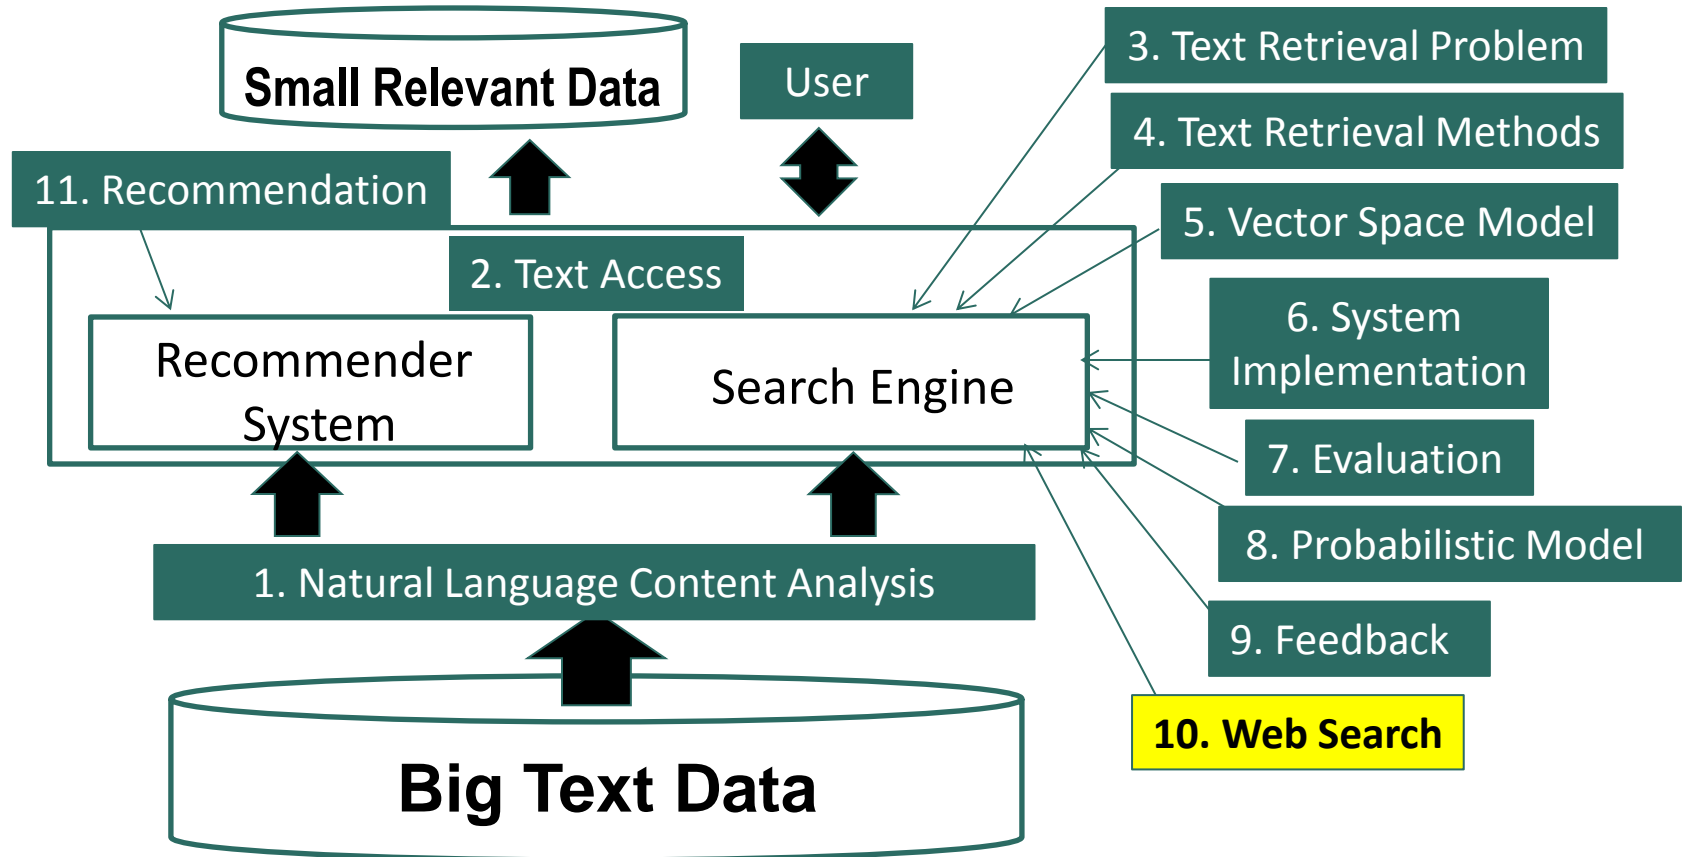


# Summary

- Web scale indexing requires
  - Storing the index on multiple machines (GFS)
  - Creating the index in parallel (MapReduce)
- Both GFS and MapReduce are general infrastructures

# Web Search: Link Analysis

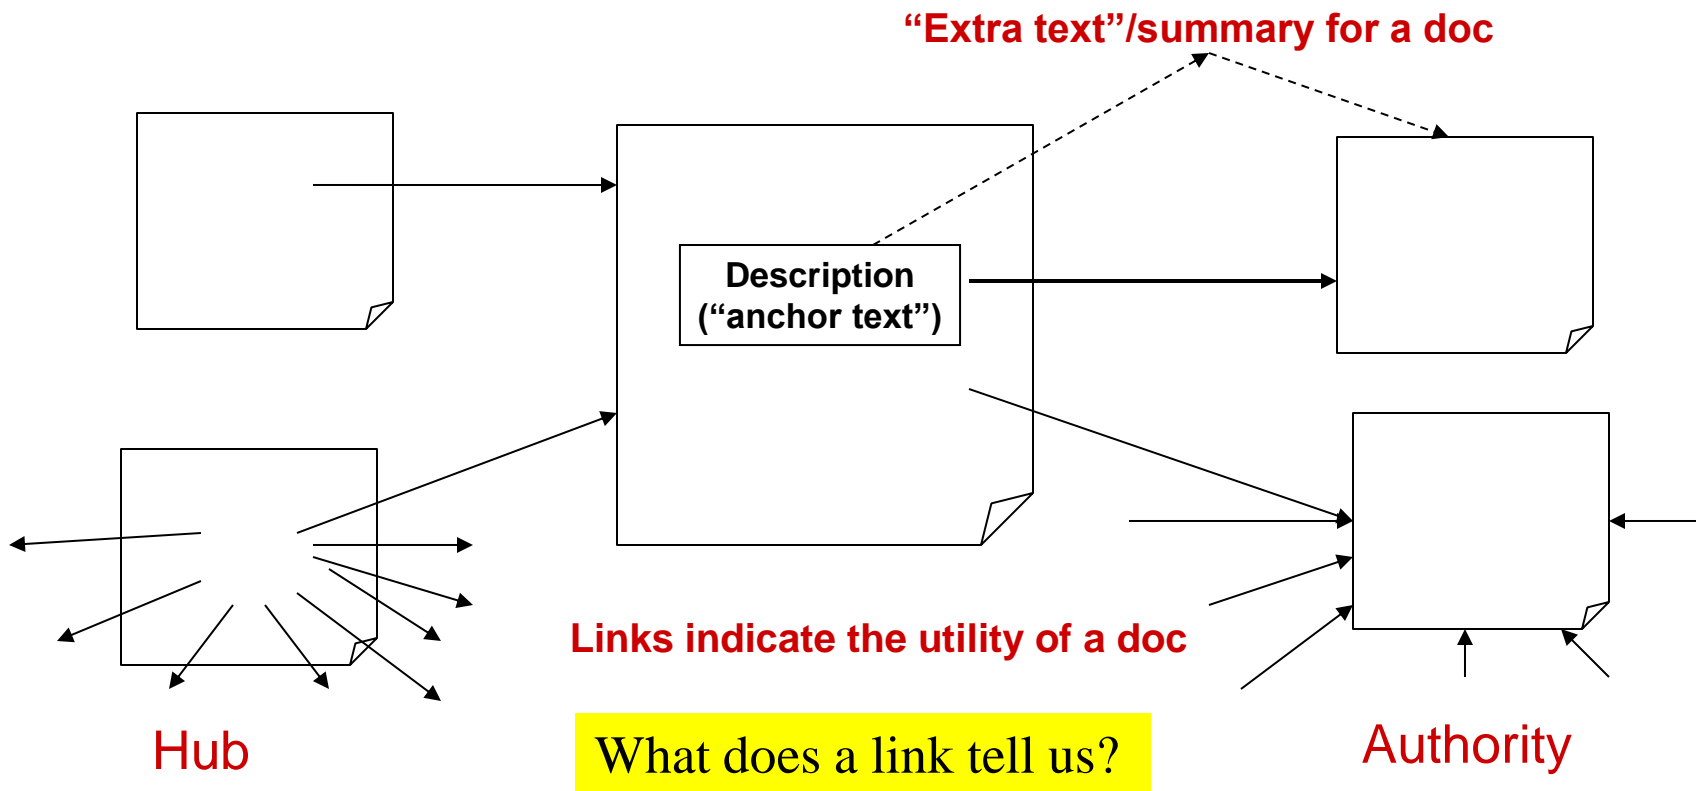
# Web Search: Link Analysis



# Ranking Algorithms for Web Search

- Standard IR models apply but aren't sufficient
  - Different information needs
  - Documents have additional information
  - Information quality varies a lot
- Major extensions
  - Exploiting links to improve scoring
  - Exploiting clickthroughs for massive implicit feedback
  - In general, rely on machine learning to combine all kinds of features

# Exploiting Inter-Document Links




# PageRank: Capturing Page “Popularity”

- Intuitions
  - Links are like citations in literature
  - A page that is cited often can be expected to be more useful in general
- PageRank is essentially “citation counting”, but improves over simple counting
  - Consider “indirect citations” (being cited by a highly cited paper counts a lot...)
  - Smoothing of citations (every page is assumed to have a non-zero pseudo citation count)
- PageRank can also be interpreted as random surfing (thus capturing popularity)

# The PageRank Algorithm

Random surfing model: At any page,  
 With prob.  $\alpha$ , randomly jumping to another page  
 With prob.  $(1-\alpha)$ , randomly picking a link to follow.

**$p(d_i)$ : PageRank score of  $d_i$  = average probability of visiting page  $d_i$**



Transition matrix

$$M = \begin{bmatrix} 0 & 0 & 1/2 & 1/2 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 1/2 & 1/2 & 0 & 0 \end{bmatrix}$$

$M_{ij}$  = probability of going  
from  $d_i$  to  $d_j$

$$\sum_{j=1}^N M_{ij} = 1$$

probability of visiting page  $d_j$  at time  $t+1$

probability of at page  $d_i$  at time  $t$

**“Equilibrium Equation”:**

$$p_{t+1}(d_j) = (1-\alpha) \underbrace{\sum_{i=1}^N M_{ij} p_t(d_i)}_{\text{Reach } d_j \text{ via following a link}} + \alpha \underbrace{\sum_{i=1}^N \frac{1}{N} p_t(d_i)}_{\text{Reach } d_j \text{ via random jumping}}$$

$N = \# \text{ pages}$

Reach  $d_j$  via following a link

Reach  $d_j$  via random jumping

**dropping the time index**

$$p(d_j) = \sum_{i=1}^N \left[ \frac{1}{N} \alpha + (1-\alpha) M_{ij} \right] p(d_i)$$




$$\bar{p} = (\alpha \mathbf{I} + (1-\alpha) M)^T \bar{p}$$

$$\mathbf{I}_{ij} = 1/N$$

We can solve the equation with an iterative algorithm

# PageRank: Example



$$p(d_j) = \sum_{i=1}^N \left[ \frac{1}{N} \alpha + (1 - \alpha) M_{ij} \right] p(d_i)$$

$$\vec{p} = (\alpha I + (1 - \alpha) M)^T \vec{p}$$

$$A = (1 - 0.2)M + 0.2I = 0.8 \times \begin{bmatrix} 0 & 0 & 1/2 & 1/2 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 1/2 & 1/2 & 0 & 0 \end{bmatrix} + 0.2 \times \begin{bmatrix} 1/4 & 1/4 & 1/4 & 1/4 \\ 1/4 & 1/4 & 1/4 & 1/4 \\ 1/4 & 1/4 & 1/4 & 1/4 \\ 1/4 & 1/4 & 1/4 & 1/4 \end{bmatrix}$$

$$\begin{bmatrix} p^{n+1}(d_1) \\ p^{n+1}(d_2) \\ p^{n+1}(d_3) \\ p^{n+1}(d_4) \end{bmatrix} = A^T \begin{bmatrix} p^n(d_1) \\ p^n(d_2) \\ p^n(d_3) \\ p^n(d_4) \end{bmatrix} = \begin{bmatrix} 0.05 & 0.85 & 0.05 & 0.45 \\ 0.05 & 0.05 & 0.85 & 0.45 \\ 0.45 & 0.05 & 0.05 & 0.05 \\ 0.45 & 0.05 & 0.05 & 0.05 \end{bmatrix} \times \begin{bmatrix} p^n(d_1) \\ p^n(d_2) \\ p^n(d_3) \\ p^n(d_4) \end{bmatrix}$$

$$p^{n+1}(d_1) = 0.05 * p^n(d_1) + 0.85 * p^n(d_2) + 0.05 * p^n(d_3) + 0.45 * p^n(d_4)$$

**Initial value  $p(d)=1/N$ ,      iterate until converge**

Do you see how scores are propagated over the graph?




# PageRank in Practice

- Computation can be quite efficient since  $M$  is usually sparse
- Normalization doesn't affect ranking, leading to some variants of the formula
- The zero-outlink problem:  $p(d_i)$ 's don't sum to 1
  - One possible solution = page-specific damping factor ( $\alpha=1.0$  for a page with no outlink)
- Many extensions (e.g., topic-specific PageRank)
- Many other applications (e.g., social network analysis)

# HITS: Capturing Authorities & Hubs

- Intuitions
  - Pages that are widely cited are good authorities
  - Pages that cite many other pages are good hubs
- The key idea of HITS (Hypertext-Induced Topic Search)
  - Good authorities are cited by good hubs
  - Good hubs point to good authorities
  - Iterative reinforcement...
- Many applications in graph/network analysis

# The HITS Algorithm



$$A = \begin{bmatrix} 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \end{bmatrix}$$

“Adjacency matrix”

Initial values:  $a(d_i)=h(d_i)=1$

$$h(d_i) = \sum_{d_j \in OUT(d_i)} a(d_j)$$

$$a(d_i) = \sum_{d_j \in IN(d_i)} h(d_j)$$

Iterate

Normalize:

$$\bar{h} = A\bar{a}; \quad \bar{a} = A^T \bar{h}$$

$$\bar{h} = AA^T \bar{h}; \quad \bar{a} = A^T A \bar{a}$$

$$\sum_i a(d_i)^2 = \sum_i h(d_i)^2 = 1$$

# Summary

- Link information is very useful
  - Anchor text
  - PageRank
  - HITS
- Both PageRank and HITS have many applications in analyzing other graphs or networks