


Probabilistic Topic Models: Mixture of Unigram Language Models

Probabilistic Topic Models: Mixture of Unigram LMs



Factoring out Background Words

d

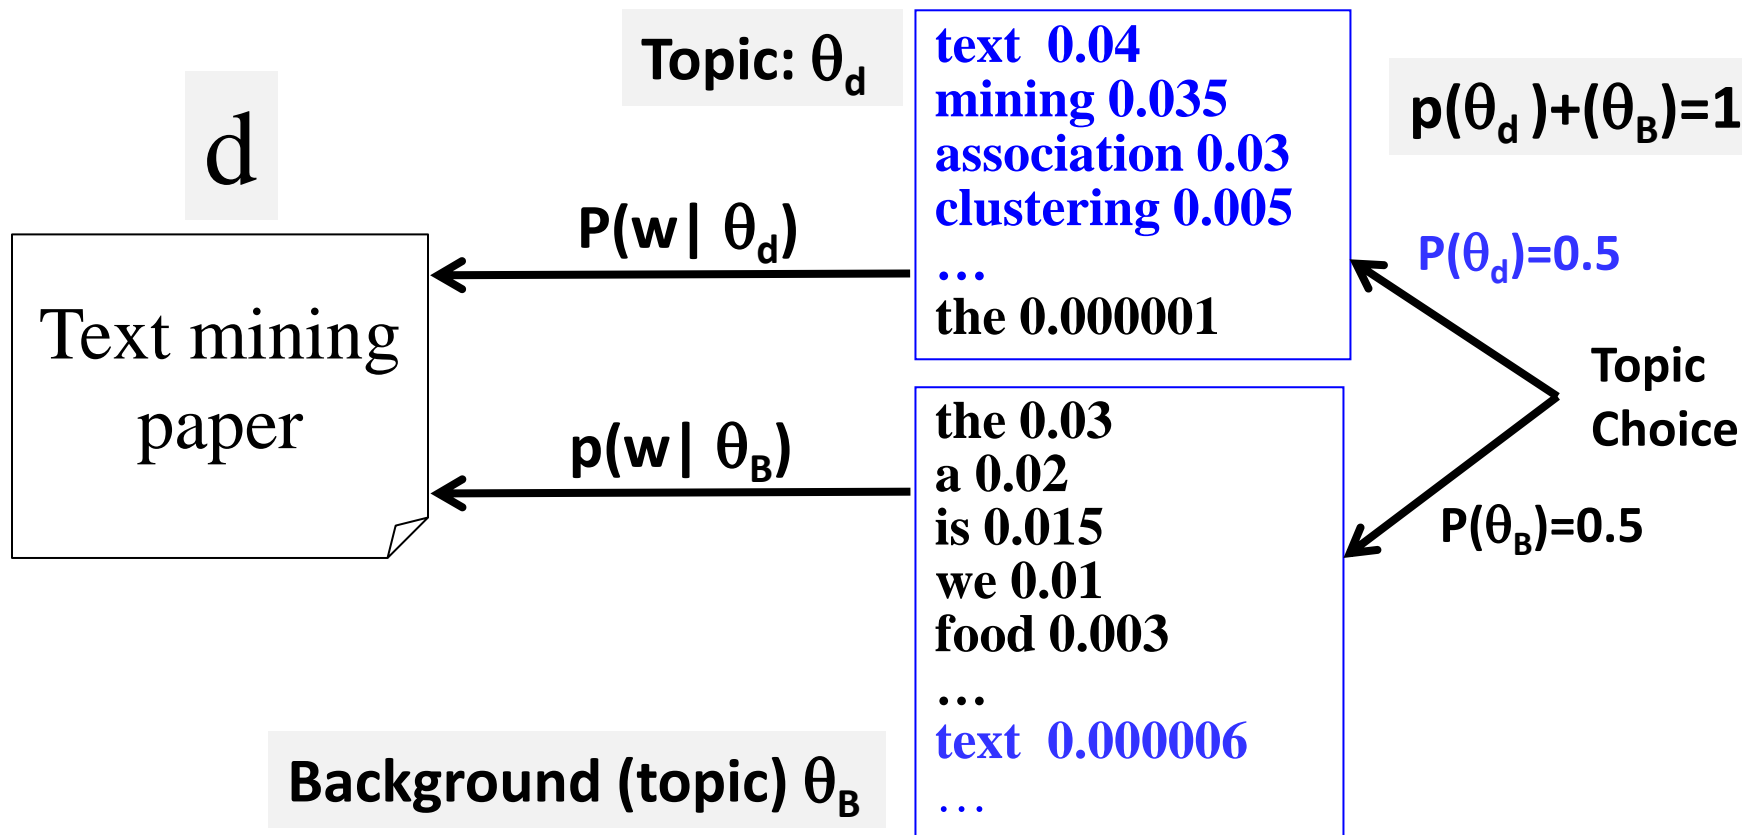
Text mining
paper

$p(w | \theta)$

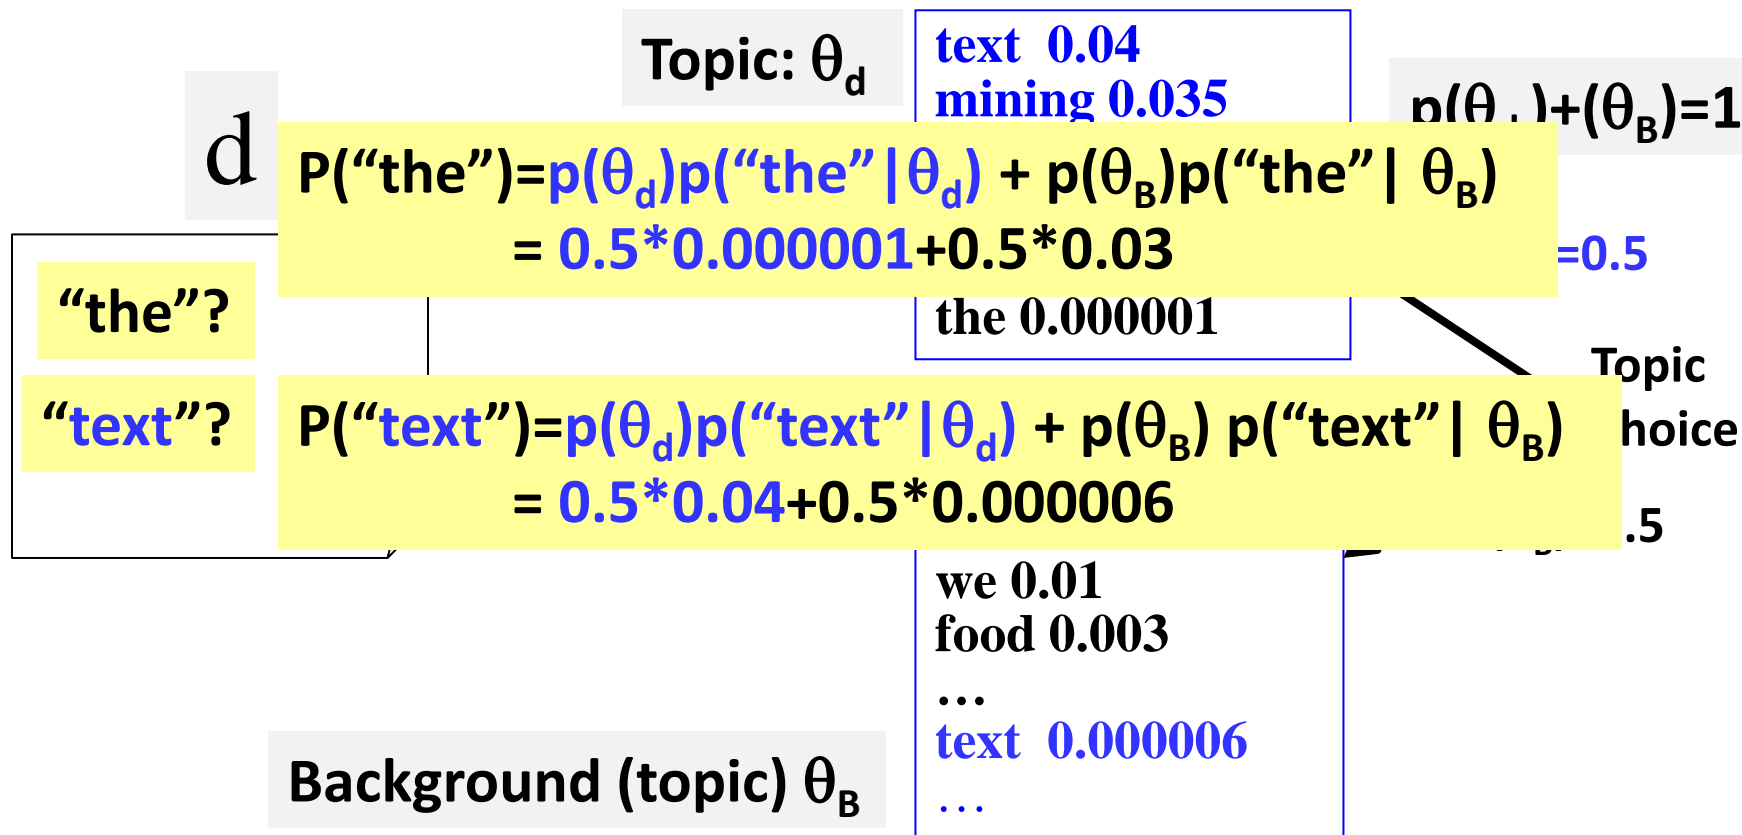
the 0.031
a 0.018
...
text 0.04
mining 0.035
association 0.03
clustering 0.005
computer 0.0009
...
food 0.000001
...

How can we get rid of
these common words?

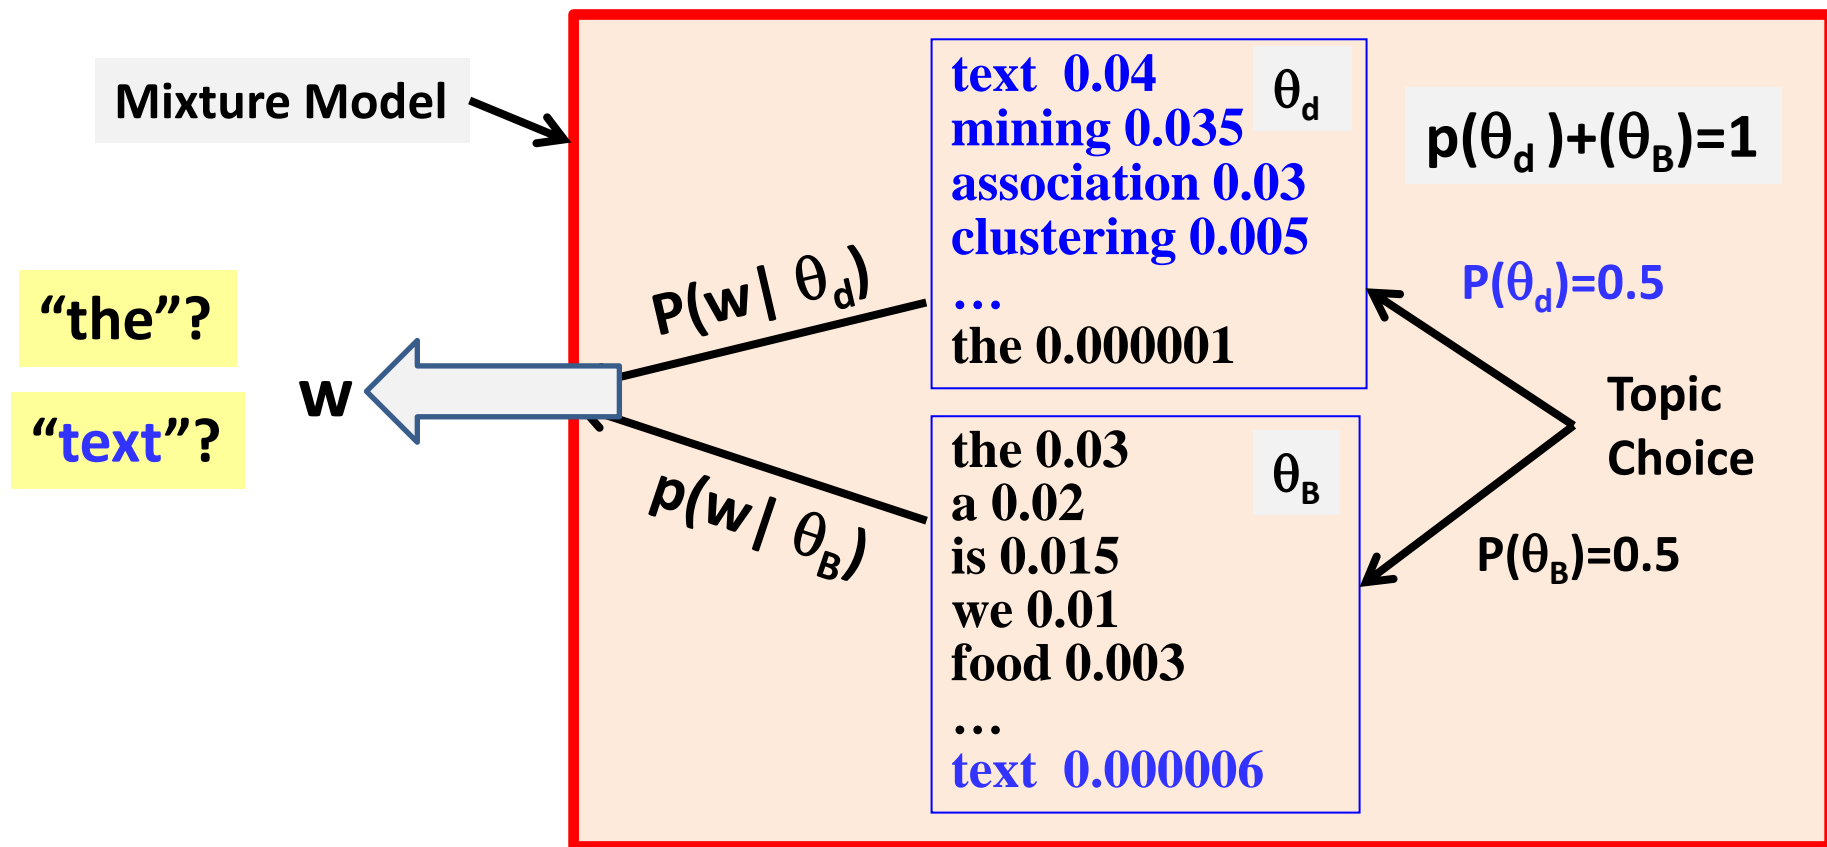
Generate d Using Two Word Distributions



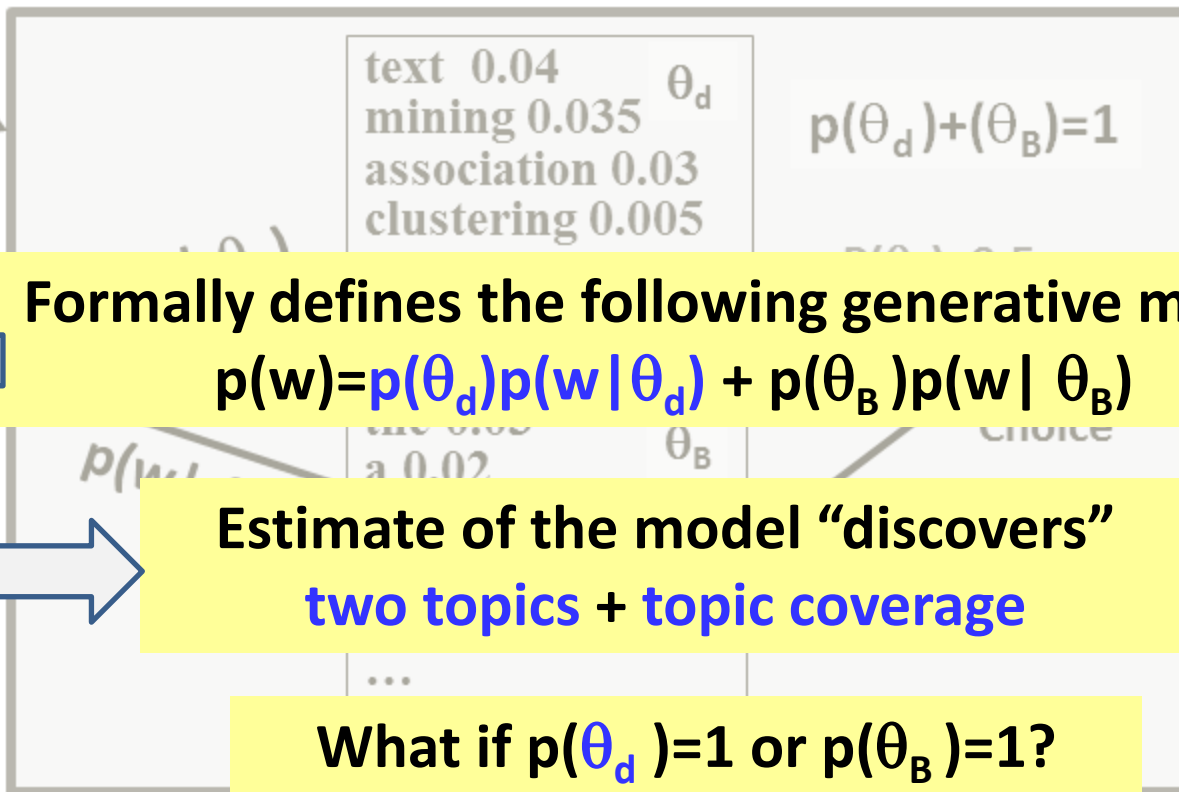
What's the probability of observing a word w?



The Idea of a Mixture Model



As a Generative Model...



Mixture of Two Unigram Language Models

- **Data:** Document d
- **Mixture Model: parameters** $\Lambda = (\{p(w | \theta_d)\}, \{p(w | \theta_B)\}, p(\theta_B), p(\theta_d))$
 - Two unigram LMs: θ_d (the topic of d); θ_B (background topic)
 - Mixing weight (topic choice): $p(\theta_d) + p(\theta_B) = 1$

- **Likelihood function:**

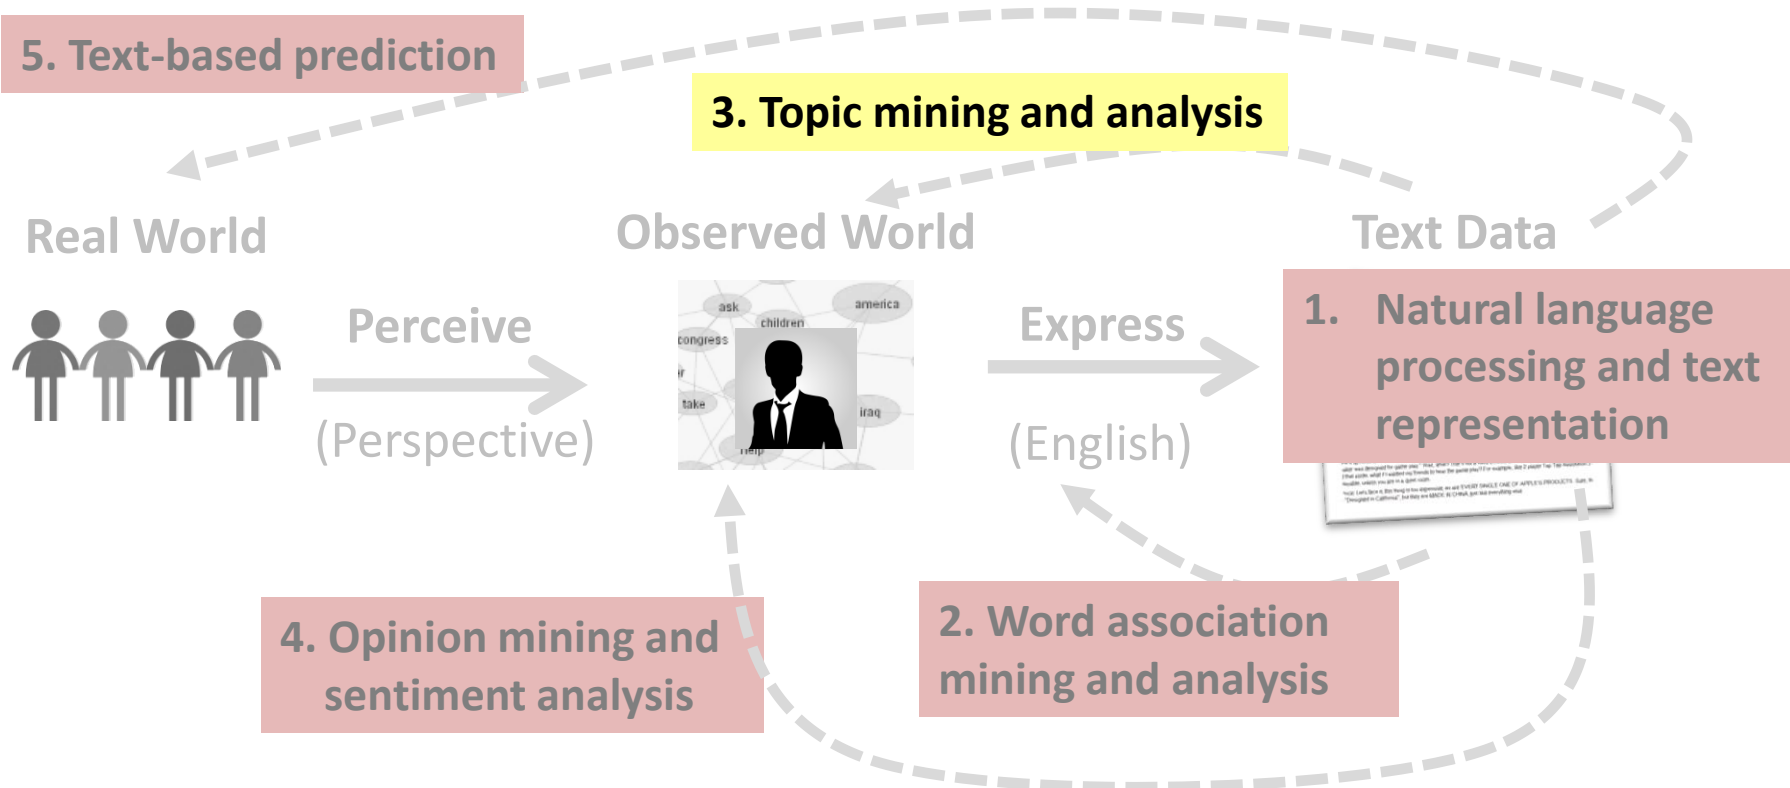
$$\begin{aligned} p(d | \Lambda) &= \prod_{i=1}^{|d|} p(x_i | \Lambda) = \prod_{i=1}^{|d|} [p(\theta_d)p(x_i | \theta_d) + p(\theta_B)p(x_i | \theta_B)] \\ &= \prod_{i=1}^M [p(\theta_d)p(w_i | \theta_d) + p(\theta_B)p(w_i | \theta_B)]^{c(w,d)} \end{aligned}$$

- **ML Estimate:** $\Lambda^* = \arg \max_{\Lambda} p(d | \Lambda)$

$$\text{Subject to} \quad \sum_{i=1}^M p(w_i | \theta_d) = \sum_{i=1}^M p(w_i | \theta_B) = 1 \quad p(\theta_d) + p(\theta_B) = 1$$

Probabilistic Topic Models: Mixture Model Estimation

Probabilistic Topic Models: Mixture Model Estimation



Back to Factoring out Background Words

Text Mining Paper

d

... text mining...
is... clustering...
we.... Text.. the

$P(w | \theta_d)$

text 0.04 θ_d
mining 0.035
association 0.03
clustering 0.005
...
the 0.000001

$p(\theta_d) + p(\theta_B) = 1$

$P(\theta_d) = 0.5$

Topic
Choice

$P(\theta_B) = 0.5$

$p(w | \theta_B)$

the 0.03 θ_B
a 0.02
is 0.015
we 0.01
food 0.003
...
text 0.000006

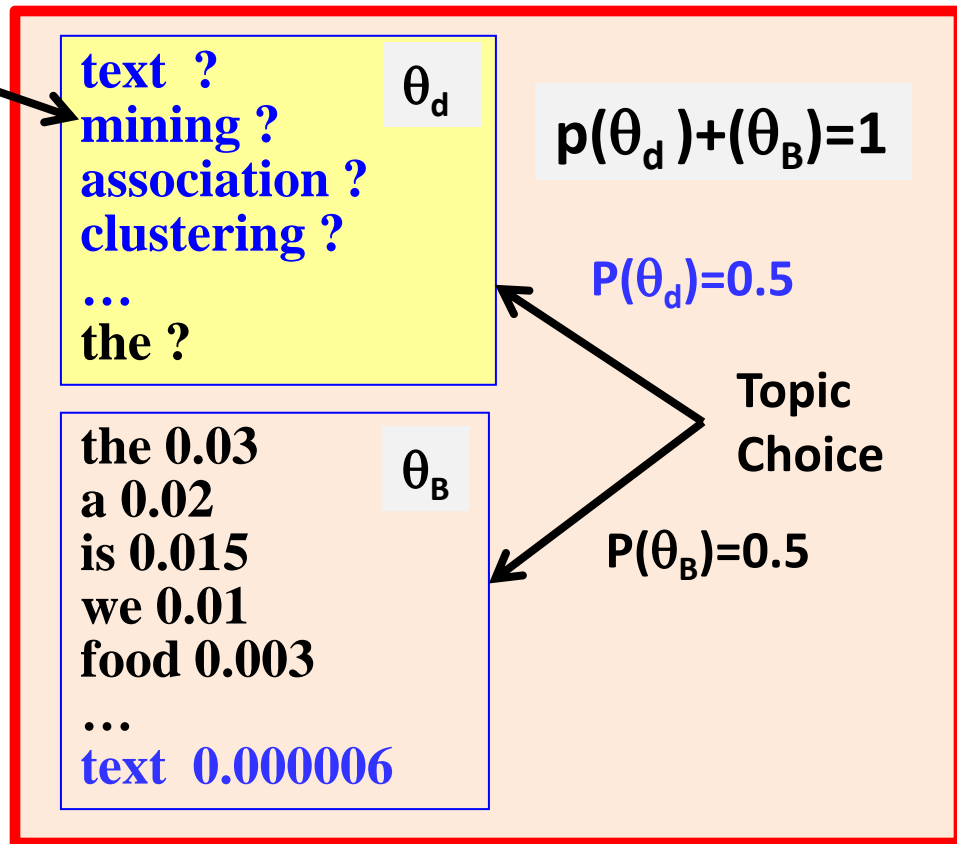
Estimation of One Topic: $P(w | \theta_d)$

Adjust θ_d to maximize $p(d | \Lambda)$
(all other parameters are known)

Would the ML estimate demote
background words in θ_d ?

d

... text mining...
is... clustering...
we.... Text.. the



Behavior of a Mixture Model

$d =$ text the

Likelihood:

$$\begin{aligned} P(\text{"text"}) &= p(\theta_d)p(\text{"text"} | \theta_d) + p(\theta_B)p(\text{"text"} | \theta_B) \\ &= 0.5 * p(\text{"text"} | \theta_d) + 0.5 * 0.1 \end{aligned}$$

$$P(\text{"the"}) = 0.5 * p(\text{"the"} | \theta_d) + 0.5 * 0.9$$

$$\begin{aligned} p(d | \Lambda) &= p(\text{"text"} | \Lambda) p(\text{"the"} | \Lambda) \\ &= [0.5 * p(\text{"text"} | \theta_d) + 0.5 * 0.1] \times \\ &\quad [0.5 * p(\text{"the"} | \theta_d) + 0.5 * 0.9] \end{aligned}$$

text ?
the ? θ_d

$P(\theta_d) = 0.5$

$P(\theta_B) = 0.5$

the 0.9
text 0.1 θ_B

How can we set $p(\text{"text"} | \theta_d)$ & $p(\text{"the"} | \theta_d)$ to maximize it?

Note that $p(\text{"text"} | \theta_d) + p(\text{"the"} | \theta_d) = 1$

“Collaboration” and “Competition” of θ_d and θ_B

$$\begin{aligned} p(d|\Lambda) &= p(\text{“text”}|\Lambda) p(\text{“the”}|\Lambda) \\ &= [0.5 * p(\text{“text”}|\theta_d) + 0.5 * 0.1] \times \\ &\quad [0.5 * p(\text{“the”}|\theta_d) + 0.5 * 0.9] \end{aligned}$$

Note that $p(\text{“text”}|\theta_d) + p(\text{“the”}|\theta_d) = 1$

If $x + y = \text{constant}$, then xy reaches maximum when $x = y$.

$$0.5 * p(\text{“text”}|\theta_d) + 0.5 * 0.1 = 0.5 * p(\text{“the”}|\theta_d) + 0.5 * 0.9$$

$$\Rightarrow p(\text{“text”}|\theta_d) = 0.9 \gg p(\text{“the”}|\theta_d) = 0.1 !$$

$d =$ text the

text ?
the ? θ_d

$P(\theta_d) = 0.5$

$P(\theta_B) = 0.5$

the 0.9
text 0.1 θ_B

Behavior 1: if $p(w1|\theta_B) > p(w2|\theta_B)$, then $p(w1|\theta_d) < p(w2|\theta_d)$

Response to Data Frequency

d =

text the

$$p(d|\Lambda) = [0.5 * p(\text{"text"}|\theta_d) + 0.5 * 0.1] \\ \times [0.5 * p(\text{"the"}|\theta_d) + 0.5 * 0.9]$$

$$\rightarrow p(\text{"text"}|\theta_d) = 0.9 \gg p(\text{"the"}|\theta_d) = 0.1 !$$

d' =

text the
the the
the ...the

$$p(d'|\Lambda) = [0.5 * p(\text{"text"}|\theta_d) + 0.5 * 0.1] \\ \times [0.5 * p(\text{"the"}|\theta_d) + 0.5 * 0.9] \\ \times [0.5 * p(\text{"the"}|\theta_d) + 0.5 * 0.9] \\ \times [0.5 * p(\text{"the"}|\theta_d) + 0.5 * 0.9]$$

...

What if we increase $p(\theta_B)$?

$$\times [0.5 * p(\text{"the"}|\theta_d) + 0.5 * 0.9]$$

What's the optimal solution now? $p(\text{"the"}|\theta_d) > 0.1$? or $p(\text{"the"}|\theta_d) < 0.1$?

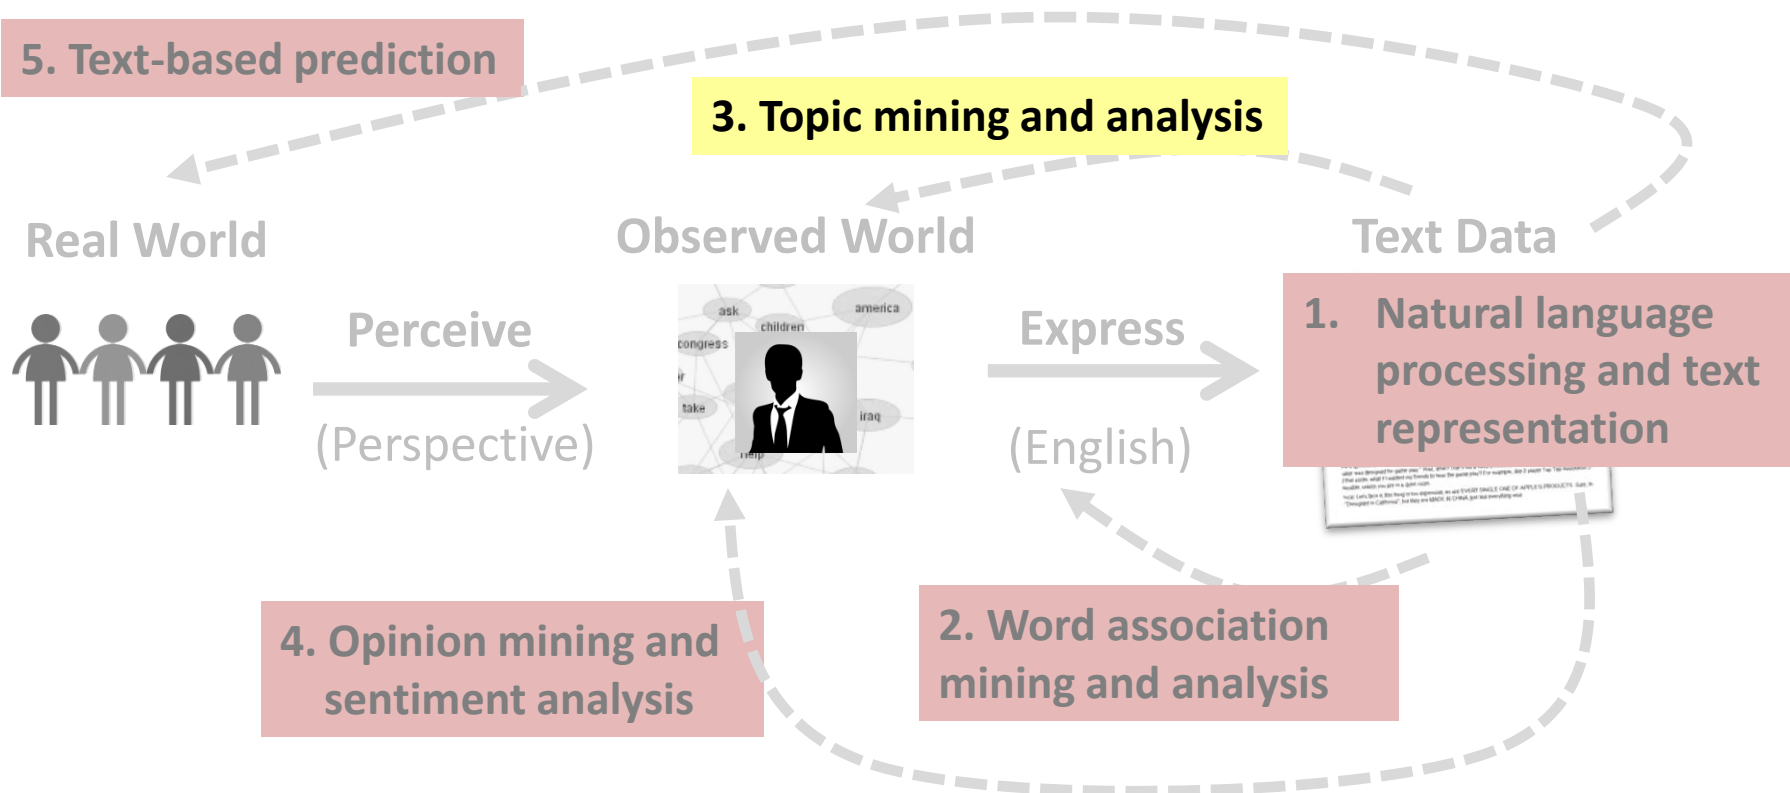
Behavior 2: high frequency words get higher $p(w|\theta_d)$

Summary

- General behavior of a mixture model:
 - Every component model attempts to assign high probabilities to highly frequent words in the data (to “collaboratively maximize likelihood”)
 - Different component models tend to “bet” high probabilities on different words (to avoid “competition” or “waste of probability”)
 - The probability of choosing each component “regulates” the collaboration/competition between the component models
- Fixing one component to a background word distribution (i.e., background language model):
 - Helps “get rid of background words” in other component
 - Is an example of imposing a prior on the model parameters (prior = one model must be exactly the same as the background LM)

Probabilistic Topic Models: Expectation-Maximization Algorithm

Probabilistic Topic Models: Expectation-Maximization (EM) Algorithm



Estimation of One Topic: $P(w | \theta_d)$

How to set θ_d to maximize $p(d | \Lambda)$?
(all other parameters are known)

d

... text mining...
is... clustering...
we.... Text.. the



text ?
mining ?
association ?
clustering ?
...
the ?

θ_d

$$p(\theta_d) + p(\theta_B) = 1$$


$$P(\theta_d) = 0.5$$

Topic
Choice

$$P(\theta_B) = 0.5$$

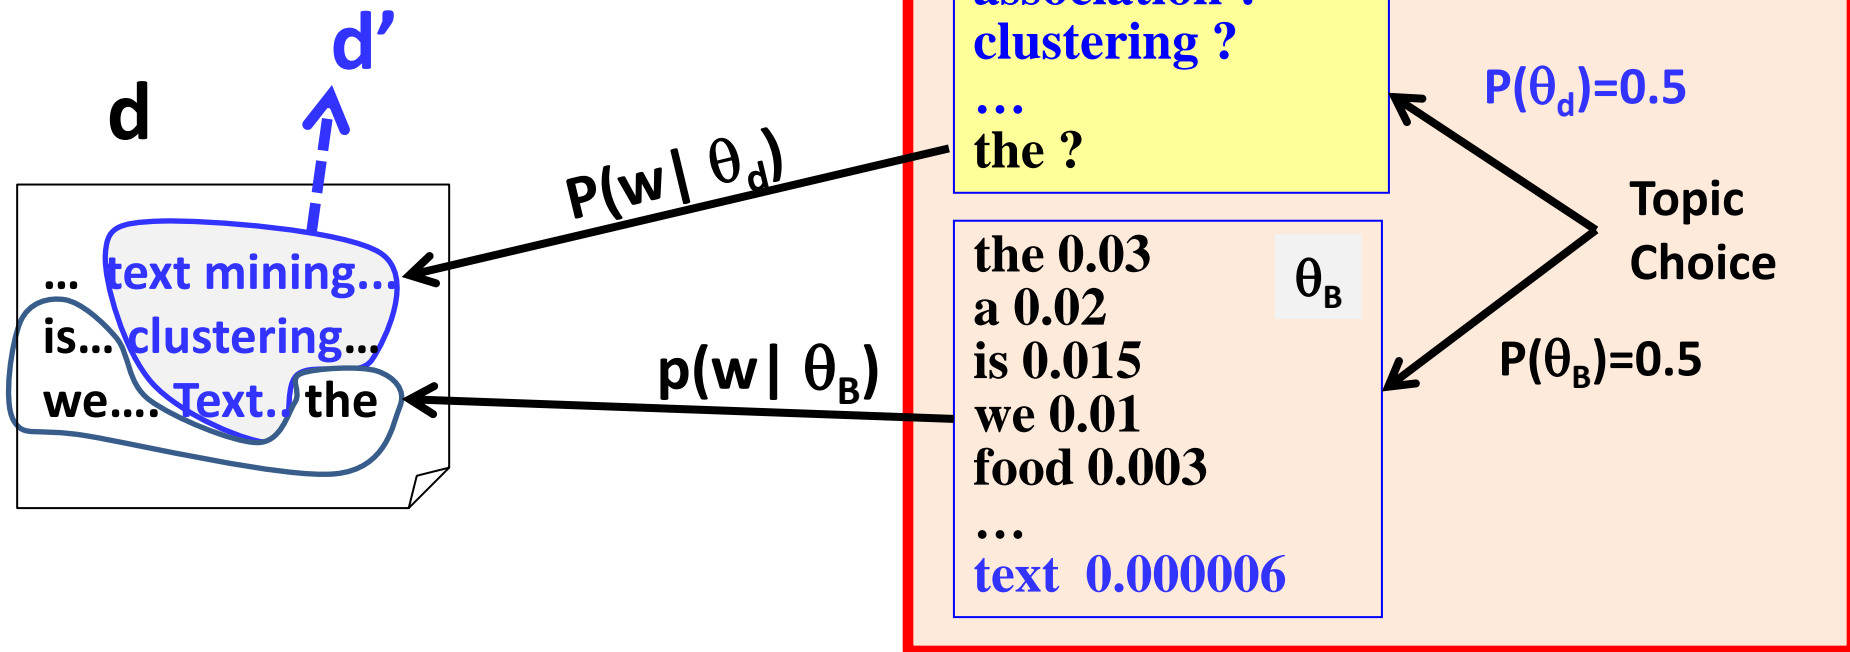
the 0.03
a 0.02
is 0.015
we 0.01
food 0.003
...
text 0.000006

θ_B



If we know which word is from which distribution...

$$p(w_i | \theta_d) = \frac{c(w_i, d')}{\sum_{w' \in V} c(w', d')}$$



Given all the parameters, infer the distribution a word is from...

Is “**text**” more likely from θ_d or θ_B ?

From θ_d ($Z=0$)?

$p(\theta_d)p(\text{"text"}|\theta_d)$

From θ_B ($Z=1$)?

$p(\theta_B)p(\text{"text"}|\theta_B)$

$P(w|\theta_d)$

text 0.04
mining 0.035
association 0.03
clustering 0.005
...
the 0.000001

θ_d

$p(w|\theta_B)$

the 0.03
a 0.02
is 0.015
we 0.01
food 0.003
...
text 0.000006

θ_B

$p(\theta_d)+p(\theta_B)=1$

$P(\theta_d)=0.5$

$P(\theta_B)=0.5$

Topic
Choice

$p(z = 0 | w = \text{"text"}) =$

$$\frac{p(\theta_d)p(\text{"text"}|\theta_d)}{p(\theta_d)p(\text{"text"}|\theta_d) + p(\theta_B)p(\text{"text"}|\theta_B)}$$

The Expectation-Maximization (EM) Algorithm

Hidden Variable:

$z \in \{0, 1\}$

	z
the	1
paper	1
presents	1
a	1
text	0
mining	0
algorithm	0
for	1
clustering	0
...	...

Initialize $p(w|\theta_d)$ with random values.

Then iteratively improve it using E-step & M-step.

Stop when likelihood doesn't change.

$$p^{(n)}(z = 0 | w) = \frac{p(\theta_d)p^{(n)}(w | \theta_d)}{p(\theta_d)p^{(n)}(w | \theta_d) + p(\theta_B)p(w | \theta_B)}$$

E-step

How likely w is from θ_d

$$p^{(n+1)}(w | \theta_d) = \frac{c(w, d)p^{(n)}(z = 0 | w)}{\sum_{w' \in V} c(w', d)p^{(n)}(z = 0 | w')}$$

M-step

EM Computation in Action

E-step
$$p^{(n)}(z = 0 | w) = \frac{p(\theta_d)p^{(n)}(w | \theta_d)}{p(\theta_d)p^{(n)}(w | \theta_d) + p(\theta_B)p(w | \theta_B)}$$

M-step
$$p^{(n+1)}(w | \theta_d) = \frac{c(w, d)p^{(n)}(z = 0 | w)}{\sum_{w' \in V} c(w', d)p^{(n)}(z = 0 | w')}$$

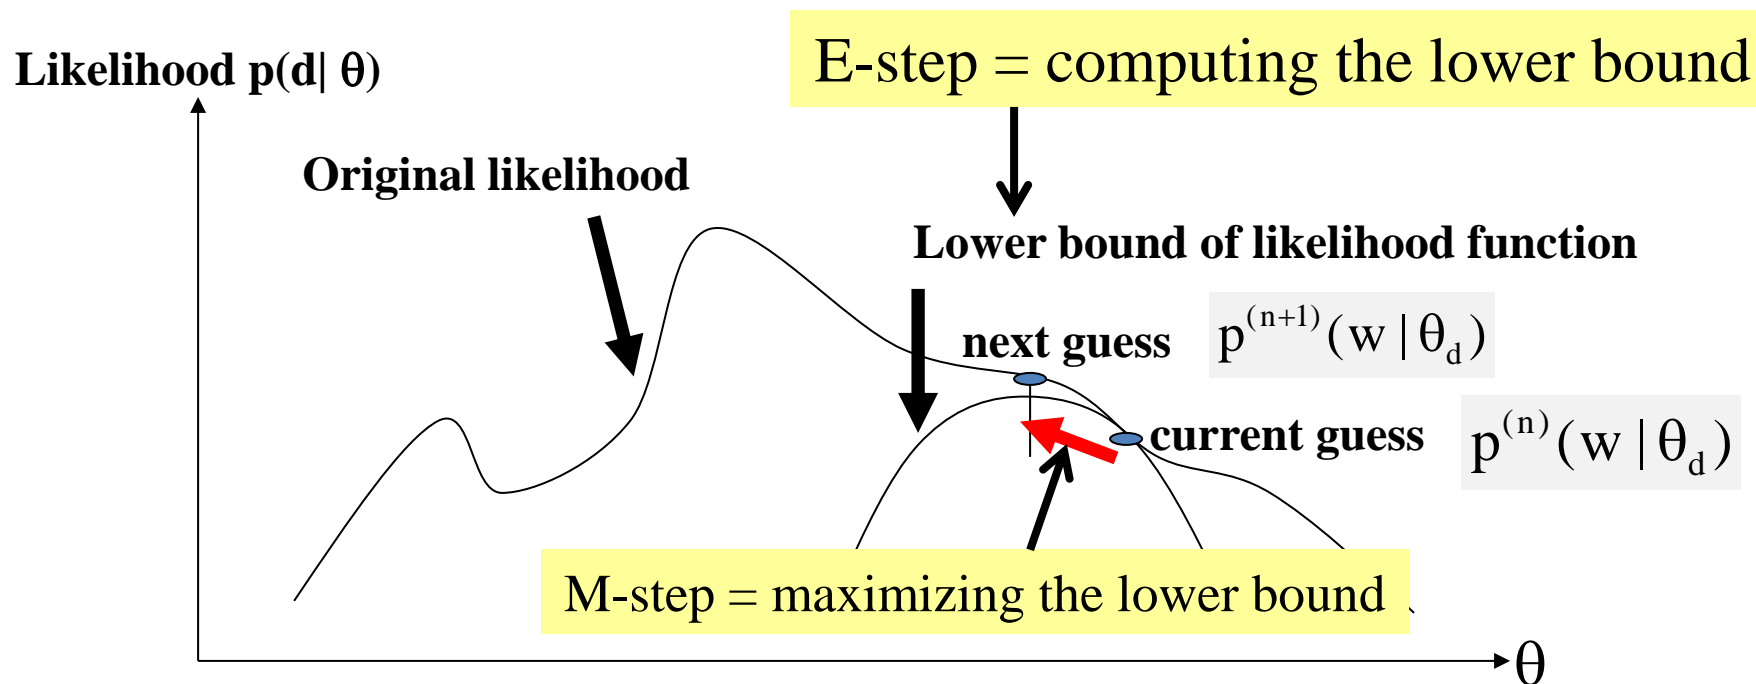
Assume
 $p(\theta_d) = p(\theta_B) = 0.5$
 and $p(w | \theta_B)$ is known

Word	#	$p(w \theta_B)$	Iteration 1		Iteration 2		Iteration 3	
			$P(w \theta)$	$p(z=0 w)$	$P(w \theta)$	$P(z=0 w)$	$P(w \theta)$	$P(z=0 w)$
The	4	0.5	0.25	0.33	0.20	0.29	0.18	0.26
Paper	2	0.3	0.25	0.45	0.14	0.32	0.10	0.25
Text	4	0.1	0.25	0.71	0.44	0.81	0.50	0.93
Mining	2	0.1	0.25	0.71	0.22	0.69	0.22	0.69
Log-Likelihood			-16.96		-16.13		-16.02	

Likelihood increasing

“By products”: Are they also useful?

EM As Hill-Climbing → Converge to Local Maximum

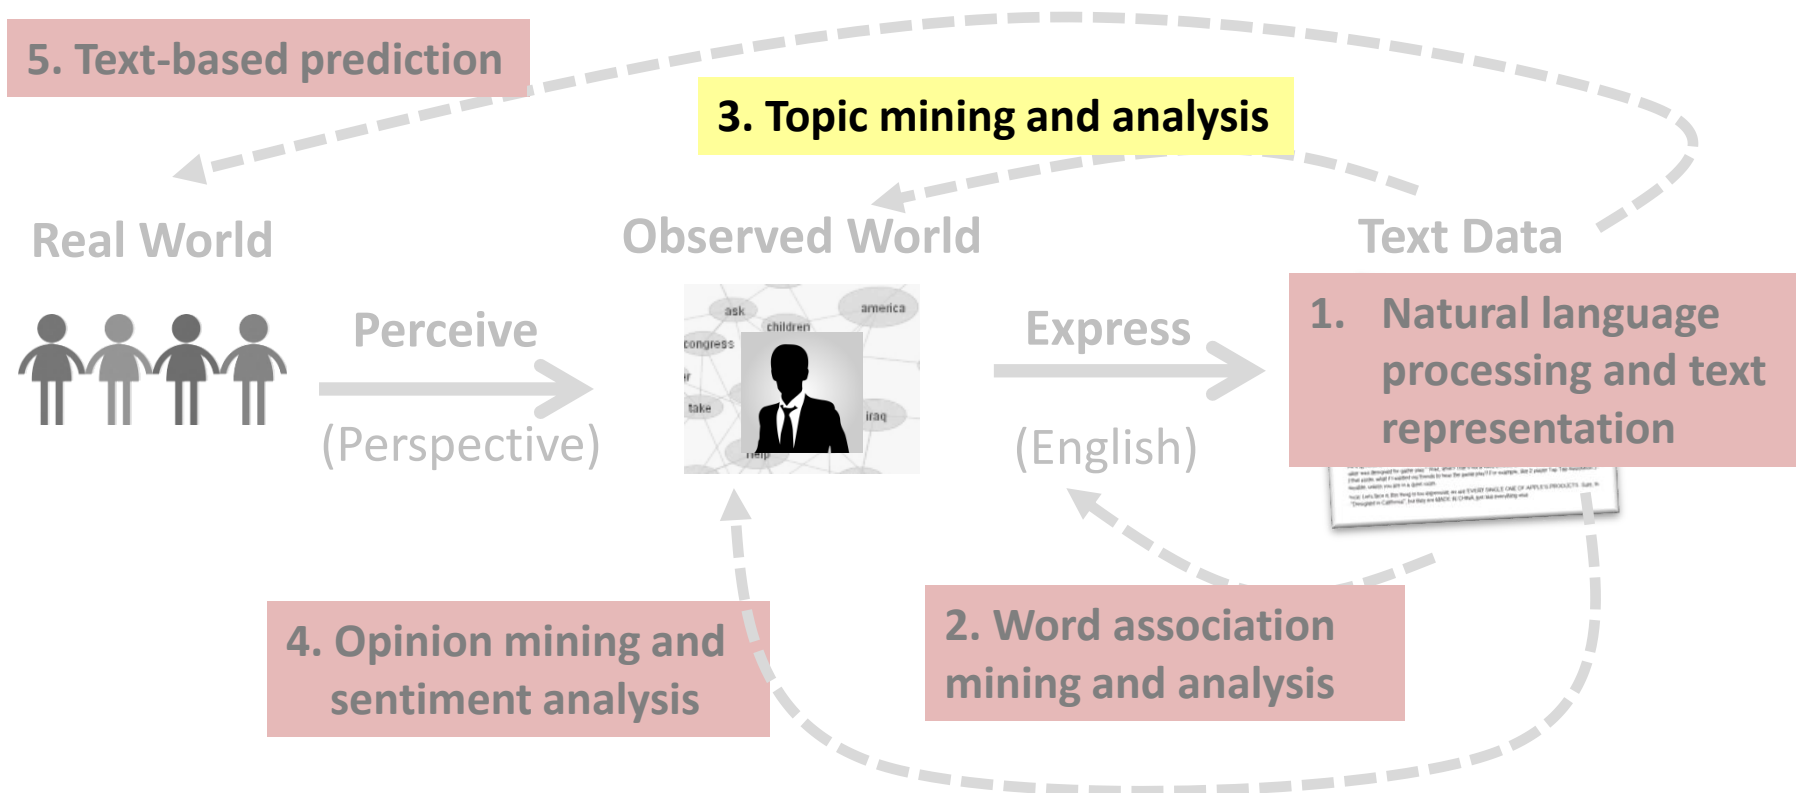


Summary

- Expectation-Maximization (EM) algorithm
 - General algorithm for computing ML estimate of mixture models
 - Hill-climbing, so can only converge to a local maximum (depending on initial points)
- E-step: “augment” data by predicting values of useful hidden variables
- M-step: exploit the “augmented data” to improve estimate of parameters (“improve” is guaranteed in terms of likelihood)
- “Data augmentation” is probabilistic ➔ Split counts of events probabilistically

Probabilistic Latent Semantic Analysis (PLSA)

Probabilistic Latent Semantic Analysis (PLSA)



Document as a Sample of Mixed Topics

Topic θ_1

government 0.3
response 0.2
...

Topic θ_2

city 0.2
new 0.1
orleans 0.05
...

...

Topic θ_k

donate 0.1
relief 0.05
help 0.02
...

Background θ_B

the 0.04
a 0.03
...

Blog article about “Hurricane Katrina”

[Criticism of government response to the hurricane primarily consisted of criticism of its response to the approach of the storm and its aftermath, specifically in the delayed response] to the [flooding of New Orleans. ... 80% of the 1.3 million residents of the greater New Orleans metropolitan area evacuated] ... [Over seventy countries pledged monetary donations or other assistance]. ...

Many applications are possible if we can “decode” the topics in text...

Mining Multiple Topics from Text

OUTPUT: $\{ \theta_1, \dots, \theta_k \}, \{ \pi_{i1}, \dots, \pi_{ik} \}$

INPUT: C, k, V

Text Data

θ_1

sports 0.02
game 0.01
basketball 0.005
football 0.004
...

θ_2

travel 0.05
attraction 0.03
trip 0.01
...

...

θ_k

science 0.04
scientist 0.03
spaceship 0.006
...

Doc 1

30%

π_{11}

Doc 2

$\pi_{21}=0\%$

...

Doc N

$\pi_{N1}=0\%$

12%

π_{12}

π_{22}

π_{N2}

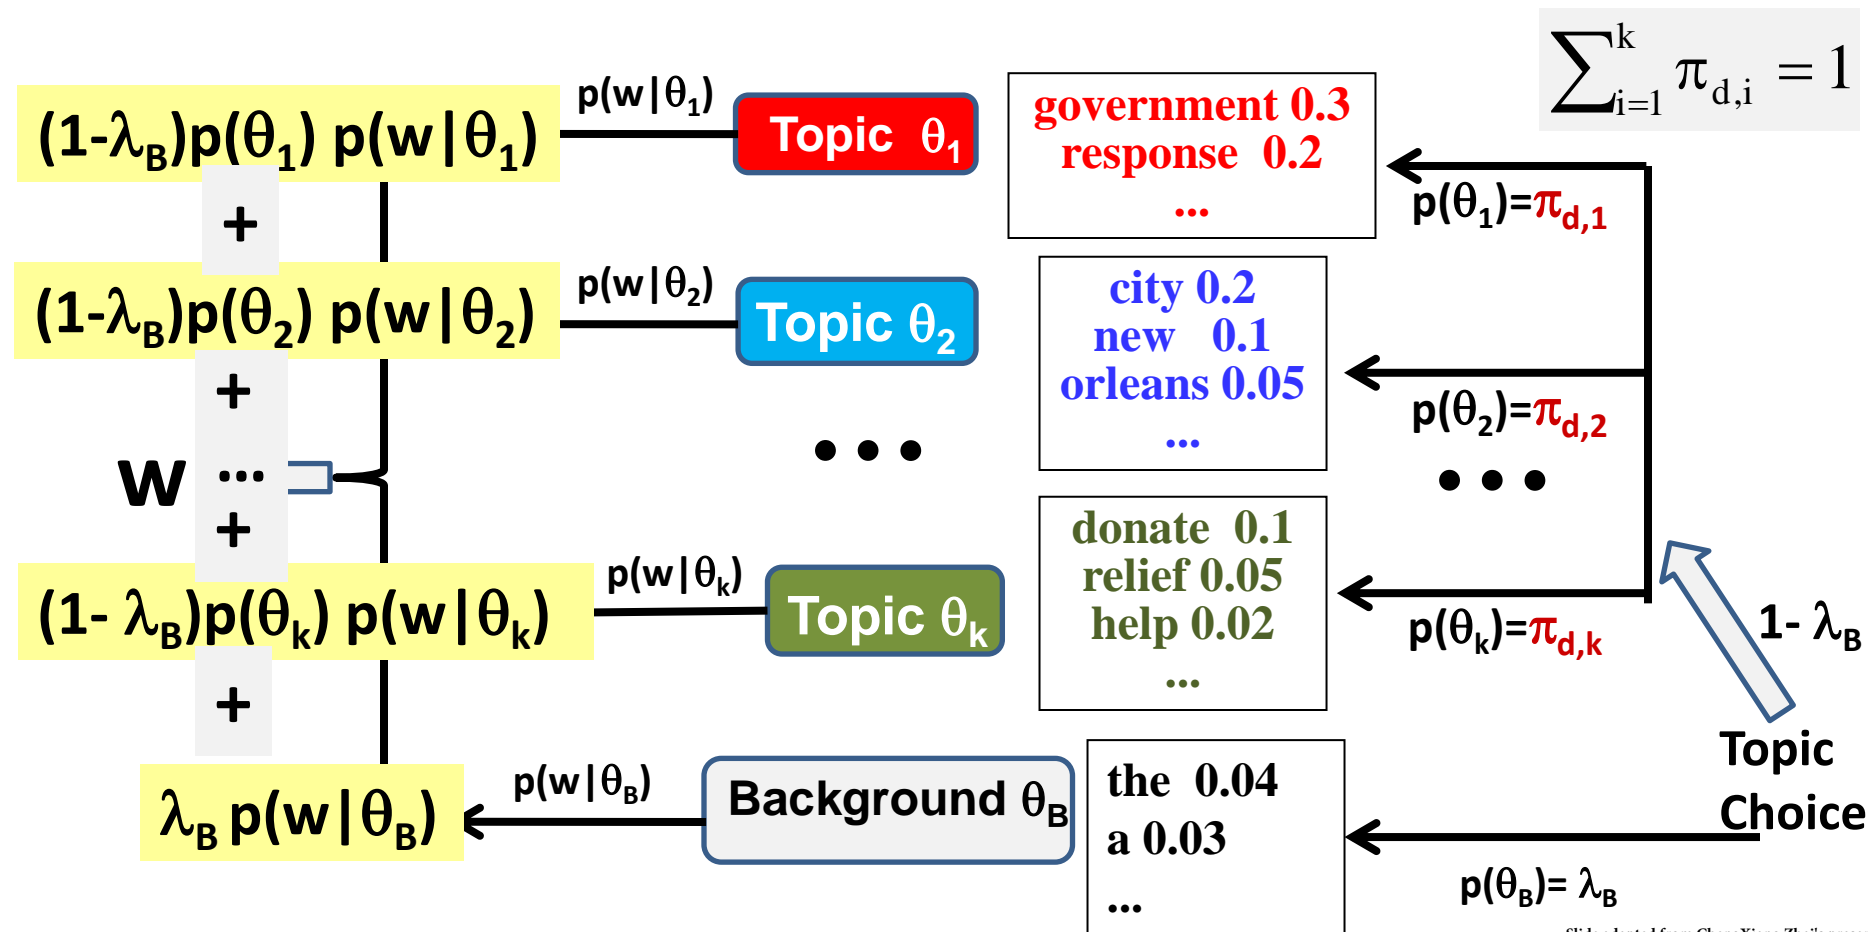
8%

π_{1k}

π_{2k}

π_{Nk}

Generating Text with Multiple Topics: $p(w)=?$



Probabilistic Latent Semantic Analysis (PLSA)

Percentage of

background words
(known)

Background
LM (known)

Coverage of topic θ_j in doc d

Prob. of word w in topic θ_j

$$p_d(w) = \lambda_B p(w | \theta_B) + (1 - \lambda_B) \sum_{j=1}^k \pi_{d,j} p(w | \theta_j)$$

$$\log p(d) = \sum_{w \in V} c(w, d) \log [\lambda_B p(w | \theta_B) + (1 - \lambda_B) \sum_{j=1}^k \pi_{d,j} p(w | \theta_j)]$$

$$\log p(C | \Lambda) = \sum_{d \in C} \sum_{w \in V} c(w, d) \log [\lambda_B p(w | \theta_B) + (1 - \lambda_B) \sum_{j=1}^k \pi_{d,j} p(w | \theta_j)]$$

Unknown Parameters: $\Lambda = (\{\pi_{d,j}\}, \{\theta_j\})$, $j=1, \dots, k$

How many unknown parameters are there in total?

ML Parameter Estimation

$$p_d(w) = \lambda_B p(w | \theta_B) + (1 - \lambda_B) \sum_{j=1}^k \pi_{d,j} p(w | \theta_j)$$

$$\log p(d) = \sum_{w \in V} c(w, d) \log [\lambda_B p(w | \theta_B) + (1 - \lambda_B) \sum_{j=1}^k \pi_{d,j} p(w | \theta_j)]$$

$$\log p(C | \Lambda) = \sum_{d \in C} \sum_{w \in V} c(w, d) \log [\lambda_B p(w | \theta_B) + (1 - \lambda_B) \sum_{j=1}^k \pi_{d,j} p(w | \theta_j)]$$

Constrained Optimization: $\Lambda^* = \arg \max_{\Lambda} p(C | \Lambda)$

$$\forall j \in [1, k], \sum_{i=1}^M p(w_i | \theta_j) = 1$$

$$\forall d \in C, \sum_{j=1}^k \pi_{d,j} = 1$$

EM Algorithm for PLSA: E-Step

Hidden Variable (=topic indicator): $z_{d,w} \in \{B, 1, 2, \dots, k\}$

Probability that **w in doc d** is generated from **topic θ_j**

$$p(z_{d,w} = j) = \frac{\pi_{d,j}^{(n)} p^{(n)}(w | \theta_j)}{\sum_{j'=1}^k \pi_{d,j'}^{(n)} p^{(n)}(w | \theta_{j'})}$$

Use of Bayes Rule

$$p(z_{d,w} = B) = \frac{\lambda_B p(w | \theta_B)}{\lambda_B p(w | \theta_B) + (1 - \lambda_B) \sum_{j=1}^k \pi_{d,j}^{(n)} p^{(n)}(w | \theta_j)}$$

Probability that **w in doc d** is generated from **background θ_B**

EM Algorithm for PLSA: M-Step

Hidden Variable (=topic indicator): $z_{d,w} \in \{B, 1, 2, \dots, k\}$

Re-estimated **probability** of doc d covering topic θ_j

ML Estimate based on
“allocated” word
counts to topic θ_j

$$\pi_{d,j}^{(n+1)} = \frac{\sum_{w \in V} c(w, d)(1 - p(z_{d,w} = B))p(z_{d,w} = j)}{\sum_{j'} \sum_{w \in V} c(w, d)(1 - p(z_{d,w} = B))p(z_{d,w} = j')}$$

$$p^{(n+1)}(w | \theta_j) = \frac{\sum_{d \in C} c(w, d)(1 - p(z_{d,w} = B))p(z_{d,w} = j)}{\sum_{w' \in V} \sum_{d \in C} c(w', d)(1 - p(z_{d,w'} = B))p(z_{d,w'} = j)}$$

Re-estimated **probability** of word w for topic θ_j

Computation of the EM Algorithm

- Initialize all unknown parameters randomly
- Repeat until likelihood converges

– E-step $p(z_{d,w} = j) \propto \pi_{d,j}^{(n)} p^{(n)}(w | \theta_j)$ $\sum_{j=1}^k p(z_{d,w} = j) = 1$

$p(z_{d,w} = B) \propto \lambda_B p(w | \theta_B) \leftarrow$

– M-step

What's the normalizer for this one?

$$\pi_{d,j}^{(n+1)} \propto \sum_{w \in V} c(w, d) (1 - p(z_{d,w} = B)) p(z_{d,w} = j)$$

$$\forall d \in C, \sum_{j=1}^k \pi_{d,j} = 1$$

$$p^{(n+1)}(w | \theta_j) \propto \sum_{d \in C} c(w, d) (1 - p(z_{d,w} = B)) p(z_{d,w} = j)$$

$$\forall j \in [1, k], \sum_{w \in V} p(w | \theta_j) = 1$$

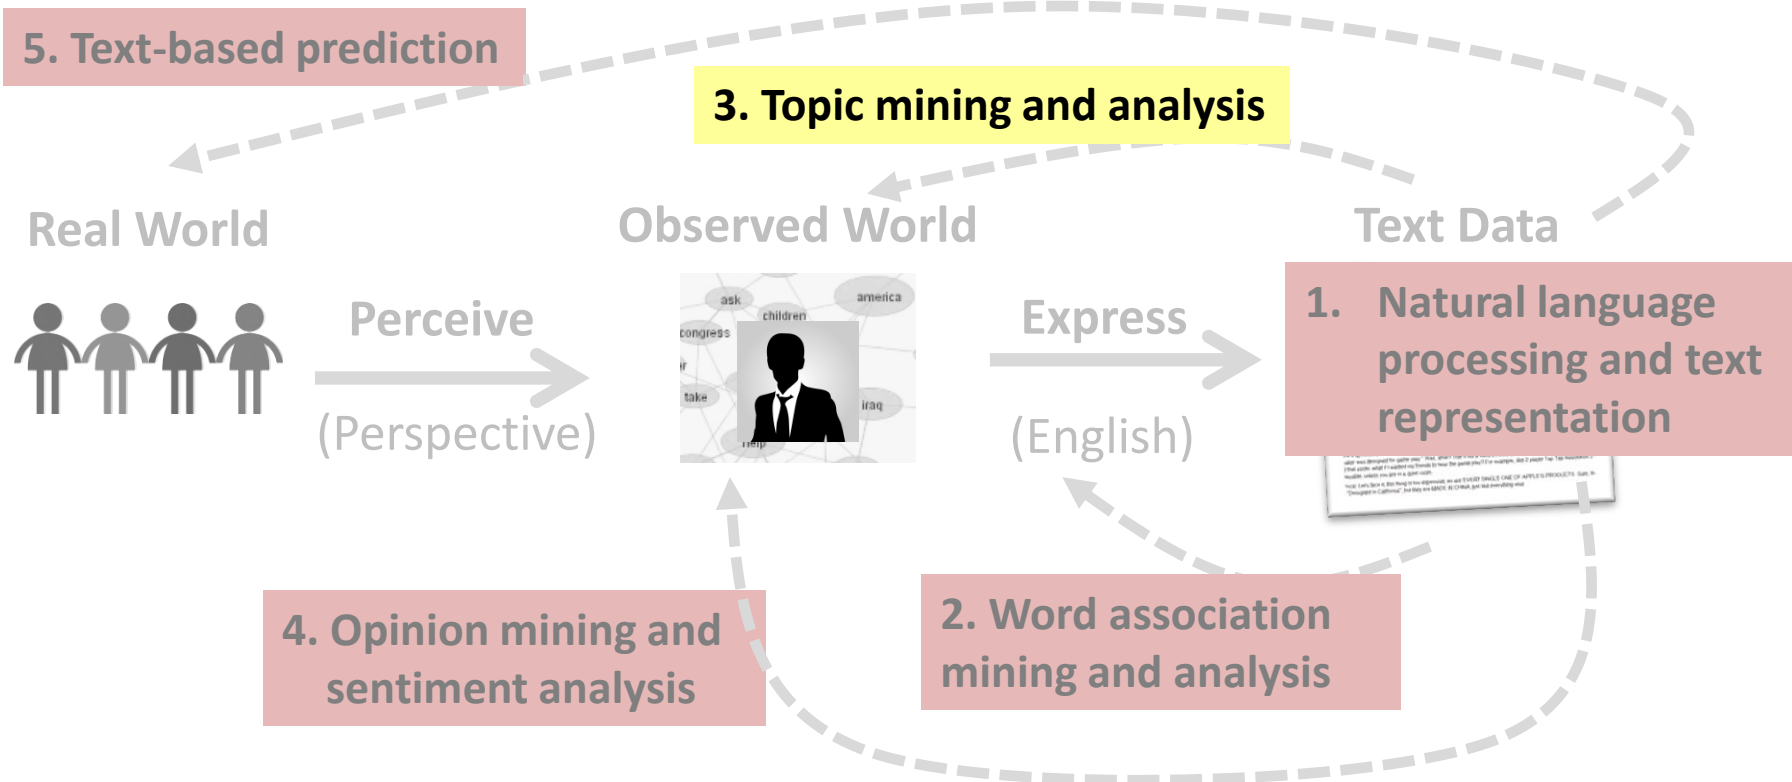
In general, accumulate counts, and then normalize

Summary

- PLSA = mixture model with k unigram LMs (k topics)
- Adding a pre-determined background LM helps discover discriminative topics
- ML estimate “discovers” topical knowledge from text data
 - k word distributions (k topics)
 - proportion of each topic in each document
- The output can enable many applications!
 - Clustering of terms and docs (treat each topic as a cluster)
 - Further associate topics with different contexts (e.g., time periods, locations, authors, sources, etc.)

Latent Dirichlet Allocation (LDA)

Latent Dirichlet Allocation (LDA)



Extensions of PLSA

- PLSA with prior knowledge → User-controlled PLSA
- PLSA as a generative model → Latent Dirichlet Allocation

PLSA with Prior Knowledge

- Users may have expectations about which topics to analyze:
 - We expect to see “retrieval models” as a topic in IR
 - We want to see aspects such as “battery” and “memory” for opinions about a laptop
- Users may have knowledge about what topics are (or are NOT) covered in a document
 - Tags = topics → A doc can only be generated using topics corresponding to the tags assigned to the document
- We can incorporate such knowledge as priors of PLSA model

Maximum a Posteriori (MAP) Estimate

$$\Lambda^* = \arg \max_{\Lambda} p(\Lambda) p(Data | \Lambda)$$

- We may use $p(\Lambda)$ to encode all kinds of preferences and constraints, e.g.,
 - $p(\Lambda) > 0$ if and only if one topic is precisely “background”: $p(w | \theta_B)$
 - $p(\Lambda) > 0$ if and only if for a particular doc d , $\pi_{d,3} = 0$ and $\pi_{d,1} = 1/2$
 - $p(\Lambda)$ favors a Λ with topics that assign high probabilities to some particular words
- The MAP estimate (with conjugate prior) can be computed using a similar EM algorithm to the ML estimate with smoothing to reflect prior preferences

EM Algorithm with Conjugate Prior on $p(w | \theta_j)$

$$p(z_{d,w} = j) = \frac{\pi_{d,j}^{(n)} p^{(n)}(w | \theta_j)}{\sum_{j'=1}^k \pi_{d,j'}^{(n)} p^{(n)}(w | \theta_{j'})}$$

$$p(z_{d,w} = B) = \frac{\lambda_B p(w | \theta_B)}{\lambda_B p(w | \theta_B) + (1 - \lambda_B) \sum_{j=1}^k \pi_{d,j}^{(n)} p^{(n)}(w | \theta_j)}$$

$$\pi_{d,j}^{(n+1)} = \frac{\sum_{w \in V} c(w, d) (1 - p(z_{d,w} = B)) p(z_{d,w} = j)}{\sum_{j'} \sum_{w \in V} c(w, d) (1 - p(z_{d,w} = B)) p(z_{d,w} = j')}$$

$$p^{(n+1)}(w | \theta_j) = \frac{\sum_{d \in C} c(w, d) (1 - p(z_{d,w} = B)) p(z_{d,w} = j) + \mu p(w | \theta'_j)}{\sum_{w' \in V} \sum_{d \in C} c(w', d) (1 - p(z_{d,w'} = B)) p(z_{d,w'} = j) + \mu}$$

Prior: $p(w | \theta'_j)$

battery 0.5
life 0.5

Pseudo counts of w
from prior θ'

What if $\mu=0$? What if $\mu=+\infty$?

Sum of all pseudo counts

We may also set any parameter to a constant (including 0) as needed

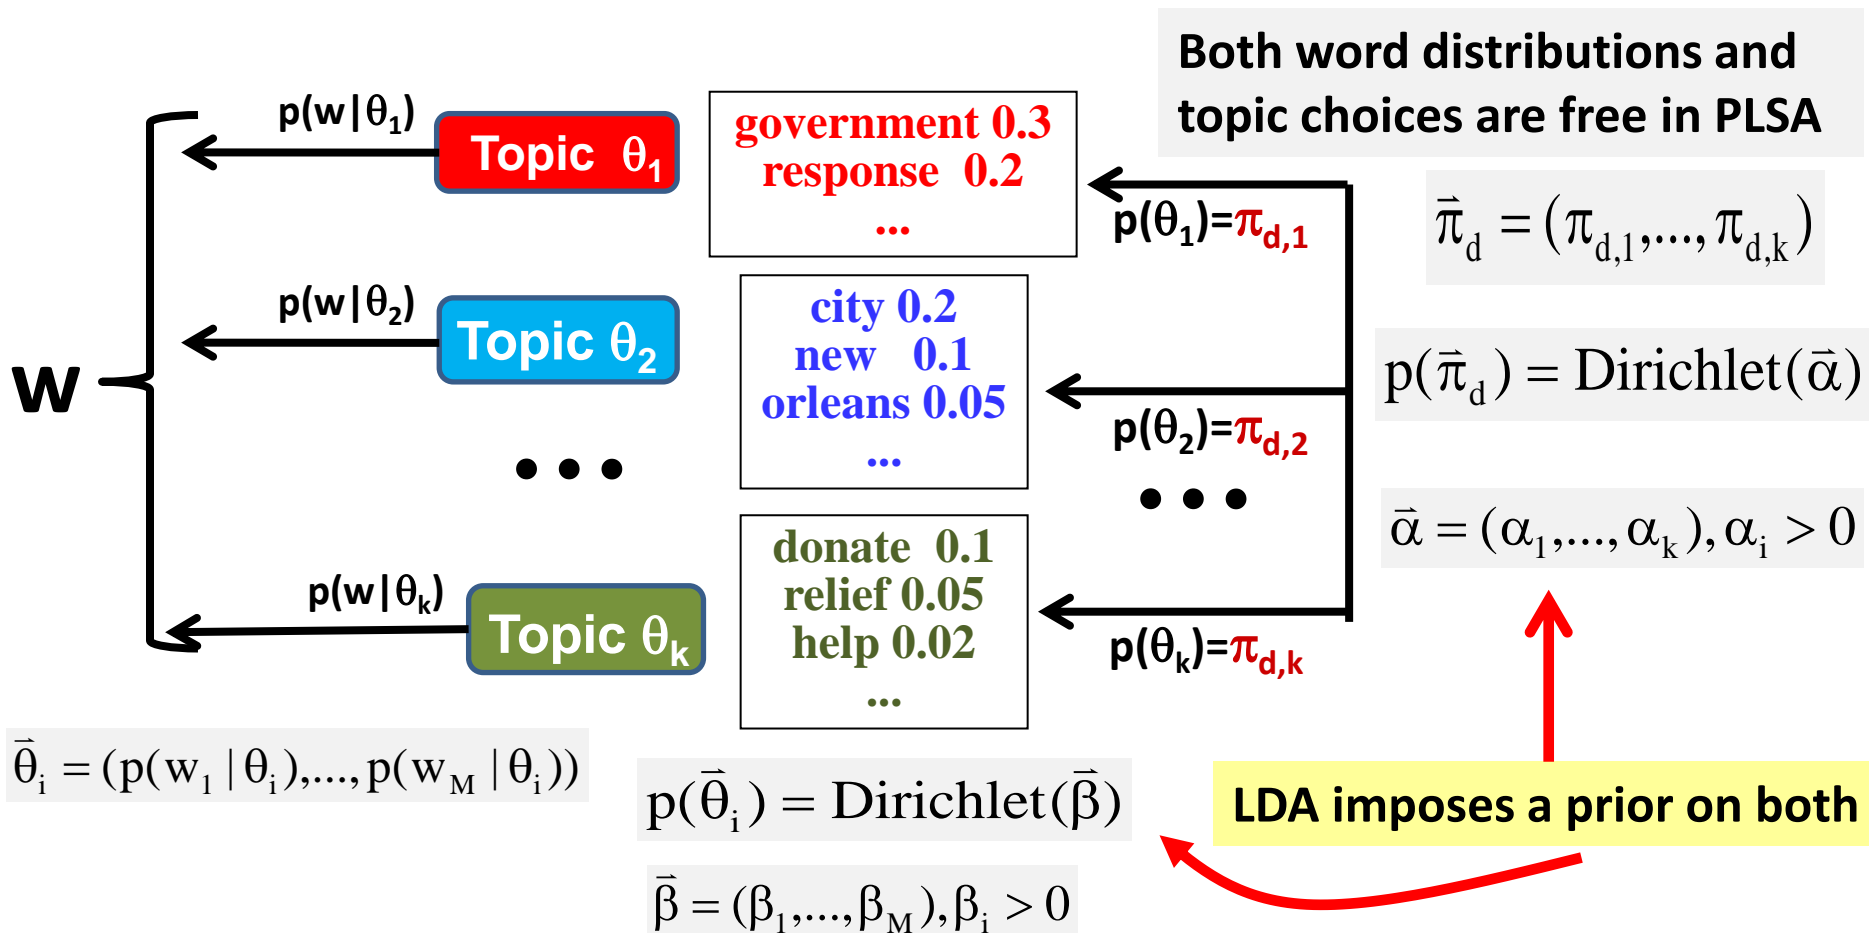
Deficiency of PLSA

- Not a generative model
 - Can't compute probability of a new document
 - Heuristic workaround is possible, though
- Many parameters → high complexity of models
 - Many local maxima
 - Prone to overfitting
- Not necessarily a problem for text mining (only interested in fitting the “training” documents)

Latent Dirichlet Allocation (LDA)

- Make PLSA a generative model by imposing a Dirichlet prior on the model parameters →
 - LDA = Bayesian version of PLSA
 - Parameters are regularized
- Can achieve the same goal as PLSA for text mining purposes
 - Topic coverage and topic word distributions can be inferred using Bayesian inference

PLSA → LDA



Likelihood Functions for PLSA vs. LDA

PLSA

$$p_d(w | \{\theta_j\}, \{\pi_{d,j}\}) = \sum_{j=1}^k \pi_{d,j} p(w | \theta_j) \quad \leftarrow \text{Core assumption in all topic models}$$

$$\log p(d | \{\theta_j\}, \{\pi_{d,j}\}) = \sum_{w \in V} c(w, d) \log \left[\sum_{j=1}^k \pi_{d,j} p(w | \theta_j) \right]$$

$$\log p(C | \{\theta_j\}, \{\pi_{d,j}\}) = \sum_{d \in C} \log p(d | \{\theta_j\}, \{\pi_{d,j}\})$$

LDA

$$p_d(w | \{\theta_j\}, \{\pi_{d,j}\}) = \sum_{j=1}^k \pi_{d,j} p(w | \theta_j)$$

$$\log p(d | \vec{\alpha}, \{\theta_j\}) = \int \left[\sum_{w \in V} c(w, d) \log \left[\sum_{j=1}^k \pi_{d,j} p(w | \theta_j) \right] \right] p(\vec{\pi}_d | \vec{\alpha}) d\vec{\pi}_d$$

$$\log p(C | \vec{\alpha}, \vec{\beta}) = \int \sum_{d \in C} \log p(d | \vec{\alpha}, \{\theta_j\}) \prod_{j=1}^k p(\theta_j | \vec{\beta}) d\theta_1 \dots d\theta_k$$

Added by LDA

PLSA component

Parameter Estimation and Inferences in LDA

- Parameters can be estimated using ML estimator

$$(\hat{\vec{\alpha}}, \hat{\vec{\beta}}) = \arg \max_{\vec{\alpha}, \vec{\beta}} \log p(C | \vec{\alpha}, \vec{\beta})$$

How many parameters in LDA vs. PLSA?

- However, $\{\theta_j\}$ and $\{\pi_{d,j}\}$ must now be computed using posterior inference
 - Computationally intractable
 - Must resort to approximate inference
 - Many different inference methods are available

Summary of Probabilistic Topic Models

- Probabilistic topic models provide a general principled way of mining and analyzing topics in text with many applications
- Basic task setup:
 - Input: Text data
 - Output: k topics + proportions of these topics covered in each document
- PLSA is the basic topic model, often adequate for most applications
- LDA improves over PLSA by imposing priors
 - Theoretically more appealing
 - Practically, LDA and PLSA perform similarly for many tasks

Suggested Readings

- Blei, D. 2012. “Probabilistic Topic Models.” *Communications of the ACM* 55 (4): 77–84. doi: 10.1145/2133806.2133826.
- Qiaozhu Mei, Xuehua Shen, and ChengXiang Zhai. “Automatic Labeling of Multinomial Topic Models.” *Proceedings of ACM KDD 2007*, pp. 490-499, DOI=10.1145/1281192.1281246.
- Yue Lu, Qiaozhu Mei, and Chengxiang Zhai. 2011. Investigating task performance of probabilistic topic models: an empirical study of PLSA and LDA. *Information Retrieval*, 14, 2 (April 2011), 178-203. DOI=10.1007/s10791-010-9141-9.